StarNEig User's Guide

v0.1 branch

Generated by Doxygen 1.8.17

Thu Nov 19 2020 18:46:59

Chapter 1

Introduction

StarNEig library aims to provide a complete task-based software stack for solving **dense nonsymmetric** (generalized) eigenvalue problems. The library is built on top of the `StarPU` runtime system and targets both **shared memory** and **distributed memory** machines. Some components of the library support **GPUs**.

The four main components of the library are:

- **Hessenberg(-triangular) reduction**: A dense matrix (or a dense matrix pair) is reduced to upper Hessenberg (or Hessenberg-triangular) form.

- **Schur reduction (QR/QZ algorithm)**: A upper Hessenberg matrix (or a Hessenberg-triangular matrix pair) is reduced to (generalized) Schur form. The (generalized) eigenvalues can be determined from the diagonal blocks.

- **Eigenvalue reordering and deflating subspaces**: Reorders a user-selected set of (generalized) eigenvalues to the upper left corner of an updated (generalized) Schur form.

- **Computation of eigenvectors**: Computes (generalized) eigenvectors for a user-selected set of (generalized) eigenvalues.

Please cite the following article when referring to StarNEig:

Mirko Myllykoski, Carl Christian Kjelgaard Mikkelsen: *Task-based, GPU-accelerated and Robust Library for Solving Dense Nonsymmetric Eigenvalue Problems*, Concurrency and Computation: Practice and Experience, 2020, doi: `10.1002/cpe.5915`

Please see publications and authors.

Performance

Performance comparisons against MAGMA (GPU) and ScaLAPACK (distributed memory), and strong scalability on shared and distributed memory machines:
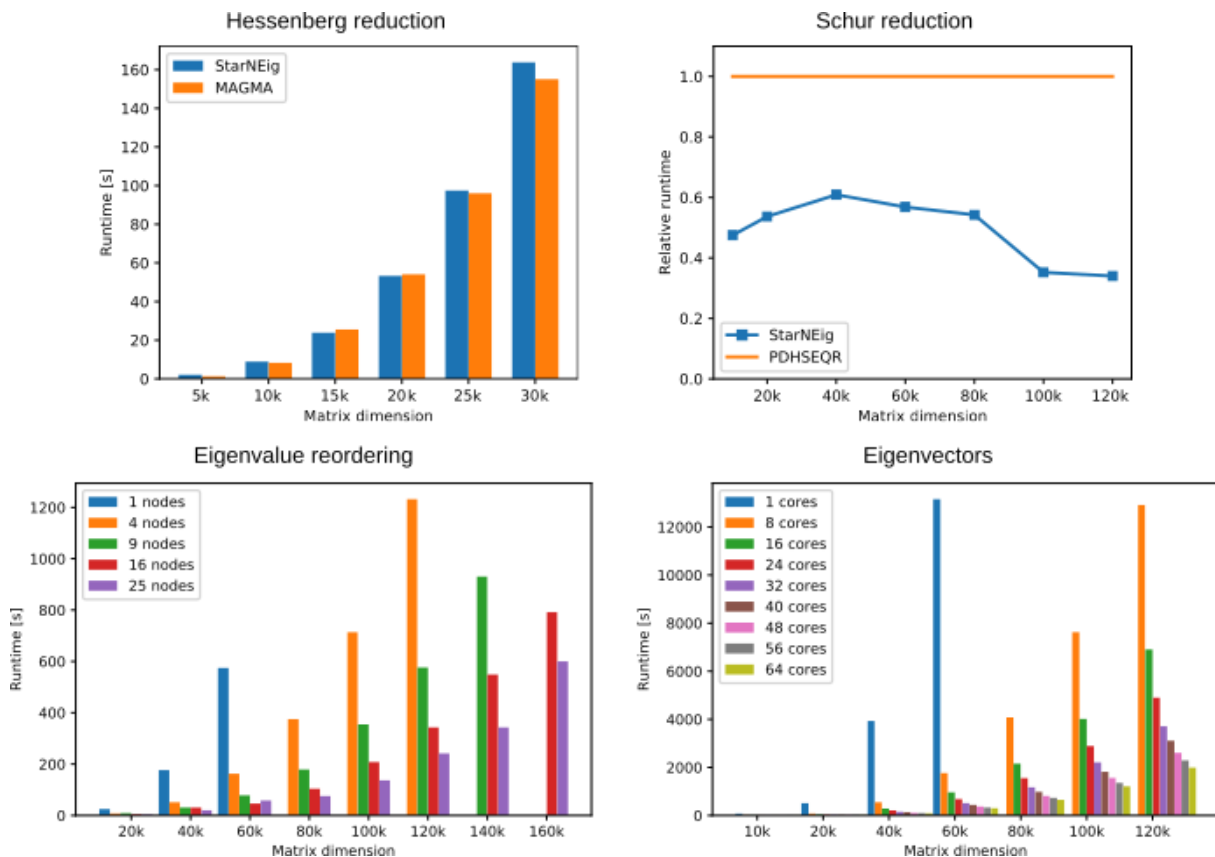
Also, see following publications:

- Mirko Myllykoski: *A Task-based Multi-shift QR/QZ Algorithm with Aggressive Early Deflation*, arXiv‐:2007.03576

- Mirko Myllykoski, Carl Christian Kjelgaard Mikkelsen: *Task-based, GPU-accelerated and Robust Library for Solving Dense Nonsymmetric Eigenvalue Problems*, Concurrency and Computation: Practice and Experience, 2020, doi: 10.1002/cpe.5915

- Mirko Myllykoski, Carl Christian Kjelgaard Mikkelsen, Angelika Schwarz, Bo Kågström: *D2.7 Eigenvalue solvers for nonsymmetric problems*, public NLAFET deliverable, 2019 (download)

Current status (stable series)

The library currently supports only real arithmetic (real input and output matrices but **real and/or complex** eigenvalues and eigenvectors). In addition, some interface functions are implemented as LAPACK and ScaLAPACK wrapper functions.

Standard eigenvalue problems:

Component	Shared memory	Distributed memory	CUDA
Hessenberg reduction	**Complete**	ScaLAPACK	**Single GPU**
Schur reduction	**Complete**	**Complete**	*Experimental*
Eigenvalue reordering	**Complete**	**Complete**	*Experimental*
Eigenvectors	**Complete**	—	—

Generalized eigenvalue problems:

Component	Shared memory	Distributed memory	CUDA
HT reduction	LAPACK	3rd party	—
Schur reduction	**Complete**	**Complete**	*Experimental*
Eigenvalue reordering	**Complete**	**Complete**	*Experimental*
Eigenvectors	**Complete**	—	—

Please see changelog and known problems.

Documentation:

The StarNEig User's Guide is available in both HTML and PDF formats at `https://nlafet.github.io/`← `StarNEig`. The PDF version is also available under `releases`.

Installation

Prebuild binary packages (Ubuntu 18.04 and Ubuntu 20.04) are available under `releases` and can be installed with the following command:
```
$ sudo apt install ./StarNEig-v0.xx.yy-ubuntu-vv.uu.deb
```

The binary packages rely on mainstream StarPU packages and do not necessary provide full functionality.

For full functionality, it is recommended that StarNEig (and StarPU) are compiled from the source code, see below and/or the StarNEig User's Guide. Please consider using one of the tested `release` versions.

Dependencies

Library dependencies:

- Linux

- CMake 3.3 or newer

- Portable Hardware Locality (hwloc)

- Starpu 1.2 or 1.3

    - Newer versions require the user set the `STARPU_LIBRARIES`, `STARPU_MPI_LIBRARIES` and `STARPU_INCLUDE_PATH` environmental variables.

- OpenBLAS, MKL, GotoBLAS or single-threaded BLAS library

- LAPACK

- MPI (optional)

- CUDA + cuBLAS (optional)

- ScaLAPACK + BLACS (optional)

Test program and example code dependencies:

- pkg-config

- GNU Scientific Library (optional)

- MAGMA (optional)

Configure, build and install

To compile, validate and install StarNEig, execute in the same directory as this `README.md` file:
```
$ mkdir build
$ cd build/
$ cmake ../
$ make
$ make test
$ sudo make install
```

Example

The following example demonstrates how a dense matrix `A` is reduced to real Schur form:
```c
// my_program.c
#include <starneig/starneig.h>
#include <stdlib.h>
#include <time.h>
int main()
{
    int n = 3000;
    srand((unsigned) time(NULL));
    // generate a random matrix A
    int ldA = ((n/8)+1)*8;
    double *A = malloc(n*ldA*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            A[j*ldA+i] = 2.0*rand()/RAND_MAX - 1.0;
    // generate an identity matrix Q
    int ldQ = ((n/8)+1)*8;
    double *Q = malloc(n*ldA*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            Q[j*ldQ+i] = i == j ? 1.0 : 0.0;
    // allocate space for the eigenvalues
    double *real = malloc(n*sizeof(double));
    double *imag = malloc(n*sizeof(double));
    // initialize the StarNEig library
    starneig_node_init(STARNEIG_USE_ALL, STARNEIG_USE_ALL, STARNEIG_HINT_SM);
    // reduce matrix A to real Schur form S = Q^T A Q
    starneig_SEP_SM_Reduce(
        n, A, ldA, Q, ldQ, real, imag, NULL, NULL, NULL, NULL);
    // de-initialize the StarNEig library
    starneig_node_finalize();
    free(A); free(Q); free(real); free(imag);
    return 0;
}
```

Compile:
```
$ gcc -o my_program my_program.c -lstarneig
```

or:
```
$ gcc -o my_program my_program.c $(pkg-config --cflags starneig) $(pkg-config --libs starneig)
```

Chapter 2

Installation

Prebuild binary packages (Ubuntu 18.04 and Ubuntu 20.04) are available under `releases` and can be installed with the following command:

```
$ sudo apt install ./StarNEig-v0.xx.yy-ubuntu-vv.uu.deb
```

The binary packages rely on mainstream StarPU packages and do not necessary provide full functionality.

For full functionality, it is recommended that StarNEig (and StarPU) are compiled from the source code, see below and/or the StarNEig User's Guide. Please consider using one of the tested `release` versions.

Dependencies

Library dependencies:

- Linux

- CMake 3.3 or newer

- Portable Hardware Locality (hwloc)

- Starpu 1.2 or 1.3

    - Newer versions require the user set the `STARPU_LIBRARIES`, `STARPU_MPI_LIBRARIES` and `STARPU_INCLUDE_PATH` environmental variables.

- OpenBLAS, MKL, GotoBLAS or single-threaded BLAS library

- LAPACK

- MPI (optional)

- CUDA (optional)

- ScaLAPACK + BLACS (optional)

Test program and example code dependencies:

- pkg-config

- GNU Scientific Library (optional)

- MAGMA (optional)

Attention

> StarNEig supports OpenBLAS, MKL and GotoBLAS. For optimal performance, a multi-threaded variant of one of the listed BLAS libraries must be provided. StarNEig will automatically set the BLAS library to single-threaded more when necessary. If a different BLAS library is provided, then the user is responsible for setting the BLAS library to *single-threaded* mode. However, the use of a non-supported BLAS library can still impact the performance negatively.

StarPU 1.3.4 installation

1. Download StarPU 1.3.4 (or newer) from http://starpu.gforge.inria.fr/files/

2. Unzip the package and create/enter directory `starpu-1.3.4/build`

3. Configure: `$ ../configure`

4. Compile: `$ make`

5. Install: `$ sudo make install`

The default installation path is `/usr/local` but this can be changed during the configuration phase (`$ ../configure --prefix=...`). It is something necessary to append the `CPATH`, `LIBRARY_PATH`, and `LD_LIBRARY_PATH` environmental variables by adding the following to ∼`/.profile`:
```
export CPATH=$CPATH:/usr/local/include/
export LIBRARY_PATH=$LIBRARY_PATH:/usr/local/lib/
export LD_LIBRARY_PATH=$LD_LIBRARY_PATH:/usr/local/lib/
```

See the StarPU handbook for further instructions: http://starpu.gforge.inria.fr/doc/html/↩ BuildingAndInstallingStarPU.html

Configuration

It is recommended that a user builds the library in a separate build directory:
```
$ cd path_to_the_top_directory/
$ mkdir build
$ cd build
```

The library is configured with the `cmake` command. In most cases, it is not necessary to give this command any additional arguments:
```
$ cmake ../
...
-- Configuring done
-- Generating done
-- Build files have been written to: /.../build
```

However, the library can be customized with various options. For example, the example codes and documentation generation can be enabled by setting the `STARNEIG_ENABLE_EXAMPLES` and `STARNEIG_ENABLE_DOCS` options:
```
$ cmake -DSTARNEIG_ENABLE_EXAMPLES=ON -DSTARNEIG_ENABLE_DOCS=ON ../
```

The *installation path* can be changed during the configuration phase:
```
$ cmake -DCMAKE_INSTALL_PREFIX=/path/to/somewhere/ ../
```

Remarks

The library can be compiled separately from the other software components:
```
$ mkdir build-src
$ cd build-src/
$ cmake ../src/
$ make
```

It may sometimes be necessary to compile CUDA source files with a different compiler than what `cmake` uses by default. For example, some CUDA version do not support GCC compilers that are newer than GCC 5 release series. In that case `cmake` can be configured to use a GCC 5 release series compiler:
```
$ cmake -DCUDA_HOST_COMPILER=/usr/bin/gcc-5 -DCUDA_PROPAGATE_HOST_FLAGS=OFF ../
```

List of StarNEig library specific configuration options:

- `STARNEIG_ENABLE_OPTIMIZATION`: Enables extra compiler optimizations (`ON` by default).

- – Can break portability.

- `STARNEIG_ENABLE_EXAMPLES`: Enables examples (`OFF` by default).

- `STARNEIG_ENABLE_DOCS`: Enables documentation generation (`OFF` by default).

- `STARNEIG_ENABLE_TESTS:`: Enables test program (`ON` by default).

- `STARNEIG_ENABLE_FULL_TESTS`: Enables additional tests (`OFF` by default).

- `STARNEIG_ENABLE_REFERENCE:`: Enables reference MPI implementations (`OFF` by default).

    - – A user must initialize and update the included GIT submodules.

- `STARNEIG_DISABLE_MPI`: Explicitly disables the MPI support even when the system would support it (`OFF` by default).

    - – Must be set before `cmake` is run for the first time.

- `STARNEIG_DISABLE_CUDA`: Explicitly disables the CUDA support even when the system would support it (`OFF` by default).

    - – Must be set before `cmake` is run for the first time.

- `STARNEIG_DISABLE_BLACS`: Explicitly disables the ScaLAPACK/BLACS support even when the system would support it (`OFF` by default).

    - – Must be set before `cmake` is run for the first time.

- `STARNEIG_ENABLE_MESSAGES`: Enable basic verbose messages (`ON` by default).

- `STARNEIG_ENABLE_VERBOSE`: Enable additional verbose messages (`OFF` by default).

- `STARNEIG_ENABLE_EVENTS`: Enable event traces (`OFF` by default).

- `STARNEIG_ENABLE_EVENT_PARSER`: Enable event parser (`OFF` by default).

- `STARNEIG_ENABLE_SANITY_CHECKS`: Enables additional satiny checks. (`OFF` by default).

    - – These checks are very expensive and should not be enabled unless absolutely necessary.

- `STARNEIG_ENABLE_PRUNING`: Enable task graph pruning (`ON` by default).

- `STARNEIG_ENABLE_MRM`: Enable multiple linear regression performance models (`OFF` by default).

- `STARNEIG_ENABLE_CUDA_REORDER_WINDOW`: Enable CUDA-based reorder_window codelet (`OFF` by default).

- `STARNEIG_ENABLE_INTEGER_SCALING`: Enable integer-based scaling factors (`ON` by default).

The following **environmental variables** can be used to configure the used libraries:

- `BLAS_LIBRARIES`: BLAS library.

- `LAPACK_LIBRARIES`: LAPACK library.

- `HWLOC_LIBRARIES`: Portable Hardware Locality (hwloc) library.

- `MPI_LIBRARIES`: C MPI library.

- `MPI_Fortran_LIBRARIES`: Fortran MPI library.

- `SCALAPACK_LIBRARIES`: ScaLAPACK library.

- `BLACS_LIBRARIES`: BLACS library.

- `STARPU_LIBRARIES`: StarPU library.

- `STARPU_MPI_LIBRARIES`: StarPU-MPI library.

- `GSL_LIBRARIES`: GNU Scientific Library.

- `MAGMA_LIBRARIES`: MAGMA library.

- `MISC_LIBRARIES`: Miscellaneous libraries.

For example, if a user has a custom build ATLAS BLAS library and a matching LAPACK library that are not detected by the build system, then the user might define `BLAS_LIBRARIES=/usr/local/atlas/lib/libsatlas.so` and `LAPACK_LIBRARIES=/usr/local/atlas/lib/liblapack.so` before calling `cmake`.

The following environmental variables can be used to configure include paths for the used libraries:

- `OMP_INCLUDE_PATH`: OpenMP include path.

- `BLAS_INCLUDE_PATH`: BLAS include path.

- `MKL_INCLUDE_PATH`: MKL include path.

- `HWLOC_INCLUDE_PATH`: Portable Hardware Locality (hwloc) include path.

- `MPI_INCLUDE_PATH`: MPI include path.

- `STARPU_INCLUDE_PATH`: StarPU include path.

- `GSL_INCLUDE_PATH`: GNU Scientific Library include path.

- `MAGMA_INCLUDE_PATH`: MAGMA include path.

- `MISC_INCLUDE_PATH`: Miscellaneous include paths.

Compile

The library (and other software components) are compiled with the `make` command:
```
$ make
...
[ 99%] Building C object test/CMakeFiles/starneig-test.dir/3rdparty/matrixmarket/mmio.c.o
[100%] Linking C executable ../starneig-test
[100%] Built target starneig-test
```

Test

The automated tests can be executed as follows:
```
$ make test
Running tests...
Test project /.../build
      Start  1: simple-hessenberg
 1/18 Test  #1: simple-hessenberg ..................  Passed   15.19 sec
      Start  2: simple-hessenberg-mpi
 2/18 Test  #2: simple-hessenberg-mpi ..............  Passed   49.59 sec
...
      Start 17: simple-full-chain-generalized
17/18 Test #17: simple-full-chain-generalized .......  Passed  180.50 sec
      Start 18: simple-full-chain-generalized-mpi
18/18 Test #18: simple-full-chain-generalized-mpi ...  Passed  195.39 sec
100% tests passed, 0 tests failed out of 18
Total Test time (real) = 1219.47 sec
```

The `STARNEIG_ENABLE_FULL_TESTS` `cmake` option can be used to enable additional tests.

Install

The library and the related header files are installed by executing:
```
$ sudo make install
```

Chapter 3

Basic usage

Header files

The principal interface functions can be found from the starneig/node.h header file. The library provides separate header files for shared memory (starneig/sep_sm.h, starneig/gep_sm.h) and distributed memory (starneig/sep_dm.h, starneig/gep_dm.h). In most cases, a user can simply include all header files as follows:

```
#include <starneig/starneig.h>
```

Certain header files and interface functions exist only when the library is compiled with MPI and ScaLAPACK / BLACS support. The configuration of the installed library can be found from the starneig/configuration.h header file. See module Library configuration for further information.

Library initialization

Each node must call the starneig_node_init() interface function to initialize the library and the starneig_node_finalize() interface function to shutdown the library:

```
starneig_node_init(cores, gpus, flags);
...
starneig_node_finalize();
```

The starneig_node_init() interface function initializes StarPU (and cuBLAS) and pauses all worker threads. The cores argument specifies the total number of used CPU cores. In distributed memory mode, one of these C←PU cores is automatically allocated for the StarPU-MPI communication thread. The gpus argument specifies the total number of used GPUs. One or more CPU cores are automatically allocated for GPU devices. The flags (starneig_flag_t) argument can provide additional configuration information.

A node can also be configured with default values:

```
starneig_node_init(STARNEIG_USE_ALL, STARNEIG_USE_ALL, STARNEIG_DEFAULT);
```

This tells the library to use all available CPU cores and GPUs.

See module Intra-node execution environment for further information.

Error handling

Most interface functions return one of the following values:

- STARNEIG_SUCCESS (0): The interface function was executed successfully.

- A negative number −i: The i'th interface function argument was invalid.

- A positive number i: The interface function encountered an error or a warning was raised. See module Error codes for further information.

All return values (starneig_error_t) are defined in the starneig/error.h header file.

Remarks

The library may call the exit() and abort() functions if an interface function encounters a fatal error from which it cannot recover.

Performance models

The StarPU performance models must be calibrated before the software can function efficiently on heterogeneous platforms (CPUs + GPUs). The calibration is triggered automatically if the models are not calibrated well enough for a given problem size. This may impact the execution time negatively during the first run. Please see the StarPU handbook for further information: http://starpu.gforge.inria.fr/doc/html/Scheduling.↩html

Compilation and linking

During compilation, the starneig library library must be linked with the user's software:
```
$ gcc -o my_program my_program.c -lstarneig
```

StarNEig provides a pkg-config file for easing the compilation. A user may integrate it to their Makefile:
```
CFLAGS   += $$(pkg-config --cflags starneig)
LDLIBS   += $$(pkg-config --libs starneig)
%.o: %.c
    $(CC) -c -o $@ $< $(CFLAGS)
my_program: my_program.c
    $(CC) -o $@ $^ $(LDLIBS)
```

Or their CMakeLists.txt file:
```
find_package(PkgConfig REQUIRED)
pkg_search_module(STARNEIG REQUIRED starneig)
include_directories (${STARNEIG_INCLUDE_DIRS})
link_directories (${STARNEIG_LIBRARY_DIRS})
set (CMAKE_C_FLAGS "${STARNEIG_C_FLAGS} ${CMAKE_C_FLAGS}")
add_executable (my_program my_program.c)
target_link_libraries (my_program ${STARNEIG_LIBRARIES})
```

Chapter 4

Distributed memory

The STARNEIG_HINT_DM initialization flag tells the library to configure itself for distributed memory computation. The flag is indented to be only a hint and the library will automatically reconfigure itself for the correct computation mode. A user is allowed to mix shared memory and distributed memory functions without reinitializing the library. The library is intended to be run in a **hybrid configuration** (each MPI rank is mapped to several CPU cores). Failing to do so leads to CPU core oversubscription. It is generally a good idea to map each MPI rank to a full node or a NUMA island / CPU socket:

```
# OpenMPI / one rank per node:
$ mpirun -n RANKS --map-by ppr:1:node --bind-to none ...
# OpenMPI / one rank per socket:
$ mpirun -n RANKS --map-by ppr:1:socket --bind-to socket ...
# Intel MPI / one rank per node:
$ mpirun -n RANKS -binding "pin=on;domain=node" ...
# Intel MPI / one rank per socket:
$ mpirun -n RANKS -binding "pin=on;domain=socket" ...
```

Attention

> StarPU attempts to bind the worker threads to the available CPU cores. This may sometimes conflict with the MPI library and/or the batch system CPU core allocation. StarNEig library attempts to correct for by factoring in the CPU core binding mask. However, if neither the MPI library nor the batch system enforces such a binding mask, it is possible that several StarPU worker threads end up bound to a same CPU core. In such a situation, it is recommended that a user disables the StarPU thread binding explicitly:
> ```
> STARPU_WORKERS_NOBIND=1 mpirun ...
> ```

This is particularly important when several ranks / processes are mapped to a same node.

The library assumes that the MPI library is already initialized when the starneig_node_init() interface function is called with the STARNEIG_HINT_DM flag or when the library reconfigures itself for distributed memory after a user has called a distributed memory interface function. The MPI library should be initialized either in the serialized mode:

```
int thread_support;
MPI_Init_thread(
    &argc, (char ***)&argv, MPI_THREAD_SERIALIZED, &thread_support);
if (thread_support < MPI_THREAD_SERIALIZED) {
    fprintf(stderr,
        "MPI_THREAD_SERIALIZED is not supported. Aborting...\n");
    abort();
}
```

Or in the multi-threaded mode:

```
int thread_support;
MPI_Init_thread(
    &argc, (char ***)&argv, MPI_THREAD_MULTIPLE, &thread_support);
if (thread_support < MPI_THREAD_SERIALIZED) {
    fprintf(stderr,
        "MPI_THREAD_SERIALIZED is not supported. Aborting...\n");
    abort();
} else if (thread_support < MPI_THREAD_MULTIPLE) {
    fprintf(stderr,
        "Warning: MPI_THREAD_MULTIPLE is not supported.\n");
}
```

A user is allowed to change the library MPI communicator with the starneig_mpi_set_comm() interface function. This interface function should be called **before** the library is initialized.

Data distribution

Distributed matrices are represented using two opaque objects:

- *Data distribution* (starneig_distr_t)

- *Distributed matrix* (starneig_distr_matrix_t)

Each matrix is divided into rectangular blocks of uniform size (excluding the last block row and column):

(0,0)	(0,1)	(0,2)	(0,3)	(0,4)	(0,5)	(0,6)
(1,0)	(1,1)	(1,2)	(1,3)	(1,4)	(1,5)	(1,6)
(2,0)	(2,1)	(2,2)	(2,3)	(2,4)	(2,5)	(2,6)
(3,0)	(3,1)	(3,2)	(3,3)	(3,4)	(3,5)	(3,6)
(4,0)	(4,1)	(4,2)	(4,3)	(4,4)	(4,5)	(4,6)

Figure 4.1 A matrix divided into rectangular blocks of uniform size.

The blocks are indexed using a two-dimensional index space. A data distribution encapsulates an arbitrary mapping from this two-dimensional block index space to the one-dimensional MPI rank space:

Figure 4.2 An example of a block mapping.

In the above example, the rank 0 owns the blocks (0,1), (1,2), (1,5), (1,6), (2,6), (3,0) and (3,5). Naturally, a data distribution can describe a two-dimensional block cyclic distribution that is very common with ScaLAPACK subroutines:

Figure 4.3 An example of a row-major ordered two-dimensional block cyclic mapping.

A data distribution can be created using one of the following interface functions:

- starneig_distr_init() creates a default distribution (row-major ordered two-dimensional block cyclic distribution with a squarish mesh).

- starneig_distr_init_mesh() creates a row-major or column-major ordered two-dimensional block cyclic distribution with desired number of rows and columns in the mesh.

- starneig_distr_init_func() creates an arbitrary distribution defined by a function.

Fox example,
```
starneig_distr_t distr = starneig_distr_init_mesh(4, 6, STARNEIG_ORDER_DEFAULT);
```

would create a two-dimensional block cyclic distribution with 4 rows and 6 columns in the mesh. Alternatively, a user can create an equivalent data distribution using the starneig_distr_init_func() interface function:
```
// additional distribution function argument structure
struct block_cyclic_arg {
    int rows;   // the number of rows in the process mesh
    int cols;   // the number of columns in the process mesh
};
// distribution function (row-major ordered 2D block cyclic distribution)
int block_cyclic_func(int i, int j, void *arg)
{
    struct block_cyclic_arg *mesh = (struct block_cyclic_arg *) arg;
    return (i % mesh->rows) * mesh->cols + j % mesh->cols;
}
void func(...)
{
    ...
    struct block_cyclic_arg arg = { .rows = 4, .cols = 6 };
    starneig_distr_t distr =
        starneig_distr_init_func(&block_cyclic_func, &arg, sizeof(arg));
    ...
}
```

A data distribution is destroyed with the starneig_distr_destroy() interface function.

Remarks

Certain interface functions (starneig_SEP_DM_Hessenberg(), starneig_SEP_DM_Reduce(), starneig_GEP_DM_HessenbergTr and starneig_GEP_DM_Reduce()) are wrappers for / use several ScaLAPACK subroutines. The involved matrices should thus have a two-dimensional block cyclic data distribution. The library will automatically convert the matrices to a compatible format but this requires extra memory.

Distributed matrix

A distributed matrix is created using the starneig_distr_matrix_create() interface function. The function call will automatically allocate the required local resources. For example,

```
starneig_distr_t distr =
    starneig_distr_init_mesh(4, 6, STARNEIG_ORDER_DEFAULT);
starneig_distr_matrix_t dA =
    starneig_distr_matrix_create(m, n, bm, bn, STARNEIG_REAL_DOUBLE, distr);
```

would create a $m \times n$ double-precision real matrix that is distributed in a two-dimensional block cyclic fashion in $bm \times bn$ blocks. Or,

```
starneig_distr_matrix_t dB =
    starneig_distr_matrix_create(n, n, -1, -1, STARNEIG_REAL_DOUBLE, NULL);
```

would create a $n \times n$ double-precision real matrix with a default data distribution (`NULL` argument) and a default block size (`-1, -1`).

Attention

> StarNEig library is designed to use much larger distributed blocks than ScaLAPACK. Selecting a too small distributed block size will be detrimental to the performance.

A user may access the locally owned blocks using the starneig_distr_matrix_get_blocks() interface function. A distributed matrix is destroyed using the starneig_distr_matrix_destroy() interface function. This will deallocate all local resources. See module Distributed Memory / Distributed matrices for further information.

Remarks

> Certain interface functions (starneig_SEP_DM_Hessenberg(), starneig_SEP_DM_Reduce(), starneig_GEP_DM_HessenbergTr and starneig_GEP_DM_Reduce()) are wrappers for / use several ScaLAPACK subroutines. The involved matrices should thus be distributed in square blocks. In addition, the ScaLAPACK subroutines usually perform better when the block size is relatively small. The library will automatically convert the matrices to a compatible format but this requires extra memory.

Copy, scatter and gather

An entire distributed matrix can be copied with the starneig_distr_matrix_copy() interface function:

```
starneig_distr_matrix_t dA, dB;
...
starneig_distr_matrix_copy(dB, dA);
```

This copies distributed matrix `dB` to a distributed matrix `dA`. A region (submatrix) of a distributed matrix can be copied to a second distributed matrix using the starneig_distr_matrix_copy_region() interface function.

A local matrix can be converted to a "single owner" distributed matrix with the starneig_distr_matrix_create_local() interface function:

```
int owner = 0; // MPI rank that owns the local matrix
double *A;      // local pointer
int ldA;        // matching leading dimension
...
starneig_distr_matrix_t lA = starneig_distr_matrix_create_local(
    n, n, STARNEIG_REAL_DOUBLE, owner, A, ldA);
```

This creates a wrapper object, i.e., the pointer `A` and the distributed matrix `lA` point to the same data on the `owner` node. The created distributed matrix is associated with a data distribution that indicated that the whole matrix is owned by the node `owner`. The used block size is $n \times n$.

Copying from a "single owner" distributed matrix to a distributed matrix performs a *scatter* operation and copying from a distributed matrix to a "single owner" distributed matrix performs a *gather* operation.

ScaLAPACK compatibility layer

The library provides a ScaLAPACK compatibility layer:

- BLACS contexts are encapsulated inside starneig_blacs_context_t objects.

- BLACS descriptors are encapsulated inside starneig_blacs_descr_t objects.

A two-dimensional block cyclic data distribution can be converted to a BLACS context and vice versa using the starneig_distr_to_blacs_context() and starneig_blacs_context_to_distr() interface functions, respectively. Similarly, a distributed matrix that uses a two-dimensional block cyclic data distribution can be converted to a BLACS descriptor (and a local buffer) and vice versa using the starneig_distr_matrix_to_blacs_descr() and starneig_blacs_descr_to_distr_matrix() interface functions, respectively. The conversion is performed in-place and a user is allowed to mix StarNEig interface functions with ScaLAPACK style subroutines/functions without reconversion.

For example,
```
starneig_distr_matrix_t dA = starneig_distr_matrix_create(...);
...
// convert the data distribution to a BLACS context
starneig_distr_t distr = starneig_distr_matrix_get_distr(A);
starneig_blacs_context_t context = starneig_distr_to_blacs_context(distr);
// convert the distributed matrix to a BLACS descriptor and a local buffer
starneig_blacs_descr_t descr_a;
double *local_a;
starneig_distr_matrix_to_blacs_descr(dA, context, &descr_a, (void **)&local_a);
...
// a ScaLAPACK subroutine for reducing general distributed matrix to upper
// Hessenberg form
extern void pdgehrd_(int const *, int const *, int const *, double *,
    int const *, int const *, starneig_blacs_descr_t const *, double *,
    double *, int const *, int *);
pdgehrd_(&n, &ilo, &ihi, local_a, &ia, &ja, &descr_a, tau, ...);
```

converts a distributed matrix dA to a BLACS descriptor descr_a and a local pointer local_a. The descriptor and the local array are then fed to a ScaLAPACK subroutine. A user must make sure that the live time of the distributed matrix dA is at least as long as the live time of the matching BLACS descriptor descr_a. See modules ScaLAPACK compatibility / BLACS helpers and ScaLAPACK compatibility / BLACS matrices for further information.

Chapter 5

Standard eigenvalue problem

Given a matrix $A \in \mathbb{R}^{n \times n}$, the *standard eigenvalue problem* (SEP) consists of finding *eigenvalues* $\lambda_i \in \mathbb{C}$ and associated *eigenvectors* $0 \neq v_i \in \mathbb{C}^n$ such that

$$Av_i = \lambda_i v_i, \text{ for } i = 1, 2, \ldots, n.$$

The eigenvalues are the n (potentially complex) roots of the polynomial $\det(A - \lambda I) = 0$ of degree n. There is often a full set of n linearly independent eigenvectors, but if there are *multiple* eigenvalues (i.e., if $\lambda_i = \lambda_j$ for some $i \neq j$) then there might not be a full set of independent eigenvectors.

The library provides 12 interface functions for the standard case.

Hessenberg reduction

Given a general matrix A, the starneig_SEP_SM_Hessenberg() and starneig_SEP_DM_Hessenberg() interface functions compute a Hessenberg decomposition

$$A = U * H * U^T,$$

where H is upper Hessenberg and U is orthogonal. On exit, A is overwritten by H and Q (which is an orthogonal matrix on entry) is overwritten by

$$Q \leftarrow Q * U.$$

This is done in order to greatly accelerate the subsequent computation of a Schur decomposition since when working on H of size $n \times n$, the amount of work in each iteration of the QR algorithm is reduced from $\mathcal{O}(n^3)$ to $\mathcal{O}(n^2)$ flops.

Schur reduction

Given a Hessenberg decomposition

$$A = Q * H * Q^T,$$

of a general matrix A, the starneig_SEP_SM_Schur() and starneig_SEP_DM_Schur() interface functions compute a Schur decomposition

$$A = Q * (U * S * U^T) * Q^T$$

where S is upper quasi-triangular with 1×1 and 2×2 blocks on the diagonal (Schur matrix) and U is orthogonal. The 1×1 blocks correspond to the real eigenvalues and each 2×2 block corresponds to a pair of complex conjugate eigenvalues. On exit, H is overwritten by S and Q is overwritten by

$$Q \leftarrow Q * U.$$

Eigenvalue reordering

Given a Schur decomposition

$$A = Q * S * Q^T$$

of a general matrix A and a selection of eigenvalues, the starneig_SEP_SM_ReorderSchur() and starneig_SEP_DM_ReorderSchur() interface functions attempt to compute an updated Schur decomposition

$$A = Q * \left(U * \begin{bmatrix} \hat{S}_{11} & \hat{S}_{12} \\ 0 & \hat{S}_{22} \end{bmatrix} * U^T \right) * Q^T,$$

where the selected eigenvalues appear in $\hat{S}_{11}$. On exit, S is overwritten by $\hat{S}$ and Q is overwritten by

$$Q \leftarrow Q * U.$$

Reordering may in rare cases fail. In such cases the output is guaranteed to be a Schur decomposition and all (if any) selected eigenvalues that are correctly placed are marked in the selection array on exit. Reordering may perturb the eigenvalues and the eigenvalues after reordering are returned.

Combined reduction to Schur form and eigenvalue reordering

Given a general matrix A, the starneig_SEP_SM_Reduce() and starneig_SEP_DM_Reduce() interface functions compute a (reordered) Schur decomposition

$$A = U * S * U^T.$$

Optionally, the interface functions attempt to reorder selected eigenvalues to the top left corner of the Schur matrix S. On exit, A is overwritten by S and Q (which is an orthogonal matrix on entry) is overwritten by

$$Q \leftarrow Q * U.$$

Reordering may in rare cases fail. In such cases the output is guaranteed to be a Schur decomposition and all (if any) selected eigenvalues that are correctly placed are marked in the selection array on exit. Reordering may perturb the eigenvalues and the eigenvalues after reordering are returned.

Eigenvectors

Given a subset consisting of $m \leq n$ of the eigenvalues λ_i for $i = 1, 2, \ldots, m$ and a Schur decomposition $A = QSQ^H$, we can compute for each λ_i an *eigenvector* $v_i \neq 0$ such that $Av_i = \lambda_i v_i$ by first computing an eigenvector w_i of S and then transform it back to the original basis by pre-multiplication with Q.

Given a Schur decomposition

$$A = Q * S * Q^T$$

of a general matrix A and a selection of eigenvalues, the starneig_SEP_SM_Eigenvectors() and starneig_SEP_DM_Eigenvectors() interface functions compute and return an eigenvector for each of the selected eigenvalues.

The eigenvectors are stored as columns in the output matrix X in the same order as their corresponding eigenvalues appear in the selection array. A real eigenvector is stored as a single column. The real and imaginary parts of a complex eigenvector are stored as consecutive columns.

For a selected pair of complex conjugate eigenvalues, an eigenvector is computed only for the eigenvalue with positive imaginary part. Thus, every selected eigenvalue contributes one column to the output matrix and thus the number of selected eigenvalues is equal to the number of columns of X.

Eigenvalue selection helper

Given a Schur matrix and a predicate function, the starneig_SEP_SM_Select() and starneig_SEP_DM_Select() interface functions conveniently generate a correct selection array and count the number of selected eigenvalues. The count is useful when allocating storage for the eigenvector matrix computed by the starneig_SEP_SM_Eigenvectors() and starneig_SEP_DM_Eigenvectors() interface functions.

```c
// a predicate function that selects all eigenvalues that have a real
// part that is larger than a given value
static int predicate(double real, double imag, void *arg)
{
    double value = * (double *) arg;
    if (value < real)
        return 1;
    return 0;
}
void func(...)
{
    double *S; int ldS;
    ...
    double value = 0.5;
    int num_selected, *selected = malloc(n*sizeof(int));
    starneig_SEP_SM_Select(
        n, S, ldS, &predicate, &value, selected, &num_selected);
    ...
}
```

See modules Shared Memory / Standard EVP and Distributed Memory / Standard EVP for further information. See also examples sep_sm_full_chain.c, sep_dm_full_chain.c and sep_sm_eigenvectors.c.

Chapter 6

Generalized eigenvalue problem

Given a matrix pair (A, B), where $A, B \in \mathbb{R}^{n \times n}$, the *generalized eigenvalue problem* (GEP) consists of finding *generalized eigenvalues* $\lambda_i \in \mathbb{C}$ and associated *generalized eigenvectors* $0 \neq v_i \in \mathbb{C}^n$ such that

$$Av_i = \lambda_i Bv_i, \text{ for } i = 1, 2, \ldots, n.$$

The eigenvalues are the n (potentially complex) roots of the polynomial $\det(A - \lambda B) = 0$ of degree n. There is often a full set of n linearly independent generalized eigenvectors, but if there are *multiple eigenvalues* (i.e., if $\lambda_i = \lambda_j$ for some $i \neq j$) then there might not be a full set of independent eigenvectors.

At least in principle, a GEP can be transformed into a SEP provided that B is invertible, since

$$Av = \lambda Bv \Leftrightarrow (B^{-1}A)v = \lambda v.$$

However, in finite precision arithmetic this practice is not recommended.

The library provides 12 interface functions for the generalized case.

Hessenberg-triangular reduction

Given a general matrix pair (A, B), the starneig_GEP_SM_HessenbergTriangular() and starneig_GEP_DM_HessenbergTriangular() interface functions compute a Hessenberg-triangular decomposition

$$(A, B) = U_1 * (H, R) * U_2^T,$$

where H is upper Hessenberg, R is upper triangular, and U_1 and U_2 are orthogonal. On exit, A is overwritten by H, B is overwritten by R, and Q and Z (which are orthogonal matrices on entry) are overwritten by

$$Q \leftarrow Q * U_1 \text{ and } Z \leftarrow Z * U_2.$$

This is done in order to greatly accelerate the subsequent computation of a generalized Schur decomposition.

Generalized Schur reduction

Given a Hessenberg-triangular decomposition

$$(A, B) = Q * (H, R) * Z^T$$

of a general matrix pair (A, B), the starneig_GEP_SM_Schur() and starneig_GEP_DM_Schur() interface functions function compute a generalized Schur decomposition

$$(A, B) = Q * (U_1 * (S, T) * U_2^T) * Z^T,$$

where S is upper quasi-triangular with 1×1 and 2×2 blocks on the diagonal, T is a upper triangular matrix, and U_1 and U_2 are orthogonal. The 1×1 blocks on the diagonal of (S, T) correspond to the real generalized eigenvalues and each 2×2 block corresponds to a pair of complex conjugate generalized eigenvalues.

On exit, H is overwritten by S, R is overwritten by T, and Q and Z are overwritten by

$$Q \leftarrow Q * U_1 \text{ and } Z \leftarrow Z * U_2.$$

The computed generalized eigenvalues are returned as a pair of values (α, β) such that α/β gives the actual generalized eigenvalue. The quantity α/β may overflow.

Generalized eigenvalue reordering

Given a generalized Schur decomposition

$$(A, B) = Q * (S, T) * Z^T$$

of a general matrix pair (A, B) and a selection of generalized eigenvalues, the starneig_GEP_SM_ReorderSchur() and starneig_GEP_DM_ReorderSchur() interface functions attempt to compute an updated generalized Schur decomposition

$$(A, B) = Q * \left(U_1 * \left(\begin{bmatrix} \hat{S}_{11} & \hat{S}_{12} \\ 0 & \hat{S}_{22} \end{bmatrix}, \begin{bmatrix} \hat{T}_{11} & \hat{T}_{12} \\ 0 & \hat{T}_{22} \end{bmatrix} \right) * U_2^T \right) * Z^T.$$

where the selected generalized eigenvalues appear in $(\hat{S}_{11}, \hat{T}_{11})$.

On exit, S is overwritten by $\hat{S}$, T is overwritten by $\hat{T}$, and Q and Z are overwritten by

$$Q \leftarrow Q * U_1 \text{ and } Z \leftarrow Z * U_2.$$

Reordering may in rare cases fail. In such cases the output is guaranteed to be a Schur decomposition and all (if any) selected generalized eigenvalues that are correctly placed are marked in the selection array on exit.

Reordering may perturb the generalized eigenvalues and the generalized eigenvalues after reordering are returned. The computed generalized eigenvalues are returned as a pair of values (α, β) such that α/β gives the actual generalized eigenvalue. The quantity α/β may overflow.

Combined reduction to generalized Schur form and eigenvalue reordering

Given a general matrix pair (A, B), the starneig_GEP_SM_Reduce() and starneig_GEP_DM_Reduce() interface functions compute a (reordered) generalized Schur decomposition

$$(A, B) = U_1 * (S, T) * U_2^T.$$

Optionally, the interface functions attempt to reorder selected generalized eigenvalues to the top left corner of the generalized Schur decomposition.

On exit, A is overwritten by S, B is overwritten by T, and Q and Z (which are orthogonal matrices on entry) are overwritten by

$$Q \leftarrow Q * U_1 \text{ and } Z \leftarrow Z * U_2.$$

The computed generalized eigenvalues are returned as a pair of values (α, β) such that α/β gives the actual generalized eigenvalue. The quantity α/β may overflow.

Reordering may in rare cases fail. In such cases the output is guaranteed to be a Schur-triangular decomposition and all (if any) selected generalized eigenvalues that are correctly placed are marked in the selection array on exit.

Generalized eigenvectors

Given a subset consisting of $m \leq n$ of the eigenvalues λ_i for $i = 1, 2, \ldots, m$ and a generalized Schur decomposition $(A, B) = Q(S, T)Z^H$, we can compute for each λ_i a *generalized eigenvector* $v_i \neq 0$ such that $Av_i = \lambda_i Bv_i$ by first computing a generalized eigenvector w_i of $S - \lambda_i T$ and then transform it back to the original basis by pre-multiplication with Z.

Given a generalized Schur decomposition

$$(A, B) = Q * (S, T) * Z^T$$

of a general matrix pair (A, B) and a selection of generalized eigenvalues, the starneig_GEP_SM_Eigenvectors() and starneig_GEP_DM_Eigenvectors() interface functions compute and return a generalized eigenvector for each of the selected generalized eigenvalues.

The generalized eigenvectors are stored as columns in the output matrix X in the same order as their corresponding generalized eigenvalues appear in the selection array. A real generalized eigenvector is stored as a single column. The real and imaginary parts of a complex generalized eigenvector are stored as consecutive columns.

For a selected pair of complex conjugate generalized eigenvalues, a generalized eigenvector is computed only for the generalized eigenvalue with positive imaginary part. Thus, every selected generalized eigenvalue contributes one column to the output matrix and thus the number of selected generalized eigenvalues is equal to the number of columns of X.

Eigenvalue selection helper

Given a Schur-triangular matrix pair (S, T) and a predicate function, the starneig_GEP_SM_Select() and starneig_GEP_DM_Select() interface functions conveniently generate a correct selection array and count the number of selected generalized eigenvalues. The count is useful when allocating storage for the generalized eigenvector matrix computed by starneig_GEP_DM_Eigenvectors().

```c
// a predicate function that selects all finite generalized eigenvalues that
// have a real part that is larger than a given value
static int predicate(double real, double imag, double beta, void *arg)
{
    double value = * (double *) arg;
    if (beta != 0.0 && value < real/beta)
        return 1;
    return 0;
}
void func(...)
{
    ...
    double value = 0.5;
    int num_selected, *selected = malloc(n*sizeof(int));
    starneig_GEP_SM_Select(
        n, S, ldS, T, ldT, &predicate, &value, selected, &num_selected);
    ...
}
```

See modules Shared Memory / Generalized EVP and Distributed Memory / Generalized EVP for further information. See also examples gep_sm_full_chain.c, gep_dm_full_chain.c and gep_sm_eigenvectors.c.

Chapter 7

Expert functions

The library provides a set of configuration structures:

- starneig_hessenberg_conf : A configuration structure for Hessenberg reduction related expert interface functions.

- starneig_schur_conf : A configuration structure for Schur reduction related expert interface functions.

- starneig_reorder_conf : A configuration structure for eigenvalue reordering related interface functions.

- starneig_eigenvectors_conf : A configuration structure for eigenvector computation related interface functions.

The default parameters can generated with the following interface functions:

- starneig_hessenberg_init_conf() : Generates default parameters for Hessenberg reduction related expert interface functions.

- starneig_schur_init_conf() : Generates default parameters for Schur reduction related expert interface functions.

- starneig_reorder_init_conf() : Generates default parameters for eigenvalue reordering related interface functions.

- starneig_eigenvectors_init_conf() : Generates default parameters for eigenvector computation related interface functions.

A user is allowed to modify these default values before passing them to the expert interface function.

Only certain interface functions have expert version:

- starneig_SEP_SM_Hessenberg_expert()

- starneig_SEP_SM_Schur_expert()

- starneig_SEP_SM_ReorderSchur_expert()

- starneig_SEP_SM_Eigenvectors_expert()

- starneig_SEP_DM_Schur_expert()

- starneig_SEP_DM_ReorderSchur_expert()

- starneig_GEP_SM_Schur_expert()

- starneig_GEP_SM_ReorderSchur_expert()

- starneig_GEP_SM_Eigenvectors_expert()

- starneig_GEP_DM_Schur_expert()

- starneig_GEP_DM_ReorderSchur_expert()

See module Expert configuration structures for further information.

Chapter 8

Test program

The test program provides an unified interface for testing the entire library. Most command line arguments have default values so only few of them have to be set in most situations. General usage information can be printed as follows:

```
$ ./starneig-test
Usage: ./starneig-test (options)
Global options:
  --mpi -- Enable MPI
  --mpi-mode [serialized,multiple] -- MPI mode
  --seed (num) -- Random number generator seed
  --experiment (experiment) -- Experiment module
  --test-workers [(num),default] -- Test program StarPU worker count
  --blas-threads [(num),default] -- Test program BLAS thread count
  --lapack-threads [(num),default] -- LAPACK solver thread count
  --scalapack-threads [(num),default] -- ScaLAPACK solver thread count
Available experiment modules:
    'hessenberg' : Hessenberg reduction experiment
    'schur' : Schur reduction experiment
    'reorder' : Eigenvalue reordering experiment
    'eigenvectors' : Eigenvectors experiment
    'full-chain' : Full chain experiment
    'partial-hessenberg' : Partial Hessenberg reduction experiment
    'validator' : Validation experiment
```

The `--mpi` option enables the MPI support and `--seed (num)` option initializes the random number generator with a given seed. Available experiment modules are listed below the global options and the desired experiment module is selected with the `--experiment (experiment)` option. For example, the Hessenberg reduction specific experiment module usage information can be printed as follows:

```
$ ./starneig-test --experiment hessenberg
Usage: ./starneig-test (options)
Global options:
  --mpi -- Enable MPI
  --mpi-mode [serialized,multiple] -- MPI mode
  --seed (num) -- Random number generator seed
  --experiment (experiment) -- Experiment module
  --test-workers [(num),default] -- Test program StarPU worker count
  --blas-threads [(num),default] -- Test program BLAS thread count
  --lapack-threads [(num),default] -- LAPACK solver thread count
  --scalapack-threads [(num),default] -- ScaLAPACK solver thread count
Available experiment modules:
    'hessenberg' : Hessenberg reduction experiment
    'schur' : Schur reduction experiment
    'reorder' : Eigenvalue reordering experiment
    'eigenvectors' : Eigenvectors experiment
    'full-chain' : Full chain experiment
    'partial-hessenberg' : Partial Hessenberg reduction experiment
    'validator' : Validation experiment
Experiment module (hessenberg) specific options:
  --data-format (format) -- Data format
  --init (initializer) -- Initialization module
  --solver (solver) -- Solver module
  --hooks (hook:mode, ...) -- Hooks and modes
  --no-reinit -- Do not reinitialize after each repetition
  --repeat (num) -- Repeated experiment
  --warmup (num) -- Perform "warmups"
  --keep-going -- Try to recover from a solver failure
  --abort -- Call abort() in failure
...
```

The overall design of the test program is modular. Each experiment module is built on *initializers*, *solvers* and *hooks*. Each experiment module contains several of each, thus allowing a user to initialize the inout data in various ways and compare different solvers with each other. Each of these building blocks can be configured with various command line arguments. However, in most situations only the problem dimension `--n (num)` needs to be specified.

Hooks are used to test and validate the output of the solver. For example, `--hooks hessenberg residual print` option enables hooks that

- check whether the output matrix is in upper Hessenberg form;

- computes the related residuals using Frobenius norm (reported in terms of unit roundoff error) and checks that they are within the permissible limits; and

- prints the output matrices.

Certain general purpose initializers allow a user to read the input data from a disk (`read-mtx` and `read-raw`) and output data can be stored to a disk using a suitable post-processing hook (`store-raw`).

The test program supports various data formats. For example, shared memory experiments are usually performed using the `pencil-local` data format which stores the matrices continuously in the main memory. Distributed memory experiments are performed using either a StarNEig library specific distributed data format (`pencil-starneig`) and the BLACS data format (`pencil-starneig-blacs`). The test program will detect the correct data format automatically. The test program is also is capable of converting the data between various data formats. The related data converter modules can be in most cases configured using additional command line arguments. For example, the distributed data formats distribute the data in rectangular sections and the section size can be set with command line arguments `--section-height (num)` and `--section-width (num)`).

Performance models

The StarPU performance models must be calibrated before the software can function efficiently on heterogeneous platforms (CPUs+GPUs). The calibration is triggered automatically if the models are not calibrated well enough for a given problem size. This can impact the execution time negatively. The test program provides an easy-to-use solution for this problem:

```
$ ./starneig-test ... --warmup 3
```

The `--warmup (number)` argument causes the test program to perform a number of warm-up runs before the actual execution time measurement takes place.

Please see the StarPU handbook for further instructions: [http://starpu.gforge.inria.↩fr/doc/html/Scheduling.html](http://starpu.gforge.inria.fr/doc/html/Scheduling.html)

Examples

- Reorder a 4000 x 4000 matrix using the StarNEig implementation:

```
$ ./starneig-test --experiment reorder --n 4000
TEST: --seed 1585762840 --experiment reorder --test-workers default --blas-threads default --lapack-threads
      default --scalapack-threads default --data-format pencil-local --init default --n 4000
      --complex-distr uniform --complex-ratio 0.500000 --zero-ratio 0.010000 --inf-ratio 0.010000
      --data-distr default --section-height default --section-width default --select-ratio 0.350000
      --solver starneig --cores default --gpus default --tile-size default --window-size default
      --values-per-chain default --small-window-size default --small-window-threshold default
      --update-width default --update-height default --plan default --blueprint default --hooks
      schur:normal eigenvalues:normal analysis:normal reordering:normal residual:normal
      --eigenvalues-fail-threshold 10000 --eigenvalues-warn-threshold 1000 --reordering-fail-threshold
      10000 --reordering-warn-threshold 1000 --residual-fail-threshold 10000 --residual-warn-threshold 500
      --repeat 1 --warmup 0
THREADS: Using 6 StarPU worker threads during initialization and validation.
```

```
THREADS: Using 6 BLAS threads during initialization and validation.
THREADS: Using 6 BLAS threads in LAPACK solvers.
THREADS: Using 1 BLAS threads in ScaLAPACK solvers.
INIT...
PREPARE...
PROCESS...
[starneig][message] Setting tile size to 192.
[starneig][message] Using multi-part task insertion plan.
[starneig][message] Using two-pass backward dummy blueprint.
[starneig][message] Using "rounded" window size.
EXPERIMENT TIME = 1701 MS
FINALIZE...
EIGENVALUES CHECK: mean = 0 u, min = 0 u, max = 0 u
EIGENVALUES ANALYSIS: zeros = 36, infinities = 0, indefinites = 0
EIGENVALUES ANALYSIS: close zeros = 0, close infinities = 0, close indefinites = 0
REORDERING CHECK: Checking selected eigenvalues...
REORDERING CHECK: Checking other eigenvalues...
REORDERING CHECK: mean = 71 u, min = 0 u, max = 526 u
|Q ~A Q^T - A| / |A| = 314 u
|Q Q^T - I| / |I| = 140 u
================================================================
TIME = 1701 MS [avg 1701 MS, cv 0.00, min 1701 MS, max 1701 MS]
NO FAILED SCHUR FORM TESTS
EIGENVALUES CHECK (WARNINGS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES CHECK (FAILS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES CHECK (MEANS): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES CHECK (MIN): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES CHECK (MAX): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES ANALYSIS (ZEROS): [avg 36.0, cv 0.00, min 36, max 36]
EIGENVALUES ANALYSIS (CLOSE ZEROS): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (INFINITIES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (CLOSE INFINITIES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (INDEFINITES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (CLOSE INDEFINITES): [avg 0.0, cv 0.00, min 0, max 0]
REORDERING CHECK (WARNINGS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
REORDERING CHECK (FAILS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
REORDERING CHECK (MEANS): [avg 71 u, cv 0.00, min 71 u, max 71 u]
REORDERING CHECK (MIN): [avg 0 u, cv 0.00, min 0 u, max 0 u]
REORDERING CHECK (MAX): [avg 526 u, cv 0.00, min 526 u, max 526 u]
|Q ~A Q^T - A| / |A| = [avg 314 u, cv 0.00, min 314 u, max 314 u]
|Q Q^T - I| / |I| = [avg 140 u, cv 0.00, min 140 u, max 140 u]
```

- Reorder a 4000 x 4000 matrix stencil (A, B) using the StarNEig implementation, initialize the random number generator using the seed `1480591971`, fortify the matrix stencil (A, B) against failed swaps, and disable GPUs:

```
$ ./starneig-test --experiment reorder --n 4000 --generalized --seed 1480591971 --fortify --gpus 0
TEST: --seed 1480591971 --experiment reorder --test-workers default --blas-threads default --lapack-threads
    default --scalapack-threads default --data-format pencil-local --init default --n 4000 --generalized
    --complex-distr uniform --fortify --data-distr default --section-height default --section-width
    default --select-ratio 0.350000 --solver starneig --cores default --gpus 0 --tile-size default
    --window-size default --values-per-chain default --small-window-size default --small-window-threshold
    default --update-width default --update-height default --plan default --blueprint default --hooks
    schur:normal eigenvalues:normal analysis:normal reordering:normal residual:normal
    --eigenvalues-fail-threshold 10000 --eigenvalues-warn-threshold 1000 --reordering-fail-threshold
    10000 --reordering-warn-threshold 1000 --residual-fail-threshold 10000 --residual-warn-threshold 500
    --repeat 1 --warmup 0
THREADS: Using 6 StarPU worker threads during initialization and validation.
THREADS: Using 6 BLAS threads during initialization and validation.
THREADS: Using 6 BLAS threads in LAPACK solvers.
THREADS: Using 1 BLAS threads in ScaLAPACK solvers.
INIT...
PREPARE...
PROCESS...
[starneig][message] Setting tile size to 192.
[starneig][message] Using multi-part task insertion plan.
[starneig][message] Using two-pass backward dummy blueprint.
[starneig][message] Using "rounded" window size.
EXPERIMENT TIME = 7472 MS
FINALIZE...
EIGENVALUES CHECK: mean = 0 u, min = 0 u, max = 0 u
EIGENVALUES ANALYSIS: zeros = 0, infinities = 0, indefinites = 0
EIGENVALUES ANALYSIS: close zeros = 0, close infinities = 0, close indefinites = 0
REORDERING CHECK: Checking selected eigenvalues...
REORDERING CHECK: Checking other eigenvalues...
REORDERING CHECK: mean = 23 u, min = 0 u, max = 169 u
|Q ~A Z^T - A| / |A| = 46 u
|Q ~B Z^T - B| / |B| = 63 u
|Q Q^T - I| / |I| = 44 u
|Z Z^T - I| / |I| = 43 u
================================================================
TIME = 7472 MS [avg 7472 MS, cv 0.00, min 7472 MS, max 7472 MS]
NO FAILED SCHUR FORM TESTS
EIGENVALUES CHECK (WARNINGS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
```

```
EIGENVALUES CHECK (FAILS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES CHECK (MEANS): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES CHECK (MIN): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES CHECK (MAX): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES ANALYSIS (ZEROS): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (CLOSE ZEROS): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (INFINITIES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (CLOSE INFINITIES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (INDEFINITES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (CLOSE INDEFINITES): [avg 0.0, cv 0.00, min 0, max 0]
REORDERING CHECK (WARNINGS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
REORDERING CHECK (FAILS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
REORDERING CHECK (MEANS): [avg 23 u, cv 0.00, min 23 u, max 23 u]
REORDERING CHECK (MIN): [avg 0 u, cv 0.00, min 0 u, max 0 u]
REORDERING CHECK (MAX): [avg 169 u, cv 0.00, min 169 u, max 169 u]
|Q ~A Z^T - A| / |A| = [avg 46 u, cv 0.00, min 46 u, max 46 u]
|Q ~B Z^T - B| / |B| = [avg 63 u, cv 0.00, min 63 u, max 63 u]
|Q Q^T - I| / |I| = [avg 44 u, cv 0.00, min 44 u, max 44 u]
|Z Z^T - I| / |I| = [avg 43 u, cv 0.00, min 43 u, max 43 u]
```

- Reduce a dense matrix to upper Hessenberg form, validate the output and store the output matrices to the disk:

```
$ ./starneig-test --experiment hessenberg --n 4000 --hooks hessenberg residual store-raw --store-raw-output
    hessenberg_%s.dat
TEST: --seed 1585762935 --experiment hessenberg --test-workers default --blas-threads default
    --lapack-threads default --scalapack-threads default --data-format pencil-local --init default --n
    4000 --data-distr default --section-height default --section-width default --solver starneig --cores
    default --gpus default --tile-size default --panel-width default --hooks hessenberg:normal
    residual:normal store-raw:normal --residual-fail-threshold 10000 --residual-warn-threshold 500
    --store-raw-output hessenberg_%s.dat --repeat 1 --warmup 0
THREADS: Using 6 StarPU worker threads during initialization and validation.
THREADS: Using 6 BLAS threads during initialization and validation.
THREADS: Using 6 BLAS threads in LAPACK solvers.
THREADS: Using 1 BLAS threads in ScaLAPACK solvers.
INIT...
PREPARE...
PROCESS...
[starneig][message] Setting tile size to 336.
[starneig][message] Setting panel width to 288.
EXPERIMENT TIME = 13121 MS
FINALIZE...
|Q ~A Q^T - A| / |A| = 15 u
|Q Q^T - I| / |I| = 11 u
WRITING TO hessenberg_A.dat...
WRITING TO hessenberg_Q.dat...
WRITING TO hessenberg_CA.dat...
================================================================
TIME = 13121 MS [avg 13121 MS, cv 0.00, min 13121 MS, max 13121 MS]
NO FAILED HESSENBERG FORM TESTS
|Q ~A Q^T - A| / |A| = [avg 15 u, cv 0.00, min 15 u, max 15 u]
|Q Q^T - I| / |I| = [avg 11 u, cv 0.00, min 11 u, max 11 u]
```

- Read an upper Hessenberg matrix from the disk, reduce it to Schur form and set tile size to `128`:

```
$ ./starneig-test --experiment schur --init read-raw --input hessenberg_%s.dat --tile-size 128
TEST: --seed 1585762972 --experiment schur --test-workers default --blas-threads default --lapack-threads
    default --scalapack-threads default --data-format pencil-local --init read-raw --input
    hessenberg_%s.dat --data-distr default --section-height default --section-width default --solver
    starneig --cores default --gpus default --iteration-limit default --tile-size 128 --small-limit
    default --aed-window-size default --aed-shift-count default --aed-nibble default
    --aed-parallel-soft-limit default --aed-parallel-hard-limit default --window-size default
    --shifts-per-window default --update-width default --update-height default --left-threshold default
    --right-threshold default --inf-threshold default --hooks schur:normal eigenvalues:normal
    known-eigenvalues:normal analysis:normal residual:normal --eigenvalues-fail-threshold 10000
    --eigenvalues-warn-threshold 1000 --known-eigenvalues-fail-threshold 1000000
    --known-eigenvalues-warn-threshold 10000 --residual-fail-threshold 10000 --residual-warn-threshold
    500 --repeat 1 --warmup 0
THREADS: Using 6 StarPU worker threads during initialization and validation.
THREADS: Using 6 BLAS threads during initialization and validation.
THREADS: Using 6 BLAS threads in LAPACK solvers.
THREADS: Using 1 BLAS threads in ScaLAPACK solvers.
INIT...
READING FROM hessenberg_A.dat...
READING A 4000 X 4000 MATRIX ...
READING FROM hessenberg_Q.dat...
READING A 4000 X 4000 MATRIX ...
READING FROM hessenberg_CA.dat...
READING A 4000 X 4000 MATRIX ...
PREPARE...
PROCESS...
[starneig][message] Using AED windows size 320.
```

```
[starneig][message] Using 240 shifts.
EXPERIMENT TIME = 9479 MS
FINALIZE...
EIGENVALUES CHECK: mean = 0 u, min = 0 u, max = 0 u
KNOWN EIGENVALUES CHECK: The stored pencil does not contain the known eigenvalues. Skipping.
EIGENVALUES ANALYSIS: zeros = 0, infinities = 0, indefinites = 0
EIGENVALUES ANALYSIS: close zeros = 0, close infinities = 0, close indefinites = 0
|Q ~A Q^T − A| / |A| = 68 u
|Q Q^T − I| / |I| = 89 u
================================================================
TIME = 9479 MS [avg 9479 MS, cv 0.00, min 9479 MS, max 9479 MS]
NO FAILED SCHUR FORM TESTS
EIGENVALUES CHECK (WARNINGS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES CHECK (FAILS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES CHECK (MEANS): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES CHECK (MIN): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES CHECK (MAX): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES ANALYSIS (ZEROS): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (CLOSE ZEROS): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (INFINITIES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (CLOSE INFINITIES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (INDEFINITES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (CLOSE INDEFINITES): [avg 0.0, cv 0.00, min 0, max 0]
|Q ~A Q^T − A| / |A| = [avg 68 u, cv 0.00, min 68 u, max 68 u]
|Q Q^T − I| / |I| = [avg 89 u, cv 0.00, min 89 u, max 89 u]
```

Distributed memory example

- Reorder a 4000 x 4000 matrix using the StarNEig implementation, use two MPI ranks, use three CPU cores per rank, distribute the matrix in 1024 x 1024 sections, and use tile size 256:

```
$ mpirun −n 2 −−map−by :PE=3 ./starneig−test −−mpi −−experiment reorder −−n 4000 −−section−height 1024
      −−section−width 1024 −−tile−size 256
```

Rank 0 output:
```
MPI INIT...
TEST: −−mpi −−mpi−mode serialized −−seed 1585763077 −−experiment reorder −−test−workers default
      −−blas−threads default −−lapack−threads default −−scalapack−threads default −−data−format
      pencil−starneig−blacs −−init default −−n 4000 −−complex−distr uniform −−complex−ratio 0.500000
      −−zero−ratio 0.010000 −−inf−ratio 0.010000 −−data−distr default −−section−height 1024 −−section−width
      1024 −−select−ratio 0.350000 −−solver starneig −−cores default −−gpus default −−tile−size 256
      −−window−size default −−values−per−chain default −−small−window−size default −−small−window−threshold
      default −−update−width default −−update−height default −−plan default −−blueprint default −−hooks
      schur:normal eigenvalues:normal analysis:normal reordering:normal residual:normal
      −−eigenvalues−fail−threshold 10000 −−eigenvalues−warn−threshold 1000 −−reordering−fail−threshold
      10000 −−reordering−warn−threshold 1000 −−residual−fail−threshold 10000 −−residual−warn−threshold 500
      −−repeat 1 −−warmup 0
THREADS: Using 3 StarPU worker threads during initialization and validation.
THREADS: Using 3 BLAS threads during initialization and validation.
THREADS: Using 3 BLAS threads in LAPACK solvers.
THREADS: Using 1 BLAS threads in ScaLAPACK solvers.
INIT...
PREPARE...
PROCESS...
[starneig][message] Attempting to set tile size to 256.
[starneig][message] Setting tile size to 256.
[starneig][message] Using multi−part task insertion plan.
[starneig][message] Using two−pass backward dummy blueprint.
[starneig][message] Using "rounded" window size.
EXPERIMENT TIME = 3320 MS
FINALIZE...
EIGENVALUES CHECK: mean = 0 u, min = 0 u, max = 0 u
EIGENVALUES ANALYSIS: zeros = 37, infinities = 0, indefinites = 0
EIGENVALUES ANALYSIS: close zeros = 0, close infinities = 0, close indefinites = 0
REORDERING CHECK: Checking selected eigenvalues...
REORDERING CHECK: Checking other eigenvalues...
REORDERING CHECK: mean = 70 u, min = 0 u, max = 527 u
|Q ~A Q^T − A| / |A| = 314 u
|Q Q^T − I| / |I| = 139 u
================================================================
TIME = 3319 MS [avg 3320 MS, cv 0.00, min 3320 MS, max 3320 MS]
NO FAILED SCHUR FORM TESTS
EIGENVALUES CHECK (WARNINGS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES CHECK (FAILS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES CHECK (MEANS): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES CHECK (MIN): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES CHECK (MAX): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES ANALYSIS (ZEROS): [avg 37.0, cv 0.00, min 37, max 37]
EIGENVALUES ANALYSIS (CLOSE ZEROS): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (INFINITIES): [avg 0.0, cv 0.00, min 0, max 0]
```

```
EIGENVALUES ANALYSIS (CLOSE INFINITIES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (INDEFINITES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (CLOSE INDEFINITES): [avg 0.0, cv 0.00, min 0, max 0]
REORDERING CHECK (WARNINGS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
REORDERING CHECK (FAILS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
REORDERING CHECK (MEANS): [avg 70 u, cv 0.00, min 70 u, max 70 u]
REORDERING CHECK (MIN): [avg 0 u, cv 0.00, min 0 u, max 0 u]
REORDERING CHECK (MAX): [avg 527 u, cv 0.00, min 527 u, max 527 u]
|Q ~A Q^T - A| / |A| = [avg 314 u, cv 0.00, min 314 u, max 314 u]
|Q Q^T - I| / |I| = [avg 139 u, cv 0.00, min 139 u, max 139 u]
```

Rank 1 output:

```
MPI INIT...
TEST: --mpi --mpi-mode serialized --seed 1585763077 --experiment reorder --test-workers default
      --blas-threads default --lapack-threads default --scalapack-threads default --data-format
      pencil-starneig-blacs --init default --n 4000 --complex-distr uniform --complex-ratio 0.500000
      --zero-ratio 0.010000 --inf-ratio 0.010000 --data-distr default --section-height 1024 --section-width
      1024 --select-ratio 0.350000 --solver starneig --cores default --gpus default --tile-size 256
      --window-size default --values-per-chain default --small-window-size default --small-window-threshold
      default --update-width default --update-height default --plan default --blueprint default --hooks
      schur:normal eigenvalues:normal analysis:normal reordering:normal residual:normal
      --eigenvalues-fail-threshold 10000 --eigenvalues-warn-threshold 1000 --reordering-fail-threshold
      10000 --reordering-warn-threshold 1000 --residual-fail-threshold 10000 --residual-warn-threshold 500
      --repeat 1 --warmup 0
THREADS: Using 3 StarPU worker threads during initialization and validation.
THREADS: Using 3 BLAS threads during initialization and validation.
THREADS: Using 3 BLAS threads in LAPACK solvers.
THREADS: Using 1 BLAS threads in ScaLAPACK solvers.
INIT...
PREPARE...
PROCESS...
[starneig][message] Attempting to set tile size to 256.
[starneig][message] Setting tile size to 256.
[starneig][message] Using multi-part task insertion plan.
[starneig][message] Using two-pass backward dummy blueprint.
[starneig][message] Using "rounded" window size.
EXPERIMENT TIME = 3320 MS
FINALIZE...
EIGENVALUES CHECK: mean = 0 u, min = 0 u, max = 0 u
EIGENVALUES ANALYSIS: zeros = 37, infinities = 0, indefinites = 0
EIGENVALUES ANALYSIS: close zeros = 0, close infinities = 0, close indefinites = 0
REORDERING CHECK: Checking selected eigenvalues...
REORDERING CHECK: Checking other eigenvalues...
REORDERING CHECK: mean = 70 u, min = 0 u, max = 527 u
|Q ~A Q^T - A| / |A| = 314 u
|Q Q^T - I| / |I| = 139 u
=============================================================
TIME = 3319 MS [avg 3320 MS, cv 0.00, min 3320 MS, max 3320 MS]
NO FAILED SCHUR FORM TESTS
EIGENVALUES CHECK (WARNINGS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES CHECK (FAILS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES CHECK (MEANS): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES CHECK (MIN): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES CHECK (MAX): [avg 0 u, cv 0.00, min 0 u, max 0 u]
EIGENVALUES ANALYSIS (ZEROS): [avg 37.0, cv 0.00, min 37, max 37]
EIGENVALUES ANALYSIS (CLOSE ZEROS): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (INFINITIES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (CLOSE INFINITIES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (INDEFINITES): [avg 0.0, cv 0.00, min 0, max 0]
EIGENVALUES ANALYSIS (CLOSE INDEFINITES): [avg 0.0, cv 0.00, min 0, max 0]
REORDERING CHECK (WARNINGS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
REORDERING CHECK (FAILS): 0 runs effected [avg 0.0, cv 0.00, min 0, max 0]
REORDERING CHECK (MEANS): [avg 70 u, cv 0.00, min 70 u, max 70 u]
REORDERING CHECK (MIN): [avg 0 u, cv 0.00, min 0 u, max 0 u]
REORDERING CHECK (MAX): [avg 527 u, cv 0.00, min 527 u, max 527 u]
|Q ~A Q^T - A| / |A| = [avg 314 u, cv 0.00, min 314 u, max 314 u]
|Q Q^T - I| / |I| = [avg 139 u, cv 0.00, min 139 u, max 139 u]
```

Chapter 9

Known problems and changelog

Known problems

- If the CUDA support is enabled, then the starneig_SEP_SM_Hessenberg() interface function cannot always handle problems that do not fit into GPU's memory. The cause of this problem is not known.

- The outputs of the starneig_GEP_SM_Schur() and starneig_GEP_DM_Schur() interface functions are not always in the so-called standard format. It is possible that some diagonal entries in the right-hand side output matrix are negative.

- The starneig_GEP_SM_Eigenvectors() interface function may scale the input matrices.

Known compatibility problems

BLAS

- With some OpenBLAS versions, it is necessary to set the `OPENBLAS_NUM_THREADS` environmental variable to value `1` (`export OPENBLAS_NUM_THREADS=1`).

- Some MKL versions can cause poor scalability. The problem appears to be related to Intel's OpenMP library. Setting the `KMP_AFFINITY` environmental variable to value `disabled` fixes the problem (`export K↩MP_AFFINITY=disabled`).

- OpenBLAS version 0.3.1 has a bug that can cause an incorrect result.

- OpenBLAS versions 0.3.3-0.3.5 can cause poor scalability.

MPI

- Some older OpenMPI versions ($<$= 2.1.1) have a bug that can cause a segmentation fault during a parallel AED.

- The library has an unsolved memory leak problem with OpenMPI. Only large problem sizes are effected. It is not known whether this problem is related to StarNEig, StarPU, OpenMPI or something else. The problem is known to occur with PMIx 2.2.1, UCX 1.5.0, OpenMPI 3.1.3, and StarPU 1.2.8. The memory leak is sometimes accompanied by the following warning:
  `mpool.c:38   UCX  WARN  object 0x2652000 was not returned to mpool ucp_requests`

- The test program can trigger the following bug in UCX 1.6.1:   https://github.com/openucx/ucx/issues/4525

StarPU

- For optimal CUDA performance, StarPU version that is newer than 1.3.3 is recommended.

- StarPU versions 1.2.4 - 1.2.8 and some StarPU 1.3 snapshots cav cause poor CUDA performance. The problem can be fixed by compiling StarPU with `--disable-cuda-memcpy-peer`. It is possible that newer versions of StarPU are also effected by this problem.

- The `STARPU_MINIMUM_AVAILABLE_MEM` and `STARPU_TARGET_AVAILABLE_MEM` environmental variables can be used to fix some GPU-related memory allocation problems:
  `STARPU_MINIMUM_AVAILABLE_MEM=10 STARPU_TARGET_AVAILABLE_MEM=15 ...`

Changelog

Planned for v0.1.6

- Minor improvements to the Hessenberg reduction step

v0.1.5

- Disable OpenCL workers during initialization.

- Add `STARNEIG_USE_ALL`.

v0.1.4

- Add deb packages.

v0.1.3

- Restore older Hessenberg reduction implementation from v0.1-beta2

- Rename `aed_shift_count` parameter as `shift_count`. Rename the default value `STARNEIG_S`↩
  `CHUR_DEFAULT_AED_SHIFT_COUNT` as `STARNEIG_SCHUR_DEFAULT_SHIFT_COUNT`.

- Improved performance models.

v0.1.2

- Improved performance models.

- Updates to the documentation.

- Rename `STARPU_LIBRARIES_BASE` and `STARPU_LIBRARIES_MPI` environmental variables as `S`↩
  `TARPU_LIBRARIES` and `STARPU_MPI_LIBRARIES`, respectively.

v0.1.1:

- Fix a bug that may cause the code to hang in distributed memory Schur reduction.

v0.1.0:

- First stable release of the library.

v0.1-beta.6:

- Fix `starneig_node_enable_pinning()` and `starneig_node_disable_pinning()` functions.

- Built `pdgghrd` as a separate library.

- Deprecate several precompiler defines and interface functions. Add `<starneig/distr_helpers.h>` header file.

v0.1-beta.5:

- Improve the performance of the Hessenberg reduction phase by limiting the number of submitted tasks. This should reduce the task scheduling overhead.

- Allocate pinned memory by default when CUDA support is enabled. Add `starneig_enable_↩ pinning()` and `starneig_disable_pinning()`.

v0.1-beta.4:

- Fix a problem where infinite eigenvalues were detected too late.

- Add an option to choose between the norm stable deflation condition (`STARNEIG_SCHUR_NORM_STAB↩ LE_THRESHOLD`) and and the LAPACK style deflation condition (`STARNEIG_SCHUR_LAPACK_THRE↩ SHOLD`).

v0.1-beta.3:

- Re-implemented Hessenberg reduction.

v0.1-beta.2:

- Fix an installation-related bug.

- Fix a MPI-related compile error.

- Remove unused code.

v0.1-beta.1:

- First beta release of the library.

Chapter 10

License, authors and literature

License

Authors

The following people have contributed to the StarNEig library:

- Angelika Schwarz (`angies@cs.umu.se`)
    - Standard eigenvectors
- Bo Kågström (`bokg@cs.umu.se`)

- **–** Coordinator and scientific director for the NLAFET project

  – Documentation

- Carl Christian Kjelgaard Mikkelsen (`spock@cs.umu.se`)

  – Generalized eigenvectors

- Lars Karlsson (`larsk@cs.umu.se`)

  – Miscellaneous user interface functions

  – Documentation

- Mirko Myllykoski (`mirkom@cs.umu.se`)

  – Hessenberg reduction

  – Schur reduction (standard and generalized)

  – Eigenvalue reordering (standard and generalized)

  – Miscellaneous user interface functions

  – Test program

  – Documentation

The distributed memory Hessenberg-triangular reduction is implemented as a ScaLAPACK wrapper and `PDGGHRD` subroutine (Lars Karlsson, Björn Adlerborn) is included with the library. The test program includes Matrix Market I/O library for ANSI C (`http://math.nist.gov/MatrixMarket`).

Related publications

Research papers

- Mirko Myllykoski: *A Task-based Multi-shift QR/QZ Algorithm with Aggressive Early Deflation*, `arXiv↩ :2007.03576`

- Mirko Myllykoski, Carl Christian Kjelgaard Mikkelsen: *Task-based, GPU-accelerated and Robust Library for Solving Dense Nonsymmetric Eigenvalue Problems*, Concurrency and Computation: Practice and Experience, 2020, doi: `10.1002/cpe.5915`

- Mirko Myllykoski, Carl Christian Kjelgaard Mikkelsen: *Introduction to StarNEig — A Task-based Library for Solving Nonsymmetric Eigenvalue Problems*, In Parallel Processing and Applied Mathematics, 13th International Conference, PPAM 2019, Bialystok, Poland, September 8–11, 2019, Revised Selected Papers, Part I, Lecture Notes in Computer Science, Vol. 12043, Wyrzykowski R., Deelman E., Dongarra J., Karczewski K. (eds), Springer International Publishing, pp. 70-81, 2020, doi: `10.1007/978-3-030-43229-4_7`

- Carl Christian Kjelgaard Mikkelsen, Mirko Myllykoski: *Parallel Robust Computation of Generalized Eigenvectors of Matrix Pencils*, presented at PPAM 2019, In Parallel Processing and Applied Mathematics, 13th International Conference, PPAM 2019, Bialystok, Poland, September 8–11, 2019, Revised Selected Papers, Part I, Lecture Notes in Computer Science, Vol. 12043, Wyrzykowski R., Deelman E., Dongarra J., Karczewski K. (eds), Springer International Publishing, pp. 58-69, 2020, doi: `10.1007/978-3-030-43229-4_6`

- Carl Christian Kjelgaard Mikkelsen, Angelika Schwarz, Lars Karlsson: *Parallel Robust Solution of Triangular Linear Systems*, Concurrency and Computation: Practice and Experience, 31 (19), 2019, doi: `10.↩ 1002/cpe.5064`

- Angelika Schwarz, Lars Karlsson: *Scalable eigenvector computation for the non-symmetric eigenvalue problem*, Parallel Computing, 85, pp. 131-140, 2019, doi: `10.1016/j.parco.2019.04.001`

- Mirko Myllykoski: *A Task-Based Algorithm for Reordering the Eigenvalues of a Matrix in Real Schur Form*, In Parallel Processing and Applied Mathematics, 12th International Conference, PPAM 2017, Lublin, Poland, September 10-13, 2017, Revised Selected Papers, Part I, Lecture Notes in Computer Science, Vol. 10777, Wyrzykowski R., Dongarra J., Deelman E., Karczewski K. (eds), Springer International Publishing, pp. 207-216, 2018, doi: 10.1007/978-3-319-78024-5_19

- Carl Christian Kjelgaard Mikkelsen, Lars Karlsson. *Blocked Algorithms for Robust Solution of Triangular Linear Systems*, In Parallel Processing and Applied Mathematics, 12th International Conference, PPAM 2017, Lublin, Poland, September 10-13, 2017, Revised Selected Papers, Part I, Lecture Notes in Computer Science, Vol. 10777, Wyrzykowski R., Dongarra J., Deelman E., Karczewski K. (eds), Springer International Publishing, pp. 207-216, 2018, doi: 10.1007/978-3-319-78024-5_7

Reports, deliverables etc

- Angelika Schwarz, Carl Christian Kjelgaard Mikkelsen, Lars Karlsson: *Robust Parallel Eigenvector Computation For the Non-Symmetric Eigenvalue Problem*, Report UMINF 20.02, Department of Computing Science, Umeå University, SE-901 87 Umeå, Sweden, 2020 (download)

- Angelika Schwarz: *Towards efficient overflow-free solvers for systems of triangular type*, Licentiate thesis, Department of computing science, Umeå University, ISSN: 0348-0542, 2019

- Mirko Myllykoski, Carl Christian Kjelgaard Mikkelsen, Angelika Schwarz, Bo Kågström: *D2.7 Eigenvalue solvers for nonsymmetric problems*, public NLAFET deliverable, 2019 (download)

- Lars Karlsson, Mahmoud Eljammaly, Mirko Myllykoski: *D6.5 Evaluation of auto-tuning techniques*, public NLAFET deliverable, 2019 (download)

- Bo Kågström et al.: *D7.8 Release of the NLAFET library*, public NLAFET deliverable, 2019 (download)

- Mirko Myllykoski, Lars Karlsson, Bo Kågström, Mahmoud Eljammaly, Srikara Pranesh, Mawussi Zounon: *D2.6 Prototype Software for Eigenvalue Problem Solvers*, public NLAFET deliverable, 2018 (download)

- Mirko Myllykoski, Carl Christian Kjelgaard Mikkelsen, Lars Karlsson, Bo Kågström: *Task-Based Parallel Algorithms for Reordering of Matrices in Real Schur Forms*, NLAFET Working Note WN-11, 2017. Also as Report UMINF 17.11, Department of Computing Science, Umeå University, SE-901 87 Umeå, Sweden (download)

- Carl Christian Kjelgaard Mikkelsen, Mirko Myllykoski, Björn Adlerborn, Lars Karlsson, Bo Kågström: *D2.5 Eigenvalue Problem Solvers*, public NLAFET deliverable, 2017 (download)

Youtube videos

- Task-based Schur reduction, N = 20 000, shared memory, 28 cores

- Task-based Hessenberg reduction, N = 20 000, shared memory, 6 cores

Chapter 11

Todo List

Global starneig_GEP_DM_Eigenvectors (int selected[], starneig_distr_matrix_t S, starneig_distr_matrix_t T, starneig_distr_matrix_t Z, starneig_distr_matrix_t X)

This interface function is not implemented.

Global starneig_GEP_DM_Eigenvectors_expert (struct starneig_eigenvectors_conf ∗conf, int selected[], starneig_distr_matrix_t S, starneig_distr_matrix_t T, starneig_distr_matrix_t Z, starneig_distr_matrix_t X)

This interface function is not implemented.

Global starneig_SEP_DM_Eigenvectors (int selected[], starneig_distr_matrix_t S, starneig_distr_matrix_t Q, starneig_distr_matrix_t X)

This interface function is not implemented.

Global starneig_SEP_DM_Eigenvectors_expert (struct starneig_eigenvectors_conf ∗conf, int selected[], starneig_distr_matrix_t S, starneig_distr_matrix_t Q, starneig_distr_matrix_t X)

This interface function is not implemented.

Chapter 12

Deprecated List

Global starneig_broadcast (int root, size_t size, void ∗buffer)

The starneig_broadcast() function has been replaced with the starneig_mpi_broadcast() function. This function will be removed in a future release of the library.

Global starneig_create_blacs_matrix (int rows, int cols, int row_blksz, int col_blksz, starneig_datatype_t type, starneig_blacs_context_t context, starneig_blacs_descr_t ∗descr, void ∗∗local)

The starneig_create_blacs_matrix() function has been replaced with the starneig_blacs_create_matrix() function. This function will be removed in a future release of the library.

Global starneig_descinit (struct starneig_blacs_descr ∗descr, int m, int n, int sm, int sn, int irsrc, int icsrc, starneig_blacs_context_t context, int ld)

The starneig_descinit() function has been replaced with the starneig_blacs_descinit() function. This function will be removed in a future release of the library.

Global starneig_destroy_blacs_matrix (starneig_blacs_descr_t ∗descr, void ∗∗local)

The starneig_destroy_blacs_matrix() function has been replaced with the starneig_blacs_destroy_matrix() function. This function will be removed in a future release of the library.

Global starneig_numroc (int n, int nb, int iproc, int isrcproc, int nprocs)

The starneig_numroc() function has been replaced with the starneig_blacs_numroc() function. This function will be removed in a future release of the library.

Chapter 13

Module Documentation

13.1   Library configuration

Configuration of the installed library.

Macros

- #define STARNEIG_VERSION_MAJOR

    StarNEig major version.
- #define STARNEIG_VERSION_MINOR

    StarNEig minor version.
- #define STARNEIG_VERSION_PATCH

    StarNEig patch level.
- #define STARNEIG_ENABLE_MPI

    MPI support enabled.
- #define STARNEIG_ENABLE_CUDA

    CUDA support enabled.
- #define STARNEIG_ENABLE_BLACS

    BLACS support enabled.
- #define STARNEIG_SEP_DM_HESSENBERG

    Distributed memory Hessenberg reduction enabled.
- #define STARNEIG_GEP_DM_HESSENBERGTRIANGULAR

    Distributed memory Hessenberg-triangular reduction enabled.
- #define STARNEIG_SEP_DM_REDUCE

    Full distributed memory reduction enabled.
- #define STARNEIG_GEP_DM_REDUCE

    Full distributed memory generalized reduction enabled.

13.1.1   Detailed Description

Configuration of the installed library.

13.1.2 Macro Definition Documentation

13.1.2.1 STARNEIG_ENABLE_MPI

`#define STARNEIG_ENABLE_MPI`

MPI support enabled.

Defined if the library was compiled with MPI support.

13.1.2.2 STARNEIG_ENABLE_CUDA

`#define STARNEIG_ENABLE_CUDA`

CUDA support enabled.

Defined if the library was compiled with CUDA support.

13.1.2.3 STARNEIG_ENABLE_BLACS

`#define STARNEIG_ENABLE_BLACS`

BLACS support enabled.

Defined if the library was compiled with ScaLAPACK compatibility layer.

13.1.2.4 STARNEIG_SEP_DM_HESSENBERG

`#define STARNEIG_SEP_DM_HESSENBERG`

Distributed memory Hessenberg reduction enabled.

Defined if the starneig_SEP_DM_Hessenberg() function exists.

13.1.2.5 STARNEIG_GEP_DM_HESSENBERGTRIANGULAR

`#define STARNEIG_GEP_DM_HESSENBERGTRIANGULAR`

Distributed memory Hessenberg-triangular reduction enabled.

Defined if the starneig_GEP_DM_HessenbergTriangular() function exists.

13.1.2.6 STARNEIG_SEP_DM_REDUCE

`#define STARNEIG_SEP_DM_REDUCE`

Full distributed memory reduction enabled.

Defined if the starneig_SEP_DM_Reduce() function exists.

13.1.2.7 STARNEIG_GEP_DM_REDUCE

`#define STARNEIG_GEP_DM_REDUCE`

Full distributed memory generalized reduction enabled.

Defined if the starneig_GEP_DM_Reduce() function exists.

13.2 Error codes

Interface function return values and error codes.

Macros

- #define STARNEIG_SUCCESS 0

  Success.
- #define STARNEIG_GENERIC_ERROR 1

  Reneric error.
- #define STARNEIG_NOT_INITIALIZED 2

  Not initialized.
- #define STARNEIG_INVALID_CONFIGURATION 3

  Invalid configuration.
- #define STARNEIG_INVALID_ARGUMENTS 4

  Invalid argument.
- #define STARNEIG_INVALID_DISTR_MATRIX 5

  Invalid distributed matrix.
- #define STARNEIG_DID_NOT_CONVERGE 6

  Did not converge.
- #define STARNEIG_PARTIAL_REORDERING 7

  Partial reordering.
- #define STARNEIG_CLOSE_EIGENVALUES 8

  Close eigenvalues.

Typedefs

- typedef int starneig_error_t

  Interface function return value data type.

13.2.1 Detailed Description

Interface function return values and error codes.

13.2.2 Macro Definition Documentation

13.2.2.1 STARNEIG_SUCCESS

```
#define STARNEIG_SUCCESS 0
```

Success.

The interface function was executed successfully.

13.2.2.2 STARNEIG_GENERIC_ERROR

`#define STARNEIG_GENERIC_ERROR 1`

Reneric error.

The interface function encountered a generic error.

13.2.2.3 STARNEIG_NOT_INITIALIZED

`#define STARNEIG_NOT_INITIALIZED 2`

Not initialized.

The library was not initialized when the interface function was called.

13.2.2.4 STARNEIG_INVALID_CONFIGURATION

`#define STARNEIG_INVALID_CONFIGURATION 3`

Invalid configuration.

The interface function encountered an invalid configuration argument.

13.2.2.5 STARNEIG_INVALID_ARGUMENTS

`#define STARNEIG_INVALID_ARGUMENTS 4`

Invalid argument.

The interface function encountered an invalid argument.

13.2.2.6 STARNEIG_INVALID_DISTR_MATRIX

`#define STARNEIG_INVALID_DISTR_MATRIX 5`

Invalid distributed matrix.

One or more of the involved distributed matrices have an invalid distribution, invalid dimensions and/or an invalid distributed block size.

13.2.2.7 STARNEIG_DID_NOT_CONVERGE

`#define STARNEIG_DID_NOT_CONVERGE 6`

Did not converge.

The interface function encountered a situation where the QR/QZ algorithm did not converge. The matrix (pair) may be partially in (generalized) Schur form.

13.2.2.8 STARNEIG_PARTIAL_REORDERING

`#define STARNEIG_PARTIAL_REORDERING 7`

Partial reordering.

The interface function failed to reorder the (generalized) Schur form. The (generalized) Schur form may be partially reordered.

13.2.2.9 STARNEIG_CLOSE_EIGENVALUES

`#define STARNEIG_CLOSE_EIGENVALUES 8`

Close eigenvalues.

The interface function encountered a situation where two selected eigenvalues were close to each other. The computed result may be inaccurate.

13.3 Intra-node execution environment

Interface to configure the intra-node execution environment.

Functions

- void starneig_node_init (int cores, int gpus, starneig_flag_t flags)

  Initializes the intra-node execution environment.
- int starneig_node_initialized ()

  Checks whether the intra-node execution environment is initialized.
- int starneig_node_get_cores ()

  Returns the number of cores (threads) per MPI rank.
- void starneig_node_set_cores (int cores)

  Changes the number of CPUs cores (threads) to use per MPI rank.
- int starneig_node_get_gpus ()

  Returns the number of GPUs per MPI rank.
- void starneig_node_set_gpus (int gpus)

  Changes the number of GPUs to use per MPI rank.
- void starneig_node_finalize ()

  Deallocates resources associated with the intra-node configuration.

Library initialization flags

- typedef unsigned starneig_flag_t

  Library initialization flag data type.
- #define STARNEIG_USE_ALL -1

  Use all resources.
- #define STARNEIG_DEFAULT 0x0

  Default mode.
- #define STARNEIG_HINT_SM 0x0

  Shared memory mode.
- #define STARNEIG_HINT_DM 0x1

  Distributed memory mode.
- #define STARNEIG_FXT_DISABLE 0x2

  No FxT traces mode.
- #define STARNEIG_AWAKE_WORKERS 0x4

  Awake worker mode.
- #define STARNEIG_AWAKE_MPI_WORKER 0x8

  Awake MPI worker mode.
- #define STARNEIG_FAST_DM (STARNEIG_HINT_DM | STARNEIG_AWAKE_WORKERS | STARNEIG_AWAKE_MPI_WORK

  Fast distributed memory mode.
- #define STARNEIG_NO_VERBOSE 0x10

  No verbose mode.
- #define STARNEIG_NO_MESSAGES (STARNEIG_NO_VERBOSE | 0x20)

  No messages mode.

Pinned host memory

- void starneig_node_enable_pinning ()

  Enable CUDA host memory pinning.
- void starneig_node_disable_pinning ()

  Disables CUDA host memory pinning.

13.3.1 Detailed Description

Interface to configure the intra-node execution environment.

13.3.2 Macro Definition Documentation

13.3.2.1 STARNEIG_USE_ALL

```
#define STARNEIG_USE_ALL -1
```

Use all resources.

Tells StarNEig to use all available CPU cores / GPUs.

Examples

gep_dm_full_chain.c, gep_sm_eigenvectors.c, gep_sm_full_chain.c, sep_dm_full_chain.c, sep_sm_eigenvectors.c, and sep_sm_full_chain.c.

13.3.2.2 STARNEIG_DEFAULT

```
#define STARNEIG_DEFAULT 0x0
```

Default mode.

As a default, the library configures itself to shared memory mode.

13.3.2.3 STARNEIG_HINT_SM

```
#define STARNEIG_HINT_SM 0x0
```

Shared memory mode.

Initializes the library for shared memory computation. The library will automatically reconfigure itself for distributed memory computation if necessary

Examples

gep_sm_eigenvectors.c, gep_sm_full_chain.c, sep_sm_eigenvectors.c, and sep_sm_full_chain.c.

13.3.2.4 STARNEIG_HINT_DM

```
#define STARNEIG_HINT_DM 0x1
```

Distributed memory mode.

Initializes the library for distributed memory computation. The library will automatically reconfigure itself for shared memory computation if necessary

Examples

> sep_dm_full_chain.c.

13.3.2.5 STARNEIG_FXT_DISABLE

```
#define STARNEIG_FXT_DISABLE 0x2
```

No FxT traces mode.

Disables FXT traces.

Attention

> This flag does not work reliably with all StarPU versions.

13.3.2.6 STARNEIG_AWAKE_WORKERS

```
#define STARNEIG_AWAKE_WORKERS 0x4
```

Awake worker mode.

Keeps the StarPU worker threads awake between interface function calls. Improves the performance in certain situations but can interfere with other software.

Examples

> gep_sm_full_chain.c.

13.3.2.7 STARNEIG_AWAKE_MPI_WORKER

```
#define STARNEIG_AWAKE_MPI_WORKER 0x8
```

Awake MPI worker mode.

Keeps the StarPU-MPI communication thread awake between interface function calls. Improves the performance in certain situations but can interfere with other software.

13.3.2.8 STARNEIG_FAST_DM

```
#define STARNEIG_FAST_DM (STARNEIG_HINT_DM | STARNEIG_AWAKE_WORKERS | STARNEIG_AWAKE_MPI_WORKER)
```

Fast distributed memory mode.

Keeps the worker threads and StarPU-MPI communication thread awake between interface function calls. Improves the performance in certain situations but can interfere with other software.

Examples

  gep_dm_full_chain.c.

13.3.2.9 STARNEIG_NO_VERBOSE

```
#define STARNEIG_NO_VERBOSE 0x10
```

No verbose mode.

Disables all additional verbose messages.

13.3.2.10 STARNEIG_NO_MESSAGES

```
#define STARNEIG_NO_MESSAGES (STARNEIG_NO_VERBOSE | 0x20)
```

No messages mode.

Disables all messages (including verbose messages).

13.3.3 Function Documentation

13.3.3.1 starneig_node_init()

```
void starneig_node_init (
            int cores,
            int gpus,
            starneig_flag_t flags )
```

Initializes the intra-node execution environment.

The interface function initializes StarPU (and cuBLAS) and pauses all worker The `cores` argument specifies the **total number of used CPU cores**. In distributed memory mode, one CPU core is automatically allocated for the StarPU-MPI communication thread. One or more CPU cores are automatically allocated for GPU devices.

Parameters

in	*cores*	The number of cores (threads) to use per MPI rank. Can be set to STARNEIG_USE_ALL in which case the library uses all available cores.
in	*gpus*	The number of GPUs to use per MPI rank. Can be set to STARNEIG_USE_ALL in which case the library uses all available GPUs.
in	*flags*	Initialization flags.

Examples

> gep_dm_full_chain.c, gep_sm_eigenvectors.c, gep_sm_full_chain.c, sep_dm_full_chain.c, sep_sm_eigenvectors.c, and sep_sm_full_chain.c.

13.3.3.2 starneig_node_initialized()

```
int starneig_node_initialized ( )
```

Checks whether the intra-node execution environment is initialized.

Returns

> Non-zero if the environment is initialized, 0 otherwise.

13.3.3.3 starneig_node_get_cores()

```
int starneig_node_get_cores ( )
```

Returns the number of cores (threads) per MPI rank.

Returns

> The number of cores (threads) per MPI rank.

13.3.3.4 starneig_node_set_cores()

```
void starneig_node_set_cores (
            int cores )
```

Changes the number of CPUs cores (threads) to use per MPI rank.

Parameters

cores	The number of CPUs to use per MPI rank.

13.3.3.5 starneig_node_get_gpus()

```
int starneig_node_get_gpus ( )
```

Returns the number of GPUs per MPI rank.

Returns

> The number of GPUs per MPI rank.

13.3.3.6 starneig_node_set_gpus()

```
void starneig_node_set_gpus (
            int gpus )
```

Changes the number of GPUs to use per MPI rank.

Parameters

gpus	The number of GPUs to use per MPI rank.

13.3.3.7 starneig_node_enable_pinning()

```
void starneig_node_enable_pinning ( )
```

Enable CUDA host memory pinning.

Should be called before any memory allocations are made.

13.3.3.8 starneig_node_disable_pinning()

```
void starneig_node_disable_pinning ( )
```

Disables CUDA host memory pinning.

Should be called before any memory allocations are made.

13.4 Shared Memory / Standard EVP

Functions for solving non-symmetric standard eigenvalue problems on shared memory systems.

Computational functions

- starneig_error_t starneig_SEP_SM_Hessenberg (int n, double A[ ], int ldA, double Q[ ], int ldQ)

  Computes a Hessenberg decomposition of a general matrix.

- starneig_error_t starneig_SEP_SM_Schur (int n, double H[ ], int ldH, double Q[ ], int ldQ, double real[ ], double imag[ ])

  Computes a Schur decomposition given a Hessenberg decomposition.

- starneig_error_t starneig_SEP_SM_ReorderSchur (int n, int selected[ ], double S[ ], int ldS, double Q[ ], int ldQ, double real[ ], double imag[ ])

  Reorders selected eigenvalues to the top left corner of a Schur decomposition.

- starneig_error_t starneig_SEP_SM_Reduce (int n, double A[ ], int ldA, double Q[ ], int ldQ, double real[ ], double imag[ ], int(∗predicate)(double real, double imag, void ∗arg), void ∗arg, int selected[ ], int ∗num_↩
selected)

  Computes a (reordered) Schur decomposition of a general matrix.

- starneig_error_t starneig_SEP_SM_Eigenvectors (int n, int selected[ ], double S[ ], int ldS, double Q[ ], int ldQ, double X[ ], int ldX)

  Computes an eigenvector for each selected eigenvalue.

Helper functions

- starneig_error_t starneig_SEP_SM_Select (int n, double S[ ], int ldS, int(∗predicate)(double real, double imag, void ∗arg), void ∗arg, int selected[ ], int ∗num_selected)

  Generates a selection array for a Schur matrix using a user-supplied predicate function.

Expert computational functions

- starneig_error_t starneig_SEP_SM_Hessenberg_expert (struct starneig_hessenberg_conf ∗conf, int n, int begin, int end, double A[ ], int ldA, double Q[ ], int ldQ)

  Computes a Hessenberg decomposition of a general matrix.

- starneig_error_t starneig_SEP_SM_Schur_expert (struct starneig_schur_conf ∗conf, int n, double H[ ], int ldH, double Q[ ], int ldQ, double real[ ], double imag[ ])

  Computes a Schur decomposition given a Hessenberg decomposition.

- starneig_error_t starneig_SEP_SM_ReorderSchur_expert (struct starneig_reorder_conf ∗conf, int n, int selected[ ], double S[ ], int ldS, double Q[ ], int ldQ, double real[ ], double imag[ ])

  Reorders selected eigenvalues to the top left corner of a Schur decomposition.

- starneig_error_t starneig_SEP_SM_Eigenvectors_expert (struct starneig_eigenvectors_conf ∗conf, int n, int selected[ ], double S[ ], int ldS, double Q[ ], int ldQ, double X[ ], int ldX)

  Computes an eigenvector for each selected eigenvalue.

13.4.1 Detailed Description

Functions for solving non-symmetric standard eigenvalue problems on shared memory systems.

13.4.2 Function Documentation

13.4.2.1 starneig_SEP_SM_Hessenberg()

starneig_error_t starneig_SEP_SM_Hessenberg (
        int *n,*
        double *A[ ],*
        int *ldA,*
        double *Q[ ],*
        int *ldQ*)

Computes a Hessenberg decomposition of a general matrix.

Parameters

in	*n*	The order of A and Q.
in,out	*A*	On entry, the general matrix A. On exit, the upper Hessenberg matrix H.
in	*ldA*	The leading dimension of A.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U$.
in	*ldQ*	The leading dimension of Q.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

Examples

sep_sm_full_chain.c.

13.4.2.2 starneig_SEP_SM_Schur()

starneig_error_t starneig_SEP_SM_Schur (
        int *n,*
        double *H[ ],*
        int *ldH,*
        double *Q[ ],*
        int *ldQ,*
        double *real[ ],*
        double *imag[ ]*)

Computes a Schur decomposition given a Hessenberg decomposition.

Parameters

in	*n*	The order of H and Q.
in,out	*H*	On entry, the upper Hessenberg matrix H. On exit, the Schur matrix S.

Parameters

in	*ldH*	The leading dimension of H.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U$.
in	*ldQ*	The leading dimension of Q.
out	*real*	An array of the same size as H containing the real parts of the computed eigenvalues.
out	*imag*	An array of the same size as H containing the imaginary parts of the computed eigenvalues.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise. STARNEIG_DID_NOT_CONVERGE if the QR algorithm failed to converge.

Examples

sep_sm_full_chain.c.

13.4.2.3 starneig_SEP_SM_ReorderSchur()

```
starneig_error_t starneig_SEP_SM_ReorderSchur (
            int n,
            int selected[],
            double S[],
            int ldS,
            double Q[],
            int ldQ,
            double real[],
            double imag[] )
```

Reorders selected eigenvalues to the top left corner of a Schur decomposition.

Parameters

in	*n*	The order of S and Q.
in,out	*selected*	The selection array. On entry, the initial positions of the selected eigenvalues. On exit, the final positions of all correctly placed selected eigenvalues. In case of failure, the number of 1's in the output may be less than the number of 1's in the input.
in,out	*S*	On entry, the Schur matrix S. On exit, the updated Schur matrix $\hat{S}$.
in	*ldS*	The leading dimension of S.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U$.
in	*ldQ*	The leading dimension of Q.
out	*real*	An array of the same size as S containing the real parts of the computed eigenvalues.
out	*imag*	An array of the same size as S containing the imaginary parts of the computed eigenvalues.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise. STARNEIG_PARTIAL_REORDERING if the Schur form is not fully reordered.

See also

starneig_SEP_SM_Select

Examples

sep_sm_full_chain.c.

13.4.2.4 starneig_SEP_SM_Reduce()

```
starneig_error_t starneig_SEP_SM_Reduce (
            int n,
            double A[],
            int ldA,
            double Q[],
            int ldQ,
            double real[],
            double imag[],
            int(*)(double real, double imag, void *arg) predicate,
            void * arg,
            int selected[],
            int * num_selected )
```

Computes a (reordered) Schur decomposition of a general matrix.

Parameters

`in`	*n*	The order of A and Q.
`in,out`	*A*	On entry, the general matrix A. On exit, the Schur matrix S.
`in`	*ldA*	The leading dimension of A.
`in,out`	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U$.
`in`	*ldQ*	The leading dimension of Q.
`out`	*real*	An array of the same size as A containing the real parts of the computed eigenvalues.
`out`	*imag*	An array of the same size as A containing the imaginary parts of the computed eigenvalues.
`in`	*predicate*	A function that takes a (complex) eigenvalue as input and returns non-zero if it should be selected. For complex conjugate pairs of eigenvalues, the predicate is called only for the eigenvalue with positive imaginary part and the corresponding 2×2 block is either selected or deselected. The reordering step is skipped if the argument is a NULL pointer.
`in`	*arg*	An optional argument for the predicate function.
`out`	*selected*	The final positions of all correctly placed selected eigenvalues.
`out`	*num_selected*	The number of selected eigenvalues (a complex conjugate pair is counted as two selected eigenvalues).

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise. STARNEIG_DID_NOT_CONVERGE if the QR algorithm failed to converge. STARNEIG_PARTIAL_REORDERING if the Schur form is not fully reordered.

Examples

sep_sm_eigenvectors.c.

13.4.2.5 starneig_SEP_SM_Eigenvectors()

```
starneig_error_t starneig_SEP_SM_Eigenvectors (
        int n,
        int selected[],
        double S[],
        int ldS,
        double Q[],
        int ldQ,
        double X[],
        int ldX )
```

Computes an eigenvector for each selected eigenvalue.

Parameters

in	*n*	The order of S and Q and the number of rows of X.
in	*selected*	The selection array specifying the locations of the selected eigenvalues. The number of 1's in the array is the same as the number of columns in X.
in	*S*	The Schur matrix S.
in	*ldS*	The leading dimension of S.
in	*Q*	The orthogonal matrix Q.
in	*ldQ*	The The leading dimension of Q.
out	*X*	A matrix with n rows and one column for each selected eigenvalue. The columns represent the computed eigenvectors as previously described.
in	*ldX*	The leading dimension of X.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

starneig_SEP_SM_Select

Examples

sep_sm_eigenvectors.c.

13.4.2.6   starneig_SEP_SM_Select()

```
starneig_error_t starneig_SEP_SM_Select (
            int n,
            double S[],
            int ldS,
            int(*)(double real, double imag, void *arg) predicate,
            void * arg,
            int selected[],
            int * num_selected )
```

Generates a selection array for a Schur matrix using a user-supplied predicate function.

Parameters

in	n	The order of S.
in	S	The Schur matrix S.
in	ldS	The leading dimension of S.
in	predicate	A function that takes a (complex) eigenvalue as input and returns non-zero if it should be selected. For complex conjugate pairs of eigenvalues, the predicate is called only for the eigenvalue with positive imaginary part and the corresponding 2×2 block is either selected or deselected.
in	arg	An optional argument for the predicate function.
out	selected	The selection array. Both elements of a selected complex conjugate pair are set to 1.
out	num_selected	The number of selected eigenvalues (a complex conjugate pair is counted as two selected eigenvalues).

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

Examples

sep_sm_eigenvectors.c, and sep_sm_full_chain.c.

13.4.2.7   starneig_SEP_SM_Hessenberg_expert()

```
starneig_error_t starneig_SEP_SM_Hessenberg_expert (
            struct starneig_hessenberg_conf * conf,
            int n,
            int begin,
            int end,
            double A[],
            int ldA,
            double Q[],
            int ldQ )
```

Computes a Hessenberg decomposition of a general matrix.

Parameters

in	*conf*	Configuration structure.
in	*n*	The order of A and Q.
in	*begin*	First column to be reduced.
in	*end*	Last column to be reduced + 1.
in,out	*A*	On entry, the general matrix A. On exit, the upper Hessenberg matrix H.
in	*ldA*	The leading dimension of A.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U$.
in	*ldQ*	The leading dimension of Q.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

starneig_SEP_SM_Hessenberg

starneig_hessenberg_conf

starneig_hessenberg_init_conf

13.4.2.8 starneig_SEP_SM_Schur_expert()

```
starneig_error_t starneig_SEP_SM_Schur_expert (
            struct starneig_schur_conf * conf,
            int n,
            double H[],
            int ldH,
            double Q[],
            int ldQ,
            double real[],
            double imag[] )
```

Computes a Schur decomposition given a Hessenberg decomposition.

Parameters

in	*conf*	Configuration structure.
in	*n*	The order of H and Q.
in,out	*H*	On entry, the upper Hessenberg matrix H. On exit, the Schur matrix S.
in	*ldH*	The leading dimension of H.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U$.
in	*ldQ*	The leading dimension of Q.
out	*real*	An array of the same size as H containing the real parts of the computed eigenvalues.
out	*imag*	An array of the same size as H containing the imaginary parts of the computed eigenvalues.

Returns

[STARNEIG_SUCCESS](#) (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

[starneig_SEP_SM_Schur](#)

[starneig_schur_conf](#)

[starneig_schur_init_conf](#)

13.4.2.9 starneig_SEP_SM_ReorderSchur_expert()

```
starneig_error_t starneig_SEP_SM_ReorderSchur_expert (
            struct starneig_reorder_conf * conf,
            int n,
            int selected[],
            double S[],
            int ldS,
            double Q[],
            int ldQ,
            double real[],
            double imag[] )
```

Reorders selected eigenvalues to the top left corner of a Schur decomposition.

Parameters

in	*conf*	Configuration structure.
in	*n*	The order of S and Q.
in,out	*selected*	The selection array.
in,out	*S*	On entry, the Schur matrix S. On exit, the updated Schur matrix $\hat{S}$.
in	*ldS*	The leading dimension of S.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U$.
in	*ldQ*	The leading dimension of Q.
out	*real*	An array of the same size as S containing the real parts of the computed eigenvalues.
out	*imag*	An array of the same size as S containing the imaginary parts of the computed eigenvalues.

Returns

[STARNEIG_SUCCESS](#) (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

[starneig_SEP_SM_ReorderSchur](#)

[starneig_SEP_SM_Select](#)

[starneig_reorder_conf](#)

[starneig_reorder_init_conf](#)

13.4.2.10 starneig_SEP_SM_Eigenvectors_expert()

starneig_error_t starneig_SEP_SM_Eigenvectors_expert (
          struct starneig_eigenvectors_conf * *conf,*
          int *n,*
          int *selected[ ],*
          double *S[ ],*
          int *ldS,*
          double *Q[ ],*
          int *ldQ,*
          double *X[ ],*
          int *ldX*)

Computes an eigenvector for each selected eigenvalue.

Parameters

in	*conf*	Configuration structure.
in	*n*	The order of S and Q and the number of rows of X.
in	*selected*	The selection array specifying the locations of the selected eigenvalues. The number of 1's in the array is the same as the number of columns in X.
in	*S*	The Schur matrix S.
in	*ldS*	The leading dimension of S.
in	*Q*	The orthogonal matrix Q.
in	*ldQ*	The The leading dimension of Q.
out	*X*	A matrix with n rows and one column for each selected eigenvalue. The columns represent the computed eigenvectors as previously described.
in	*ldX*	The leading dimension of X.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

starneig_SEP_SM_Select

13.5 Shared Memory / Generalized EVP

Functions for solving non-symmetric generalized eigenvalue problems on shared memory systems.

Computational functions

- starneig_error_t starneig_GEP_SM_HessenbergTriangular (int n, double A[ ], int ldA, double B[ ], int ldB, double Q[ ], int ldQ, double Z[ ], int ldZ)

    Computes a Hessenberg-triangular decomposition of a general matrix pencil.
- starneig_error_t starneig_GEP_SM_Schur (int n, double H[ ], int ldH, double R[ ], int ldR, double Q[ ], int ldQ, double Z[ ], int ldZ, double real[ ], double imag[ ], double beta[ ])

    Computes a generalized Schur decomposition given a Hessenberg-triangular decomposition.
- starneig_error_t starneig_GEP_SM_ReorderSchur (int n, int selected[ ], double S[ ], int ldS, double T[ ], int ldT, double Q[ ], int ldQ, double Z[ ], int ldZ, double real[ ], double imag[ ], double beta[ ])

    Reorders selected generalized eigenvalues to the top left corner of a generalized Schur decomposition.
- starneig_error_t starneig_GEP_SM_Reduce (int n, double A[ ], int ldA, double B[ ], int ldB, double Q[ ], int ldQ, double Z[ ], int ldZ, double real[ ], double imag[ ], double beta[ ], int(∗predicate)(double real, double imag, double beta, void ∗arg), void ∗arg, int selected[ ], int ∗num_selected)

    Computes a (reordered) generalized Schur decomposition given a general matrix pencil.
- starneig_error_t starneig_GEP_SM_Eigenvectors (int n, int selected[ ], double S[ ], int ldS, double T[ ], int ldT, double Z[ ], int ldZ, double X[ ], int ldX)

    Computes a generalized eigenvector for each selected generalized eigenvalue.

Helper functions

- starneig_error_t starneig_GEP_SM_Select (int n, double S[ ], int ldS, double T[ ], int ldT, int(∗predicate)(double real, double imag, double beta, void ∗arg), void ∗arg, int selected[ ], int ∗num_selected)

    Generates a selection array for a Schur-triangular matrix pencil using a user-supplied predicate function.

Expert computational functions

- starneig_error_t starneig_GEP_SM_Schur_expert (struct starneig_schur_conf ∗conf, int n, double H[ ], int ldH, double R[ ], int ldR, double Q[ ], int ldQ, double Z[ ], int ldZ, double real[ ], double imag[ ], double beta[ ])

    Computes a generalized Schur decomposition given a Hessenberg-triangular decomposition.
- starneig_error_t starneig_GEP_SM_ReorderSchur_expert (struct starneig_reorder_conf ∗conf, int n, int selected[ ], double S[ ], int ldS, double T[ ], int ldT, double Q[ ], int ldQ, double Z[ ], int ldZ, double real[ ], double imag[ ], double beta[ ])

    Reorders selected eigenvalues to the top left corner of a generalized Schur decomposition.
- starneig_error_t starneig_GEP_SM_Eigenvectors_expert (struct starneig_eigenvectors_conf ∗conf, int n, int selected[ ], double S[ ], int ldS, double T[ ], int ldT, double Z[ ], int ldZ, double X[ ], int ldX)

    Computes a generalized eigenvector for each selected generalized eigenvalue.

13.5.1 Detailed Description

Functions for solving non-symmetric generalized eigenvalue problems on shared memory systems.

13.5.2 Function Documentation

13.5.2.1 starneig_GEP_SM_HessenbergTriangular()

starneig_error_t starneig_GEP_SM_HessenbergTriangular (
        int *n,*
        double *A[],*
        int *ldA,*
        double *B[],*
        int *ldB,*
        double *Q[],*
        int *ldQ,*
        double *Z[],*
        int *ldZ*)

Computes a Hessenberg-triangular decomposition of a general matrix pencil.

Remarks

This function is a wrapper for several LAPACK subroutines.

Parameters

in	*n*	The order of A, B, Q and Z.
in,out	*A*	On entry, the general matrix A. On exit, the upper Hessenberg matrix H.
in	*ldA*	The leading dimension of A.
in,out	*B*	On entry, the general matrix B. On exit, the upper triangular matrix R.
in	*ldB*	The leading dimension of B.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U_1$.
in	*ldQ*	The leading dimension of Q.
in,out	*Z*	On entry, the orthogonal matrix Z. On exit, the product matrix $Z * U_2$.
in	*ldZ*	The leading dimension of Z.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

Examples

gep_sm_full_chain.c.

13.5.2.2 starneig_GEP_SM_Schur()

starneig_error_t starneig_GEP_SM_Schur (
        int *n,*

```
        double H[],
        int ldH,
        double R[],
        int ldR,
        double Q[],
        int ldQ,
        double Z[],
        int ldZ,
        double real[],
        double imag[],
        double beta[] )
```

Computes a generalized Schur decomposition given a Hessenberg-triangular decomposition.

Parameters

in	n	The order of H, T, Q and Z.
in,out	H	On entry, the upper Hessenberg matrix H. On exit, the Schur matrix S.
in	ldH	The leading dimension of H.
in,out	R	On entry, the upper triangular matrix R. On exit, the upper triangular matrix T.
in	ldR	The leading dimension of R.
in,out	Q	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U_1$.
in	ldQ	The leading dimension of Q.
in,out	Z	On entry, the orthogonal matrix Z. On exit, the product matrix $Z * U_2$.
in	ldZ	The leading dimension of Z.
out	real	An array of the same size as H containing the real parts of the α values of the computed generalized eigenvalues.
out	imag	An array of the same size as H containing the imaginary parts of the α values of the computed generalized eigenvalues.
out	beta	An array of the same size as H containing the β values of computed generalized eigenvalues.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise. STARNEIG_DID_NOT_CONVERGE if the QZ algorithm failed to converge.

Examples

gep_sm_full_chain.c.

13.5.2.3 starneig_GEP_SM_ReorderSchur()

starneig_error_t starneig_GEP_SM_ReorderSchur (
        int n,
        int selected[],
        double S[],
        int ldS,
        double T[],
        int ldT,

```
        double Q[],
        int ldQ,
        double Z[],
        int ldZ,
        double real[],
        double imag[],
        double beta[] )
```

Reorders selected generalized eigenvalues to the top left corner of a generalized Schur decomposition.

Parameters

in	*n*	The order of H, T, Q and Z.
in,out	*selected*	The selection array. On entry, the initial positions of the selected generalized eigenvalues. On exit, the final positions of all correctly placed selected generalized eigenvalues. In case of failure, the number of 1's in the output may be less than the number of 1's in the input.
in,out	*S*	On entry, the Schur matrix S. On exit, the updated Schur matrix $\hat{S}$.
in	*ldS*	The leading dimension of S.
in,out	*T*	On entry, the upper triangular T. On exit, the updates upper triangular matrix $\hat{T}$.
in	*ldT*	The leading dimension of T.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U_1$.
in	*ldQ*	The leading dimension of Q.
in,out	*Z*	On entry, the orthogonal matrix Z. On exit, the product matrix $Z * U_2$.
in	*ldZ*	The leading dimension of Z.
out	*real*	An array of the same size as S containing the real parts of the α values of the computed generalized eigenvalues.
out	*imag*	An array of the same size as S containing the imaginary parts of the α values of the computed generalized eigenvalues.
out	*beta*	An array of the same size as S containing the β values of computed generalized eigenvalues.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise. STARNEIG_PARTIAL_REORDERING if the generalized Schur form is not fully reordered.

See also

starneig_GEP_SM_Select

Examples

gep_sm_full_chain.c.

13.5.2.4 starneig_GEP_SM_Reduce()

```
starneig_error_t starneig_GEP_SM_Reduce (
            int n,
            double A[],
            int ldA,
            double B[],
            int ldB,
            double Q[],
            int ldQ,
            double Z[],
            int ldZ,
            double real[],
            double imag[],
            double beta[],
            int(*)(double real, double imag, double beta, void *arg) predicate,
            void * arg,
            int selected[],
            int * num_selected )
```

Computes a (reordered) generalized Schur decomposition given a general matrix pencil.

Parameters

in	*n*	The order of A, B, Q and Z.
in,out	*A*	On entry, the general matrix A. On exit, the Schur matrix S.
in	*ldA*	The leading dimension of A.
in,out	*B*	On entry, the general matrix B. On exit, the upper triangular matrix T.
in	*ldB*	The leading dimension of B.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U_1$.
in	*ldQ*	The leading dimension of Q.
in,out	*Z*	On entry, the orthogonal matrix Z. On exit, the product matrix $Z * U_2$.
in	*ldZ*	The leading dimension of Z.
out	*real*	An array of the same size as A containing the real parts of the α values of the computed generalized eigenvalues.
out	*imag*	An array of the same size as A containing the imaginary parts of the α values of the computed generalized eigenvalues.
out	*beta*	An array of the same size as A containing the β values of computed generalized eigenvalues.
in	*predicate*	A function that takes a (complex) generalized eigenvalue as input and returns non-zero if it should be selected. For complex conjugate pairs of generalized eigenvalues, the predicate is called only for the generalized eigenvalue with positive imaginary part and the corresponding 2×2 block is either selected or deselected. The reordering step is skipped if the argument is a NULL pointer.
in	*arg*	An optional argument for the predicate function.
out	*selected*	The final positions of all correctly placed selected generalized eigenvalues.
out	*num_selected*	The number of selected generalized eigenvalues (a complex conjugate pair is counted as two selected generalized eigenvalues).

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise. STARNEIG_DID_NOT_CONVERGE if the QZ algorithm failed to converge. STARNEIG_PARTIAL_REORDERING if the generalized Schur form is not fully reordered.

13.5.2.5 starneig_GEP_SM_Eigenvectors()

starneig_error_t starneig_GEP_SM_Eigenvectors (
          int *n,*
          int *selected[],*
          double *S[],*
          int *ldS,*
          double *T[],*
          int *ldT,*
          double *Z[],*
          int *ldZ,*
          double *X[],*
          int *ldX*)

Computes a generalized eigenvector for each selected generalized eigenvalue.

Parameters

in	*n*	The order of S and Q and the number of rows of X.
in	*selected*	The selection array specifying the locations of the selected generalized eigenvalues. The number of 1's in the array is the same as the number of columns in X.
in	*S*	The Schur matrix S.
in	*ldS*	The leading dimension of S.
in	*T*	The upper triangular matrix T.
in	*ldT*	The leading dimension of T.
in	*Z*	The orthogonal matrix Z.
in	*ldZ*	The leading dimension of Z.
out	*X*	A matrix with n rows and one column for each selected generalized eigenvalue. The columns represent the computed generalized eigenvectors as previously described.
in	*ldX*	The leading dimension of X.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

starneig_GEP_SM_Select

Examples

gep_sm_eigenvectors.c.

13.5.2.6   starneig_GEP_SM_Select()

```
starneig_error_t starneig_GEP_SM_Select (
            int n,
            double S[],
            int ldS,
            double T[],
            int ldT,
            int(*)(double real, double imag, double beta, void *arg) predicate,
            void * arg,
            int selected[],
            int * num_selected )
```

Generates a selection array for a Schur-triangular matrix pencil using a user-supplied predicate function.

Parameters

in	n	The order of S and T.
in	S	The Schur matrix S.
in	ldS	The leading dimension of S.
in	T	The upper triangular matrix T.
in	ldT	The leading dimension of T.
in	predicate	A function that takes a (complex) generalized eigenvalue as input and returns non-zero if it should be selected. For complex conjugate pairs of generalized eigenvalues, the predicate is called only for the generallized eigenvalue with positive imaginary part and the corresponding 2×2 block is either selected or deselected.
in	arg	An optional argument for the predicate function.
out	selected	The selection array. Both elements of a selected complex conjugate pair are set to 1.
out	num_selected	The number of selected generalized eigenvalues (a complex conjugate pair is counted as two selected generalized eigenvalues).

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

Examples

gep_sm_eigenvectors.c, and gep_sm_full_chain.c.

13.5.2.7   starneig_GEP_SM_Schur_expert()

```
starneig_error_t starneig_GEP_SM_Schur_expert (
            struct starneig_schur_conf * conf,
            int n,
            double H[],
            int ldH,
            double R[],
            int ldR,
            double Q[],
```

```
int ldQ,
double Z[],
int ldZ,
double real[],
double imag[],
double beta[] )
```

Computes a generalized Schur decomposition given a Hessenberg-triangular decomposition.

Parameters

in	*conf*	Configuration structure.
in	*n*	The order of H, T, Q and Z.
in,out	*H*	On entry, the upper Hessenberg matrix H. On exit, the Schur matrix S.
in	*ldH*	The leading dimension of H.
in,out	*R*	On entry, the upper triangular matrix R. On exit, the upper triangular matrix T.
in	*ldR*	The leading dimension of R.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U_1$.
in	*ldQ*	The leading dimension of Q.
in,out	*Z*	On entry, the orthogonal matrix Z. On exit, the product matrix $Z * U_2$.
in	*ldZ*	The leading dimension of Z.
out	*real*	An array of the same size as H containing the real parts of the α values of the computed generalized eigenvalues.
out	*imag*	An array of the same size as H containing the imaginary parts of the α values of the computed generalized eigenvalues.
out	*beta*	An array of the same size as H containing the β values of computed generalized eigenvalues.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

starneig_GEP_SM_Schur

starneig_schur_conf

starneig_schur_init_conf

13.5.2.8 starneig_GEP_SM_ReorderSchur_expert()

```
starneig_error_t starneig_GEP_SM_ReorderSchur_expert (
        struct starneig_reorder_conf * conf,
        int n,
        int selected[],
        double S[],
        int ldS,
        double T[],
```

```
          int ldT,
          double Q[],
          int ldQ,
          double Z[],
          int ldZ,
          double real[],
          double imag[],
          double beta[] )
```

Reorders selected eigenvalues to the top left corner of a generalized Schur decomposition.

Parameters

in	*conf*	Configuration structure.
in	*n*	The order of H, T, Q and Z.
in,out	*selected*	The selection array.
in,out	*S*	On entry, the Schur matrix S. On exit, the updated Schur matrix $\hat{S}$.
in	*ldS*	The leading dimension of S.
in,out	*T*	On entry, the upper triangular T. On exit, the updates upper triangular matrix $\hat{T}$.
in	*ldT*	The leading dimension of T.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U_1$.
in	*ldQ*	The leading dimension of Q.
in,out	*Z*	On entry, the orthogonal matrix Z. On exit, the product matrix $Z * U_2$.
in	*ldZ*	The leading dimension of Z.
out	*real*	An array of the same size as S containing the real parts of the α values of the computed generalized eigenvalues.
out	*imag*	An array of the same size as S containing the imaginary parts of the α values of the computed generalized eigenvalues.
out	*beta*	An array of the same size as S containing the β values of computed generalized eigenvalues.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

starneig_GEP_SM_ReorderSchur

starneig_GEP_SM_Select

starneig_reorder_conf

starneig_reorder_init_conf

13.5.2.9 starneig_GEP_SM_Eigenvectors_expert()

starneig_error_t starneig_GEP_SM_Eigenvectors_expert (
          struct starneig_eigenvectors_conf * conf,

```
        int n,
        int selected[],
        double S[],
        int ldS,
        double T[],
        int ldT,
        double Z[],
        int ldZ,
        double X[],
        int ldX )
```

Computes a generalized eigenvector for each selected generalized eigenvalue.

Parameters

in	*conf*	Configuration structure.
in	*n*	The order of S and Q and the number of rows of X.
in	*selected*	The selection array specifying the locations of the selected generalized eigenvalues. The number of 1's in the array is the same as the number of columns in X.
in	*S*	The Schur matrix S.
in	*ldS*	The leading dimension of S.
in	*T*	The upper triangular matrix T.
in	*ldT*	The leading dimension of T.
in	*Z*	The orthogonal matrix Z.
in	*ldZ*	The leading dimension of Z.
out	*X*	A matrix with n rows and one column for each selected generalized eigenvalue. The columns represent the computed generalized eigenvectors as previously described.
in	*ldX*	The leading dimension of X.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

starneig_GEP_SM_Select

13.6 Distributed Memory / Distributed matrices

Data types and functions for distributed matrices.

Data Structures

- struct starneig_distr_block

    Distributed block. More...

Data distributions

- enum starneig_distr_order_t { STARNEIG_ORDER_DEFAULT, STARNEIG_ORDER_ROW_MAJOR, STARNEIG_ORDER_COL_MAJOR }

    Process mapping order.
- typedef struct starneig_distr ∗ starneig_distr_t

    Data distribution.
- starneig_distr_t starneig_distr_init ()

    Creates a default data distribution.
- starneig_distr_t starneig_distr_init_mesh (int rows, int cols, starneig_distr_order_t order)

    Creates a two-dimensional block cyclic data distribution.
- starneig_distr_t starneig_distr_init_func (int(∗func)(int row, int col, void ∗arg), void ∗arg, size_t arg_size)

    Creates a distribution using a data distribution function.
- starneig_distr_t starneig_distr_duplicate (starneig_distr_t distr)

    Duplicates a data distribution.
- void starneig_distr_destroy (starneig_distr_t distr)

    Destroys a data distribution.

Distributed matrices

- enum starneig_datatype_t { STARNEIG_REAL_DOUBLE }

    Distributed matrix element data type.
- typedef struct starneig_distr_matrix ∗ starneig_distr_matrix_t

    Distributed matrix.
- starneig_distr_matrix_t starneig_distr_matrix_create (int rows, int cols, int row_blksz, int col_blksz, starneig_datatype_t type, starneig_distr_t distr)

    Creates a distributed matrix with uninitialized matrix elements.
- starneig_distr_matrix_t starneig_distr_matrix_create_local (int rows, int cols, starneig_datatype_t type, int owner, double ∗A, int ldA)

    Creates a single-owner distributed matrix from a local matrix.
- void starneig_distr_matrix_destroy (starneig_distr_matrix_t matrix)

    Destroys a distributed matrix.
- void starneig_distr_matrix_copy (starneig_distr_matrix_t source, starneig_distr_matrix_t dest)

    Copies the contents of a distributed matrix to a second distributed matrix.
- void starneig_distr_matrix_copy_region (int sr, int sc, int dr, int dc, int rows, int cols, starneig_distr_matrix_t source, starneig_distr_matrix_t dest)

    Copies region of a distributed matrix to a second distributed matrix.

Query functions

- void starneig_distr_matrix_get_blocks (starneig_distr_matrix_t matrix, struct starneig_distr_block ∗∗blocks, int ∗num_blocks)

  Returns the locally owned distributed blocks.
- starneig_distr_t starneig_distr_matrix_get_distr (starneig_distr_matrix_t matrix)

  Returns the distribution that is associated with a distributed matrix.
- starneig_datatype_t starneig_distr_matrix_get_datatype (starneig_distr_matrix_t matrix)

  Returns the matrix element data type.
- size_t starneig_distr_matrix_get_elemsize (starneig_distr_matrix_t matrix)

  Returns the matrix element size.
- int starneig_distr_matrix_get_rows (starneig_distr_matrix_t matrix)

  Returns the number of (global) rows.
- int starneig_distr_matrix_get_cols (starneig_distr_matrix_t matrix)

  Returns the number of (global) columns.
- int starneig_distr_matrix_get_row_blksz (starneig_distr_matrix_t matrix)

  Returns the number of rows in a distribution block.
- int starneig_distr_matrix_get_col_blksz (starneig_distr_matrix_t matrix)

  Returns the number of columns in a distribution block.

13.6.1 Detailed Description

Data types and functions for distributed matrices.

13.6.2 Data Structure Documentation

13.6.2.1 struct starneig_distr_block

Distributed block.

Data Fields

int	row_blksz	The number of rows in the block.
int	col_blksz	The number of columns in the block.
int	glo_row	The topmost global row that belong to the block.
int	glo_col	The leftmost global column that belong to the block.
int	ld	The leading dimension of the local array.
void ∗	ptr	A pointer to the local array.

13.6.3 Enumeration Type Documentation

13.6.3.1 starneig_distr_order_t

```
enum starneig_distr_order_t
```

Process mapping order.

Enumerator

STARNEIG_ORDER_DEFAULT	Default ordering.
STARNEIG_ORDER_ROW_MAJOR	Row-major natural ordering.
STARNEIG_ORDER_COL_MAJOR	Column-major natural ordering.

13.6.3.2 starneig_datatype_t

enum starneig_datatype_t

Distributed matrix element data type.

Enumerator

STARNEIG_REAL_DOUBLE	Double precision real numbers.

13.6.4 Function Documentation

13.6.4.1 starneig_distr_init()

starneig_distr_t starneig_distr_init ( )

Creates a default data distribution.

Returns

A new data distribution.

13.6.4.2 starneig_distr_init_mesh()

starneig_distr_t starneig_distr_init_mesh (
        int *rows,*
        int *cols,*
        starneig_distr_order_t *order*)

Creates a two-dimensional block cyclic data distribution.

Parameters

in	*rows*	The number of rows in the mesh. Can be set to -1 in which case the library decides the value.
in	*cols*	The number of columns in the mesh. Can be set to -1 in which case the library decides the value.
in	*order*	The process mapping order.

Returns

A new data distribution.

Examples

[gep_dm_full_chain.c](), and [sep_dm_full_chain.c]().

13.6.4.3 starneig_distr_init_func()

```
starneig_distr_t starneig_distr_init_func (
            int(*)(int row, int col, void *arg) func,
            void * arg,
            size_t arg_size )
```

Creates a distribution using a data distribution function.

The distribution function maps each block to it's owner. The function takes three arguments: block's row index, blocks's column index and an optional user defined argument.

```
struct block_cyclic_arg {
    int rows;
    int cols;
};
int block_cyclic_func(int i, int j, void *arg)
{
    struct block_cyclic_arg *mesh = (struct block_cyclic_arg *) arg;
    return (i % mesh->rows) * mesh->cols + j % mesh->cols;
}
void func(...)
{
    ...
    // create a custom two-dimensional block cyclic distribution with 4 rows
    // and 6 columns in the mesh
    struct block_cyclic_arg arg = { .rows = 4, .cols = 6 };
    starneig_distr_t distr =
        starneig_distr_init_func(&block_cyclic_func, &arg, sizeof(arg));
    ...
}
```

Parameters

in	*func*	The data distribution function.
in	*arg*	An optional data distribution function argument.
in	*arg_size*	The size of the optional data distribution function argument.

Returns

A new data distribution.

13.6.4.4 starneig_distr_duplicate()

```
starneig_distr_t starneig_distr_duplicate (
            starneig_distr_t distr )
```

Duplicates a data distribution.

Parameters

`in`	*distr*	The data distribution to be duplicated.

Returns

A duplicated data distribution.

13.6.4.5 starneig_distr_destroy()

```
void starneig_distr_destroy (
            starneig_distr_t distr )
```

Destroys a data distribution.

Parameters

`in,out`	*distr*	The data distribution to be destroyed.

Examples

gep_dm_full_chain.c, and sep_dm_full_chain.c.

13.6.4.6 starneig_distr_matrix_create()

```
starneig_distr_matrix_t starneig_distr_matrix_create (
            int rows,
            int cols,
            int row_blksz,
            int col_blksz,
            starneig_datatype_t type,
            starneig_distr_t distr )
```

Creates a distributed matrix with uninitialized matrix elements.
```
// create a m X n double-precision real matrix that is distributed in a
// two-dimensional block cyclic fashion in bm X bn blocks
starneig_distr_t distr = starneig_distr_init();
starneig_distr_matrix_t dA =
    starneig_distr_matrix_create(m, n, bm, bn, STARNEIG_REAL_DOUBLE, distr);
```

Attention

StarNEig library is designed to use much larger distributed blocks than ScaLAPACK. Selecting a too small distributed block size will be detrimental to the performance.

Parameters

in	*rows*	The number of (global) rows in the matrix.
in	*cols*	The number of (global) columns in the matrix.
in	*row_blksz*	The number of rows in a distribution block. Can be set to -1 in which case the library decides the value.
in	*col_blksz*	The number of columns in a distribution block. Can be set to -1 in which case the library decides the value.
in	*type*	The matrix element data type.
in	*distr*	The data distribution. Can be left to NULL in which case the library decides the distribution.

Returns

A new distributed matrix.

Examples

gep_dm_full_chain.c, and sep_dm_full_chain.c.

13.6.4.7   starneig_distr_matrix_create_local()

```
starneig_distr_matrix_t starneig_distr_matrix_create_local (
            int rows,
            int cols,
            starneig_datatype_t type,
            int owner,
            double * A,
            int ldA )
```

Creates a single-owner distributed matrix from a local matrix.

This creates a wrapper. The contents of the local matrix may be modified by the functions that use the wrapper. The starneig_distr_matrix_destroy() function does not free the local matrix.

```
int m = 1000, n = 1000;
double *A = NULL; size_t ldA = 0;
// rank 3 initialized the local matrix
if (my_rank = 3) {
    A = initialize_matrix(m, n, &ldA);
}
// all ranks initialize the distributed matrix
starneig_distr_matrix_t lA = starneig_distr_matrix_create_local(
    m, n, STARNEIG_REAL_DOUBLE, 3, A, ldA);
```

Parameters

in	*rows*	The number of rows in the matrix.
in	*cols*	The number of columns in the matrix.
in	*type*	Matrix element data type.
in	*owner*	MPI rank that owns the distributed matrix.
in	*A*	A pointer to the local matrix. This argument is ignored the calling rank is not the same as the owner.
in	*ldA*	The leading dimension of the local matrix. This argument is ignored the calling rank is not the same as the owner.

Returns

> A new distributed matrix.

Examples

> [gep_dm_full_chain.c](), and [sep_dm_full_chain.c]().

13.6.4.8 starneig_distr_matrix_destroy()

```
void starneig_distr_matrix_destroy (
            starneig_distr_matrix_t matrix )
```

Destroys a distributed matrix.

Parameters

in,out	*matrix*	The distributed matrix to be destroyed.

Examples

> [gep_dm_full_chain.c](), and [sep_dm_full_chain.c]().

13.6.4.9 starneig_distr_matrix_copy()

```
void starneig_distr_matrix_copy (
            starneig_distr_matrix_t source,
            starneig_distr_matrix_t dest )
```

Copies the contents of a distributed matrix to a second distributed matrix.

Parameters

in	*source*	The source matrix.
out	*dest*	The destination matrix.

Examples

> [gep_dm_full_chain.c](), and [sep_dm_full_chain.c]().

13.6.4.10 starneig_distr_matrix_copy_region()

```
void starneig_distr_matrix_copy_region (
            int sr,
```

```
            int sc,
            int dr,
            int dc,
            int rows,
            int cols,
            starneig_distr_matrix_t source,
            starneig_distr_matrix_t dest )
```

Copies region of a distributed matrix to a second distributed matrix.

Parameters

in	sr	The first source matrix row to be copied.
in	sc	The first source matrix column to be copied.
in	dr	The first destination matrix row.
in	dc	The first destination matrix column.
in	rows	The number of rows to copy.
in	cols	The number of columns to copy.
in	source	The source matrix.
out	dest	The destination matrix.

13.6.4.11   starneig_distr_matrix_get_blocks()

```
void starneig_distr_matrix_get_blocks (
            starneig_distr_matrix_t matrix,
            struct starneig_distr_block ** blocks,
            int * num_blocks )
```

Returns the locally owned distributed blocks.

Attention

> A user is allowed to modify the contents of the locally owned blocks but the the returned array itself should not be modified.

Parameters

in	matrix	The distributed matrix.
out	blocks	An array that contains all locally owned distributed blocks.
out	num_blocks	The total number of locally owned distributed blocks.

13.6.4.12   starneig_distr_matrix_get_distr()

```
starneig_distr_t starneig_distr_matrix_get_distr (
            starneig_distr_matrix_t matrix )
```

Returns the distribution that is associated with a distributed matrix.

Attention

The distributed matrix maintains the ownership of the returned data distribution. A user must duplicate the data distribution if necessary.

Parameters

in	*matrix*	The distributed matrix.

Returns

The associated distribution.

13.6.4.13 starneig_distr_matrix_get_datatype()

starneig_datatype_t starneig_distr_matrix_get_datatype (
            starneig_distr_matrix_t *matrix*)

Returns the matrix element data type.

Parameters

in	*matrix*	The distributed matrix.

Returns

The matrix element data type.

13.6.4.14 starneig_distr_matrix_get_elemsize()

size_t starneig_distr_matrix_get_elemsize (
            starneig_distr_matrix_t *matrix*)

Returns the matrix element size.

Parameters

in	*matrix*	The distributed matrix.

Returns

The matrix element size.

13.6.4.15 starneig_distr_matrix_get_rows()

```
int starneig_distr_matrix_get_rows (
            starneig_distr_matrix_t matrix )
```

Returns the number of (global) rows.

Parameters

in	*matrix*	The distributed matrix.

Returns

The number of (global) rows.

13.6.4.16 starneig_distr_matrix_get_cols()

```
int starneig_distr_matrix_get_cols (
            starneig_distr_matrix_t matrix )
```

Returns the number of (global) columns.

Parameters

in	*matrix*	The distributed matrix.

Returns

The number of (global) columns.

13.6.4.17 starneig_distr_matrix_get_row_blksz()

```
int starneig_distr_matrix_get_row_blksz (
            starneig_distr_matrix_t matrix )
```

Returns the number of rows in a distribution block.

Parameters

in	*matrix*	The distributed matrix.

Returns

The number of rows in a distribution block.

13.6.4.18 starneig_distr_matrix_get_col_blksz()

```
int starneig_distr_matrix_get_col_blksz (
            starneig_distr_matrix_t matrix )
```

Returns the number of columns in a distribution block.

Parameters

in	*matrix*	The distributed matrix.

Returns

The number of columns in a distribution block.

13.7 Distributed Memory / Helper functions

Distributed memory helper functions.

MPI communicator

- void starneig_mpi_set_comm (MPI_Comm comm)

    Sets a MPI communicator for the library.
- MPI_Comm starneig_mpi_get_comm ()

    Returns the library MPI communicator.

Broadcast

- void starneig_mpi_broadcast (int root, size_t size, void ∗buffer)

    Broadcast a buffer.
- void starneig_broadcast (int root, size_t size, void ∗buffer)

    Broadcast a buffer. Deprecated.

13.7.1 Detailed Description

Distributed memory helper functions.

13.7.2 Function Documentation

13.7.2.1 starneig_mpi_set_comm()

```
void starneig_mpi_set_comm (
            MPI_Comm comm )
```

Sets a MPI communicator for the library.

Should be called before the starneig_node_init() interface function.

Parameters

in	*comm*	The library MPI communicator.

13.7.2.2 starneig_mpi_get_comm()

```
MPI_Comm starneig_mpi_get_comm ( )
```

Returns the library MPI communicator.

Returns

The library MPI communicator.

13.7.2.3 starneig_mpi_broadcast()

```
void starneig_mpi_broadcast (
            int root,
            size_t size,
            void * buffer )
```

Broadcast a buffer.

Parameters

in	*root*	The rank that is going to broadcast the buffer.
in	*size*	The size of the buffer.
in,out	*buffer*	A pointer to the buffer.

13.7.2.4 starneig_broadcast()

```
void starneig_broadcast (
            int root,
            size_t size,
            void * buffer )
```

Broadcast a buffer. Deprecated.

Deprecated The starneig_broadcast() function has been replaced with the starneig_mpi_broadcast() function. This function will be removed in a future release of the library.

13.8 Distributed Memory / Standard EVP

Functions for solving non-symmetric standard eigenvalue problems on distributed memory systems.

Computational functions

- starneig_error_t starneig_SEP_DM_Hessenberg (starneig_distr_matrix_t A, starneig_distr_matrix_t Q)

  Computes a Hessenberg decomposition of a general matrix.
- starneig_error_t starneig_SEP_DM_Schur (starneig_distr_matrix_t H, starneig_distr_matrix_t Q, double real[ ], double imag[ ])

  Computes a Schur decomposition given a Hessenberg decomposition.
- starneig_error_t starneig_SEP_DM_ReorderSchur (int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t Q, double real[ ], double imag[ ])

  Reorders selected eigenvalues to the top left corner of a Schur decomposition.
- starneig_error_t starneig_SEP_DM_Reduce (starneig_distr_matrix_t A, starneig_distr_matrix_t Q, double real[ ], double imag[ ], int(∗predicate)(double real, double imag, void ∗arg), void ∗arg, int selected[ ], int ∗num_selected)

  Computes a (reordered) Schur decomposition of a general matrix.
- starneig_error_t starneig_SEP_DM_Eigenvectors (int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t Q, starneig_distr_matrix_t X)

  Computes an eigenvector for each selected eigenvalue.

Helper functions

- starneig_error_t starneig_SEP_DM_Select (starneig_distr_matrix_t S, int(∗predicate)(double real, double imag, void ∗arg), void ∗arg, int selected[ ], int ∗num_selected)

  Generates a selection array for a Schur matrix using a user-supplied predicate function.

Expert computational functions

- starneig_error_t starneig_SEP_DM_Schur_expert (struct starneig_schur_conf ∗conf, starneig_distr_matrix_t H, starneig_distr_matrix_t Q, double real[ ], double imag[ ])

  Computes a Schur decomposition given a Hessenberg decomposition.
- starneig_error_t starneig_SEP_DM_ReorderSchur_expert (struct starneig_reorder_conf ∗conf, int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t Q, double real[ ], double imag[ ])

  Reorders selected eigenvalues to the top left corner of a Schur decomposition.
- starneig_error_t starneig_SEP_DM_Eigenvectors_expert (struct starneig_eigenvectors_conf ∗conf, int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t Q, starneig_distr_matrix_t X)

  Computes an eigenvector for each selected eigenvalue.

13.8.1 Detailed Description

Functions for solving non-symmetric standard eigenvalue problems on distributed memory systems.

13.8.2 Function Documentation

13.8.2.1 starneig_SEP_DM_Hessenberg()

starneig_error_t starneig_SEP_DM_Hessenberg (
   starneig_distr_matrix_t *A,*
   starneig_distr_matrix_t *Q*)

Computes a Hessenberg decomposition of a general matrix.

Attention

> This function is a wrapper for several ScaLAPACK subroutines. The function exists if STARNEIG_SEP_DM_HESSENBERG is defined.

Parameters

in,out	*A*	On entry, the general matrix A. On exit, the upper Hessenberg matrix H.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U$.

Returns

> STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

Examples

> sep_dm_full_chain.c.

13.8.2.2 starneig_SEP_DM_Schur()

starneig_error_t starneig_SEP_DM_Schur (
   starneig_distr_matrix_t *H,*
   starneig_distr_matrix_t *Q,*
   double *real[ ],*
   double *imag[ ]*)

Computes a Schur decomposition given a Hessenberg decomposition.

Parameters

in,out	*H*	On entry, the upper Hessenberg matrix H. On exit, the Schur matrix S.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U$.
out	*real*	An array of the same size as H containing the real parts of the computed eigenvalues.
out	*imag*	An array of the same size as H containing the imaginary parts of the computed eigenvalues.

Returns

> STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise. STARNEIG_DID_NOT_CONVERGE if the QR algorithm failed to converge.

Examples

sep_dm_full_chain.c.

13.8.2.3 starneig_SEP_DM_ReorderSchur()

starneig_error_t starneig_SEP_DM_ReorderSchur (
        int *selected[],*
        starneig_distr_matrix_t *S,*
        starneig_distr_matrix_t *Q,*
        double *real[],*
        double *imag[]*)

Reorders selected eigenvalues to the top left corner of a Schur decomposition.

Parameters

in,out	*selected*	The selection array. On entry, the initial positions of the selected eigenvalues. On exit, the final positions of all correctly placed selected eigenvalues. In case of failure, the number of 1's in the output may be less than the number of 1's in the input.
in,out	*S*	On entry, the Schur matrix S. On exit, the updated Schur matrix $\hat{S}$.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U$.
out	*real*	An array of the same size as S containing the real parts of the computed eigenvalues.
out	*imag*	An array of the same size as S containing the imaginary parts of the computed eigenvalues.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise. STARNEIG_PARTIAL_REORDERING if the Schur form is not fully reordered.

See also

starneig_SEP_DM_Select

Examples

sep_dm_full_chain.c.

13.8.2.4 starneig_SEP_DM_Reduce()

starneig_error_t starneig_SEP_DM_Reduce (
        starneig_distr_matrix_t *A,*
        starneig_distr_matrix_t *Q,*
        double *real[],*
        double *imag[],*
        int(*)(double real, double imag, void *arg) *predicate,*
        void * *arg,*
        int *selected[],*
        int * *num_selected*)

Computes a (reordered) Schur decomposition of a general matrix.

Attention

> This function uses several ScaLAPACK subroutines. The function exists if STARNEIG_SEP_DM_REDUCE is defined.

Parameters

in,out	*A*	On entry, the general matrix A. On exit, the Schur matrix S.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U$.
out	*real*	An array of the same size as A containing the real parts of the computed eigenvalues.
out	*imag*	An array of the same size as A containing the imaginary parts of the computed eigenvalues.
in	*predicate*	A function that takes a (complex) eigenvalue as input and returns non-zero if it should be selected. For complex conjugate pairs of eigenvalues, the predicate is called only for the eigenvalue with positive imaginary part and the corresponding 2×2 block is either selected or deselected. The reordering step is skipped if the argument is a NULL pointer.
in	*arg*	An optional argument for the predicate function.
out	*selected*	The final positions of all correctly placed selected eigenvalues.
out	*num_selected*	The number of selected eigenvalues (a complex conjugate pair is counted as two selected eigenvalues).

Returns

> STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise. STARNEIG_DID_NOT_CONVERGE if the QR algorithm failed to converge. STARNEIG_PARTIAL_REORDERING if the Schur form is not fully reordered.

13.8.2.5 starneig_SEP_DM_Eigenvectors()

```
starneig_error_t starneig_SEP_DM_Eigenvectors (
            int selected[ ],
            starneig_distr_matrix_t S,
            starneig_distr_matrix_t Q,
            starneig_distr_matrix_t X )
```

Computes an eigenvector for each selected eigenvalue.

Parameters

in	*selected*	The selection array specifying the locations of the selected eigenvalues. The number of 1's in the array is the same as the number of columns in X.
in	*S*	The Schur matrix S.
in	*Q*	The orthogonal matrix Q.
out	*X*	A matrix with n rows and one column for each selected eigenvalue. The columns represent the computed eigenvectors as previously described.

Returns

[STARNEIG_SUCCESS](0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

[starneig_SEP_DM_Select](#)

Todo This interface function is not implemented.

13.8.2.6 starneig_SEP_DM_Select()

```
starneig_error_t starneig_SEP_DM_Select (
        starneig_distr_matrix_t S,
        int(*)(double real, double imag, void *arg) predicate,
        void * arg,
        int selected[],
        int * num_selected )
```

Generates a selection array for a Schur matrix using a user-supplied predicate function.

Parameters

in	S	The Schur matrix S.
in	*predicate*	A function that takes a (complex) eigenvalue as input and returns non-zero if it should be selected. For complex conjugate pairs of eigenvalues, the predicate is called only for the eigenvalue with positive imaginary part and the corresponding 2×2 block is either selected or deselected.
in	*arg*	An optional argument for the predicate function.
out	*selected*	The selection array. Both elements of a selected complex conjugate pair are set to 1.
out	*num_selected*	The (global) number of selected eigenvalues (a complex conjugate pair is counted as two selected eigenvalues).

Returns

[STARNEIG_SUCCESS](0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

Examples

[sep_dm_full_chain.c](#).

13.8.2.7 starneig_SEP_DM_Schur_expert()

```
starneig_error_t starneig_SEP_DM_Schur_expert (
        struct starneig_schur_conf * conf,
```

```
        starneig_distr_matrix_t H,
        starneig_distr_matrix_t Q,
        double real[],
        double imag[] )
```

Computes a Schur decomposition given a Hessenberg decomposition.

Parameters

in	*conf*	Configuration structure.
in,out	*H*	On entry, the upper Hessenberg matrix H. On exit, the Schur matrix S.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U$.
out	*real*	An array of the same size as H containing the real parts of the computed eigenvalues.
out	*imag*	An array of the same size as H containing the imaginary parts of the computed eigenvalues.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

starneig_SEP_DM_Schur

starneig_schur_conf

starneig_schur_init_conf

13.8.2.8 starneig_SEP_DM_ReorderSchur_expert()

starneig_error_t starneig_SEP_DM_ReorderSchur_expert (
        struct starneig_reorder_conf * conf,
        int selected[],
        starneig_distr_matrix_t S,
        starneig_distr_matrix_t Q,
        double real[],
        double imag[])

Reorders selected eigenvalues to the top left corner of a Schur decomposition.

Parameters

in	*conf*	Configuration structure.
in,out	*selected*	The selection array.
in,out	*S*	On entry, the Schur matrix S. On exit, the updated Schur matrix $\hat{S}$.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U$.
out	*real*	An array of the same size as S containing the real parts of the computed eigenvalues.
out	*imag*	An array of the same size as S containing the imaginary parts of the computed eigenvalues.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

starneig_SEP_DM_ReorderSchur

starneig_SEP_DM_Select

starneig_reorder_conf

starneig_reorder_init_conf

13.8.2.9 starneig_SEP_DM_Eigenvectors_expert()

```
starneig_error_t starneig_SEP_DM_Eigenvectors_expert (
            struct starneig_eigenvectors_conf * conf,
            int selected[],
            starneig_distr_matrix_t S,
            starneig_distr_matrix_t Q,
            starneig_distr_matrix_t X )
```

Computes an eigenvector for each selected eigenvalue.

Parameters

in	*conf*	Configuration structure.
in	*selected*	The selection array specifying the locations of the selected eigenvalues. The number of 1's in the array is the same as the number of columns in X.
in	*S*	The Schur matrix S.
in	*Q*	The orthogonal matrix Q.
out	*X*	A matrix with n rows and one column for each selected eigenvalue. The columns represent the computed eigenvectors as previously described.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

starneig_SEP_DM_Select

Todo This interface function is not implemented.

13.9 Distributed Memory / Generalized EVP

Functions for solving non-symmetric generalized eigenvalue problems on distributed memory systems.

Computational functions

- starneig_error_t starneig_GEP_DM_HessenbergTriangular (starneig_distr_matrix_t A, starneig_distr_matrix_t B, starneig_distr_matrix_t Q, starneig_distr_matrix_t Z)

  Computes a Hessenberg-triangular decomposition of a general matrix pencil.
- starneig_error_t starneig_GEP_DM_Schur (starneig_distr_matrix_t H, starneig_distr_matrix_t R, starneig_distr_matrix_t Q, starneig_distr_matrix_t Z, double real[ ], double imag[ ], double beta[ ])

  Computes a generalized Schur decomposition given a Hessenberg-triangular decomposition.
- starneig_error_t starneig_GEP_DM_ReorderSchur (int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t T, starneig_distr_matrix_t Q, starneig_distr_matrix_t Z, double real[ ], double imag[ ], double beta[ ])

  Reorders selected generalized eigenvalues to the top left corner of a generalized Schur decomposition.
- starneig_error_t starneig_GEP_DM_Reduce (starneig_distr_matrix_t A, starneig_distr_matrix_t B, starneig_distr_matrix_t Q, starneig_distr_matrix_t Z, double real[ ], double imag[ ], double beta[ ], int(∗predicate)(double real, double imag, double beta, void ∗arg), void ∗arg, int selected[ ], int ∗num_↩ selected)

  Computes a (reordered) generalized Schur decomposition given a general matrix pencil.
- starneig_error_t starneig_GEP_DM_Eigenvectors (int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t T, starneig_distr_matrix_t Z, starneig_distr_matrix_t X)

  Computes a generalized eigenvector for each selected generalized eigenvalue.

Helper functions

- starneig_error_t starneig_GEP_DM_Select (starneig_distr_matrix_t S, starneig_distr_matrix_t T, int(∗predicate)(double real, double imag, double beta, void ∗arg), void ∗arg, int selected[ ], int ∗num_selected)

  Generates a selection array for a Schur-triangular matrix pencil using a user-supplied predicate function.

Expert computational functions

- starneig_error_t starneig_GEP_DM_Schur_expert (struct starneig_schur_conf ∗conf, starneig_distr_matrix_t H, starneig_distr_matrix_t R, starneig_distr_matrix_t Q, starneig_distr_matrix_t Z, double real[ ], double imag[ ], double beta[ ])

  Computes a generalized Schur decomposition given a Hessenberg-triangular decomposition.
- starneig_error_t starneig_GEP_DM_ReorderSchur_expert (struct starneig_reorder_conf ∗conf, int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t T, starneig_distr_matrix_t Q, starneig_distr_matrix_t Z, double real[ ], double imag[ ], double beta[ ])

  Reorders selected generalized eigenvalues to the top left corner of a generalized Schur decomposition.
- starneig_error_t starneig_GEP_DM_Eigenvectors_expert (struct starneig_eigenvectors_conf ∗conf, int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t T, starneig_distr_matrix_t Z, starneig_distr_matrix_t X)

  Computes a generalized eigenvector for each selected generalized eigenvalue.

13.9.1 Detailed Description

Functions for solving non-symmetric generalized eigenvalue problems on distributed memory systems.

13.9.2 Function Documentation

13.9.2.1 starneig_GEP_DM_HessenbergTriangular()

starneig_error_t starneig_GEP_DM_HessenbergTriangular (
           starneig_distr_matrix_t *A,*
           starneig_distr_matrix_t *B,*
           starneig_distr_matrix_t *Q,*
           starneig_distr_matrix_t *Z*)

Computes a Hessenberg-triangular decomposition of a general matrix pencil.

Attention

> This function is a wrapper for several ScaLAPACK subroutines. The function exists if STARNEIG_GEP_DM_HESSENBERGTRI
> is defined.

Parameters

in,out	*A*	On entry, the general matrix A. On exit, the upper Hessenberg matrix H.
in,out	*B*	On entry, the general matrix B. On exit, the upper triangular matrix R.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U_1$.
in,out	*Z*	On entry, the orthogonal matrix Z. On exit, the product matrix $Z * U_2$.

Returns

> STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

Examples

> gep_dm_full_chain.c.

13.9.2.2 starneig_GEP_DM_Schur()

starneig_error_t starneig_GEP_DM_Schur (
           starneig_distr_matrix_t *H,*
           starneig_distr_matrix_t *R,*
           starneig_distr_matrix_t *Q,*
           starneig_distr_matrix_t *Z,*
           double *real[ ],*
           double *imag[ ],*
           double *beta[ ]*)

Computes a generalized Schur decomposition given a Hessenberg-triangular decomposition.

Parameters

in,out	H	On entry, the upper Hessenberg matrix H. On exit, the Schur matrix S.
in,out	R	On entry, the upper triangular matrix R. On exit, the upper triangular matrix T.
in,out	Q	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U_1$.
in,out	Z	On entry, the orthogonal matrix Z. On exit, the product matrix $Z * U_2$.
out	*real*	An array of the same size as H containing the real parts of the α values of the computed generalized eigenvalues.
out	*imag*	An array of the same size as H containing the imaginary parts of the α values of the computed generalized eigenvalues.
out	*beta*	An array of the same size as H containing the β values of computed generalized eigenvalues.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise. STARNEIG_DID_NOT_CONVERGE if the QZ algorithm failed to converge.

Examples

gep_dm_full_chain.c.

13.9.2.3 starneig_GEP_DM_ReorderSchur()

```
starneig_error_t starneig_GEP_DM_ReorderSchur (
            int selected[],
            starneig_distr_matrix_t S,
            starneig_distr_matrix_t T,
            starneig_distr_matrix_t Q,
            starneig_distr_matrix_t Z,
            double real[],
            double imag[],
            double beta[] )
```

Reorders selected generalized eigenvalues to the top left corner of a generalized Schur decomposition.

Parameters

in,out	*selected*	The selection array. On entry, the initial positions of the selected generalized eigenvalues. On exit, the final positions of all correctly placed selected generalized eigenvalues. In case of failure, the number of 1's in the output may be less than the number of 1's in the input.
in,out	S	On entry, the Schur matrix S. On exit, the updated Schur matrix $\hat{S}$.
in,out	T	On entry, the upper triangular T. On exit, the updates upper triangular matrix $\hat{T}$.
in,out	Q	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U_1$.
in,out	Z	On entry, the orthogonal matrix Z. On exit, the product matrix $Z * U_2$.
out	*real*	An array of the same size as S containing the real parts of the α values of the computed generalized eigenvalues.
out	*imag*	An array of the same size as S containing the imaginary parts of the α values of the computed generalized eigenvalues.
out	*beta*	An array of the same size as S containing the β values of computed generalized eigenvalues.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise. STARNEIG_PARTIAL_REORDERING if the generalized Schur form is not fully reordered.

See also

starneig_GEP_DM_Select

Examples

gep_dm_full_chain.c.

13.9.2.4   starneig_GEP_DM_Reduce()

```
starneig_error_t starneig_GEP_DM_Reduce (
            starneig_distr_matrix_t A,
            starneig_distr_matrix_t B,
            starneig_distr_matrix_t Q,
            starneig_distr_matrix_t Z,
            double real[ ],
            double imag[ ],
            double beta[ ],
            int(*)(double real, double imag, double beta, void *arg) predicate,
            void * arg,
            int selected[ ],
            int * num_selected )
```

Computes a (reordered) generalized Schur decomposition given a general matrix pencil.

Attention

This function uses several ScaLAPACK subroutines. The function exists if STARNEIG_GEP_DM_REDUCE is defined.

Parameters

in,out	*A*	On entry, the general matrix A. On exit, the Schur matrix S.
in,out	*B*	On entry, the general matrix B. On exit, the upper triangular matrix T.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U_1$.
in,out	*Z*	On entry, the orthogonal matrix Z. On exit, the product matrix $Z * U_2$.
out	*real*	An array of the same size as A containing the real parts of the α values of the computed generalized eigenvalues.
out	*imag*	An array of the same size as A containing the imaginary parts of the α values of the computed generalized eigenvalues.
out	*beta*	An array of the same size as A containing the β values of computed generalized eigenvalues.
in	*predicate*	A function that takes a (complex) generalized eigenvalue as input and returns non-zero if it should be selected. For complex conjugate pairs of generalized eigenvalues, the predicate is called only for the generalized eigenvalue with positive imaginary part and the corresponding 2×2 block is either selected or deselected. The reordering step is skipped if the argument is a NULL pointer.
in	*arg*	Optional argument to the predicate function.
out	*selected*	The final positions of all correctly placed selected generalized eigenvalues.
out	*num_selected*	The number of selected generalized eigenvalues (a complex conjugate pair is counted as two selected generalized eigenvalues).

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise. STARNEIG_DID_NOT_CONVERGE if the QZ algorithm failed to converge. STARNEIG_PARTIAL_REORDERING if the generalized Schur form is not fully reordered.

13.9.2.5 starneig_GEP_DM_Eigenvectors()

```
starneig_error_t starneig_GEP_DM_Eigenvectors (
            int selected[],
            starneig_distr_matrix_t S,
            starneig_distr_matrix_t T,
            starneig_distr_matrix_t Z,
            starneig_distr_matrix_t X )
```

Computes a generalized eigenvector for each selected generalized eigenvalue.

Parameters

in	*selected*	The selection array specifying the locations of the selected generalized eigenvalues. The number of 1's in the array is the same as the number of columns in X.
in	*S*	The Schur matrix S.
in	*T*	The upper triangular matrix T.
in	*Z*	The orthogonal matrix Z.
out	*X*	A matrix with n rows and one column for each selected generalized eigenvalue. The columns represent the computed generalized eigenvectors as previously described.

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

starneig_GEP_DM_Select

Todo This interface function is not implemented.

13.9.2.6 starneig_GEP_DM_Select()

```
starneig_error_t starneig_GEP_DM_Select (
            starneig_distr_matrix_t S,
            starneig_distr_matrix_t T,
            int(*)(double real, double imag, double beta, void *arg) predicate,
            void * arg,
            int selected[],
            int * num_selected )
```

Generates a selection array for a Schur-triangular matrix pencil using a user-supplied predicate function.

Parameters

in	*S*	The Schur matrix S.
in	*T*	The upper triangular matrix T.
in	*predicate*	A function that takes a (complex) generalized eigenvalue as input and returns non-zero if it should be selected. For complex conjugate pairs of generalized eigenvalues, the predicate is called only for the generallized eigenvalue with positive imaginary part and the corresponding 2×2 block is either selected or deselected.
in	*arg*	An optional argument for the predicate function.
out	*selected*	The selection array. Both elements of a selected complex conjugate pair are set to 1.
out	*num_selected*	The number of selected generalized eigenvalues (a complex conjugate pair is counted as two selected generalized eigenvalues).

Returns

STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

Examples

gep_dm_full_chain.c.

13.9.2.7 starneig_GEP_DM_Schur_expert()

```
starneig_error_t starneig_GEP_DM_Schur_expert (
            struct starneig_schur_conf * conf,
            starneig_distr_matrix_t H,
            starneig_distr_matrix_t R,
            starneig_distr_matrix_t Q,
            starneig_distr_matrix_t Z,
            double real[],
            double imag[],
            double beta[] )
```

Computes a generalized Schur decomposition given a Hessenberg-triangular decomposition.

Parameters

in	*conf*	Configuration structure.
in,out	*H*	On entry, the upper Hessenberg matrix H. On exit, the Schur matrix S.
in,out	*R*	On entry, the upper triangular matrix R. On exit, the upper triangular matrix T.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U_1$.
in,out	*Z*	On entry, the orthogonal matrix Z. On exit, the product matrix $Z * U_2$.
out	*real*	An array of the same size as H containing the real parts of the α values of the computed generalized eigenvalues.
out	*imag*	An array of the same size as H containing the imaginary parts of the α values of the computed generalized eigenvalues.
out	*beta*	An array of the same size as H containing the β values of computed generalized eigenvalues.

Returns

[STARNEIG_SUCCESS](0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

[starneig_GEP_DM_Schur](#)

[starneig_schur_conf](#)

[starneig_schur_init_conf](#)

13.9.2.8 starneig_GEP_DM_ReorderSchur_expert()

```
starneig_error_t starneig_GEP_DM_ReorderSchur_expert (
            struct starneig_reorder_conf * conf,
            int selected[],
            starneig_distr_matrix_t S,
            starneig_distr_matrix_t T,
            starneig_distr_matrix_t Q,
            starneig_distr_matrix_t Z,
            double real[],
            double imag[],
            double beta[] )
```

Reorders selected generalized eigenvalues to the top left corner of a generalized Schur decomposition.

Parameters

in	*conf*	Configuration structure.
in,out	*selected*	The selection array.
in,out	*S*	On entry, the Schur matrix S. On exit, the updated Schur matrix $\hat{S}$.
in,out	*T*	On entry, the upper triangular T. On exit, the updates upper triangular matrix $\hat{T}$.
in,out	*Q*	On entry, the orthogonal matrix Q. On exit, the product matrix $Q * U_1$.
in,out	*Z*	On entry, the orthogonal matrix Z. On exit, the product matrix $Z * U_2$.
out	*real*	An array of the same size as S containing the real parts of the α values of the computed generalized eigenvalues.
out	*imag*	An array of the same size as S containing the imaginary parts of the α values of the computed generalized eigenvalues.
out	*beta*	An array of the same size as S containing the β values of computed generalized eigenvalues.

Returns

[STARNEIG_SUCCESS](0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

    starneig_GEP_DM_ReorderSchur

    starneig_GEP_DM_Select

    starneig_reorder_conf

    starneig_reorder_init_conf

13.9.2.9 starneig_GEP_DM_Eigenvectors_expert()

starneig_error_t starneig_GEP_DM_Eigenvectors_expert (
        struct starneig_eigenvectors_conf * conf,
        int selected[ ],
        starneig_distr_matrix_t S,
        starneig_distr_matrix_t T,
        starneig_distr_matrix_t Z,
        starneig_distr_matrix_t X)

Computes a generalized eigenvector for each selected generalized eigenvalue.

Parameters

in	*conf*	Configuration structure.
in	*selected*	The selection array specifying the locations of the selected generalized eigenvalues. The number of 1's in the array is the same as the number of columns in X.
in	*S*	The Schur matrix S.
in	*T*	The upper triangular matrix T.
in	*Z*	The orthogonal matrix Z.
out	*X*	A matrix with n rows and one column for each selected generalized eigenvalue. The columns represent the computed generalized eigenvectors as previously described.

Returns

    STARNEIG_SUCCESS (0) on success. Negative integer -i when i'th argument is invalid. Positive error code otherwise.

See also

    starneig_GEP_DM_Select

Todo This interface function is not implemented.

13.10 Expert configuration structures

Configuration structures and functions for the expert interface functions.

Data Structures

- struct starneig_hessenberg_conf

  Hessenberg reduction configuration structure. More...
- struct starneig_schur_conf

  Schur reduction configuration structure. More...
- struct starneig_reorder_conf

  Eigenvalue reordering configuration structure. More...
- struct starneig_eigenvectors_conf

  Eigenvector computation configuration structure. More...

Hessenberg reduction

- void starneig_hessenberg_init_conf (struct starneig_hessenberg_conf ∗conf)

  Initializes a Hessenberg reduction configuration structure with default parameters.
- #define STARNEIG_HESSENBERG_DEFAULT_TILE_SIZE -1

  Default tile size.
- #define STARNEIG_HESSENBERG_DEFAULT_PANEL_WIDTH -1

  Default panel width.
- #define STARNEIG_HESSENBERG_DEFAULT_PARALLEL_WORKER_SIZE -1

  Default parallel worker size.

Schur reduction

- void starneig_schur_init_conf (struct starneig_schur_conf ∗conf)

  Initializes a Schur reduction configuration structure with default parameters.
- #define STARNEIG_SCHUR_DEFAULT_INTERATION_LIMIT -1

  Default iteration limit.
- #define STARNEIG_SCHUR_DEFAULT_TILE_SIZE -1

  Default tile size.
- #define STARNEIG_SCHUR_DEFAULT_SMALL_LIMIT -1

  Default sequential QR limit.
- #define STARNEIG_SCHUR_DEFAULT_AED_WINDOW_SIZE -1

  Default AED window size.
- #define STARNEIG_SCHUR_DEFAULT_AED_NIBBLE -1

  Default nibble value.
- #define STARNEIG_SCHUR_DEFAULT_AED_PARALLEL_SOFT_LIMIT -1

  Default soft sequential AED limit.
- #define STARNEIG_SCHUR_DEFAULT_AED_PARALLEL_HARD_LIMIT -1

  Default hard sequential AED limit.
- #define STARNEIG_SCHUR_DEFAULT_SHIFT_COUNT -1

  Default shift count.
- #define STARNEIG_SCHUR_DEFAULT_WINDOW_SIZE -1

  Default bulge chasing window size.

- #define STARNEIG_SCHUR_ROUNDED_WINDOW_SIZE -2

  Rounded bulge chasing window.
- #define STARNEIG_SCHUR_DEFAULT_SHIFTS_PER_WINDOW -1

  Default number of shifts per bulge chasing window.
- #define STARNEIG_SCHUR_DEFAULT_UPDATE_WIDTH -1

  Default left-hand side update width.
- #define STARNEIG_SCHUR_DEFAULT_UPDATE_HEIGHT -1

  Default right-hand side update height.
- #define STARNEIG_SCHUR_DEFAULT_THRESHOLD -1

  Default deflation threshold.
- #define STARNEIG_SCHUR_NORM_STABLE_THRESHOLD -2

  Norm stable deflation threshold.
- #define STARNEIG_SCHUR_LAPACK_THRESHOLD -3

  LAPACK-style deflation threshold.

Eigenvalue reordering

- enum starneig_reorder_plan_t { STARNEIG_REORDER_DEFAULT_PLAN = 1, STARNEIG_REORDER_ONE_PART_PLAN = 2, STARNEIG_REORDER_MULTI_PART_PLAN = 3 }

  Reordering plan enumerator.
- enum starneig_reorder_blueprint_t {
  STARNEIG_REORDER_DEFAULT_BLUEPRINT = 1, STARNEIG_REORDER_DUMMY_INSERT_A = 2,
  STARNEIG_REORDER_DUMMY_INSERT_B = 3, STARNEIG_REORDER_CHAIN_INSERT_A = 4,
  STARNEIG_REORDER_CHAIN_INSERT_B = 5, STARNEIG_REORDER_CHAIN_INSERT_C = 6,
  STARNEIG_REORDER_CHAIN_INSERT_D = 7, STARNEIG_REORDER_CHAIN_INSERT_E = 8,
  STARNEIG_REORDER_CHAIN_INSERT_F = 9 }

  Task insertion blueprint.
- void starneig_reorder_init_conf (struct starneig_reorder_conf ∗conf)

  Initializes an eigenvalue reordering configuration structure with default parameters.
- #define STARNEIG_REORDER_DEFAULT_UPDATE_WIDTH -1

  Default left-hand side update task width.
- #define STARNEIG_REORDER_DEFAULT_UPDATE_HEIGHT -1

  Default right-hand side update task height.
- #define STARNEIG_REORDER_DEFAULT_TILE_SIZE -1

  Default tile size.
- #define STARNEIG_REORDER_DEFAULT_VALUES_PER_CHAIN -1

  Default number of selected eigenvalues per window.
- #define STARNEIG_REORDER_DEFAULT_WINDOW_SIZE -1

  Default default window size.
- #define STARNEIG_REORDER_ROUNDED_WINDOW_SIZE -2

  Default rounded window size.
- #define STARNEIG_REORDER_DEFAULT_SMALL_WINDOW_SIZE -1

  Default small window size.
- #define STARNEIG_REORDER_DEFAULT_SMALL_WINDOW_THRESHOLD -1

  Default small window threshold.

Eigenvectors

- void starneig_eigenvectors_init_conf (struct starneig_eigenvectors_conf ∗conf)

  Initializes an eigenvectors configuration structure with default parameters.
- #define STARNEIG_EIGENVECTORS_DEFAULT_TILE_SIZE -1

  Default tile size.

13.10.1 Detailed Description

Configuration structures and functions for the expert interface functions.

13.10.2 Data Structure Documentation

13.10.2.1 struct starneig_hessenberg_conf

Hessenberg reduction configuration structure.

Data Fields

int	tile_size	The matrices are divided into square tiles. This parameter defines the used tile size. If the parameter is set to STARNEIG_HESSENBERG_DEFAULT_TILE_SIZE, then the implementation will determine a suitable tile size automatically.
int	panel_width	The reduction is performed one panel at a time. This parameter defines the used panel width. If the parameter is set to STARNEIG_HESSENBERG_DEFAULT_PANEL_WIDTH, then the implementation will determine a suitable panel width automatically.
int	parallel_worker_size	The CPU variants of the panel reduction and trailing matrix update tasks are multithreaded. This parameter defines the number of cores allocated to these tasks. If the parameter is set to STARNEIG_HESSENBERG_DEFAULT_PARALLEL_WORKER_SIZE, then the implementation will determine a suitable CPU core count automatically.

13.10.2.2 struct starneig_schur_conf

Schur reduction configuration structure.

Data Fields

int	iteration_limit	The QR/QZ is an iterative algorithm. This parameter defines the maximum number of iterations the algorithm is allowed to perform. If the parameter is STARNEIG_SCHUR_DEFAULT_INTERATION_LIMIT, then the implementation will determine a suitable iteration limit automatically.
int	tile_size	The matrices are divided into square tiles. This parameter defines the used tile size. If the parameter is set to STARNEIG_SCHUR_DEFAULT_TILE_SIZE, then the implementation will determine a suitable tile size automatically.
int	small_limit	As the QR/QZ algorithm progresses, the size of the active region shrinks. Once the size of the active region is small enough, then the remaining problem is solved in a sequential manner. This parameter defines the transition point where the implementation switches to a sequential QR algorithm. If the parameter is set to STARNEIG_SCHUR_DEFAULT_SMALL_LIMIT, then the implementation will determine a suitable switching point automatically.
int	aed_window_size	The implementation relies on a so-called Aggressive Early Deflation (AED) technique to accelerate the convergence of the algorithm. Each AED is performed inside a small diagonal window. This parameter defines used AED window size. If the parameter is set to STARNEIG_SCHUR_DEFAULT_AED_WINDOW_SIZE, then the implementation will determine a suitable AED window size automatically.

Data Fields

int	aed_nibble	The implementation relies on a so-called Aggressive Early Deflation (AED) technique to accelerate the convergence of the algorithm. Each AED is performed inside a small diagonal window. If the number deflated (converged) eigenvalues is larger than `(aed_nibble / 100) ×` `size of AED window`, then the next bulge chasing step is skipped. If the parameter is set to STARNEIG_SCHUR_DEFAULT_AED_NIBBLE, then the implementation will determine a suitable value automatically.
int	aed_parallel_soft_limit	The implementation relies on a so-called Aggressive Early Deflation (AED) technique to accelerate the convergence of the algorithm. Each AED is performed inside a small diagonal window. An AED can be performed sequentially or in parallel. This parameter defines the transition point where the implementation allowed to switch to a sequential AED algorithm. The decision is made based on the size of the AED window. If the parameter is set to STARNEIG_SCHUR_DEFAULT_AED_PARALLEL_SOFT_LIMIT, then the implementation will determine a suitable switching point automatically.
int	aed_parallel_hard_limit	The implementation relies on a so-called Aggressive Early Deflation (AED) technique to accelerate the convergence of the algorithm. Each AED is performed inside a small diagonal window. An AED can be performed sequentially or in parallel. This parameter defines the transition point where the implementation switches to a sequential AED algorithm. The decision is made based on the size of the AED window. If the parameter is set to STARNEIG_SCHUR_DEFAULT_AED_PARALLEL_HARD_LIMIT, then the implementation will determine a suitable switching point automatically.
int	shift_count	The QR/QZ algorithm chases a set of 3×3 bulges across the diagonal of the Hessenberg(-triangular) decomposition. Two shifts (eigenvalue estimates) are required to generate each bulge. This parameter defines the number of shifts to use. If the parameter is set to STARNEIG_SCHUR_DEFAULT_SHIFT_COUNT, then the implementation will determine a suitable shift count automatically.
int	window_size	The QR/QZ algorithm chases a set of 3×3 bulges across the diagonal of the Hessenberg(-triangular) decomposition. The bulges are chased in batches. The related similarity transformations are initially restricted to inside a small diagonal window and the accumulated transformation are applied only later as BLAS-3 updates. This parameter defines the used bulge chasing window size. If the parameter is set to STARNEIG_SCHUR_ROUNDED_WINDOW_SIZE, then • maximum window size is set to 2 ∗ tile_size and • the windows are placed such that their lower right corners respect the boundaries of the underlying data tiles. If the parameter is set to STARNEIG_SCHUR_DEFAULT_WINDOW_SIZE, then the implementation will determine a suitable window size automatically.
int	shifts_per_window	The QR/QZ algorithm chases a set of 3×3 bulges across the diagonal of the Hessenberg(-triangular) decomposition. The bulges are chased in batches. This parameter defines the used batch size. If the parameter is set to STARNEIG_SCHUR_DEFAULT_SHIFTS_PER_WINDOW then the implementation will determine a suitable batch size automatically.

Data Fields

int	update_width	The similarity similarity transformations are initially restricted to inside a small diagonal window and the accumulated transformation are applied only later as BLAS-3 updates. This parameter defines the width of each left-hand side update task. The value should be multiple of the tile size. If the parameter is set to STARNEIG_SCHUR_DEFAULT_UPDATE_WIDTH, then the implementation will determine a suitable width automatically.
int	update_height	The similarity similarity transformations are initially restricted to inside a small diagonal window and the accumulated transformation are applied only later as BLAS-3 updates. This parameter defines the height of each right-hand side update task. The value should be multiple of the tile size. If the parameter is set to STARNEIG_SCHUR_DEFAULT_UPDATE_HEIGHT, then the implementation will determine a suitable height automatically.
double	left_threshold	The QR/QZ algorithm is allowed to set tiny matrix entires to zero as long as their magnitudes are smaller that a given threshold. This parameter defines the threshold for the left-hand side matrix (H). If the parameter is set to STARNEIG_SCHUR_DEFAULT_THRESHOLD, then the implementation will determine a suitable threshold automatically. If the parameter is set to STARNEIG_SCHUR_NORM_STABLE_THRESHOLD, then the implementation will use the threshold $u|H|_F$, where u is the unit roundoff and $|H|_F$ is the Frobenius norm of the matrix H. If the parameter is set to STARNEIG_SCHUR_LAPACK_THRESHOLD, then the implementation will use a deflation threshold that is compatible with LAPACK.
double	right_threshold	The QZ algorithm is allowed to set tiny matrix entires to zero as long as their magnitudes are smaller that a given threshold. This parameter defines the threshold for the right-hand side matrix (R) off-diagonal entires. If the parameter is set to STARNEIG_SCHUR_DEFAULT_THRESHOLD, then the implementation will determine a suitable threshold automatically. If the parameter is set to STARNEIG_SCHUR_NORM_STABLE_THRESHOLD, then the implementation will use the threshold $u|R|_F$, where u is the unit roundoff and $|H|_F$ is the Frobenius norm of the matrix R. If the parameter is set to STARNEIG_SCHUR_LAPACK_THRESHOLD, then the implementation will use a deflation threshold that is compatible with LAPACK.
double	inf_threshold	The QZ algorithm is allowed to set tiny matrix entires to zero as long as their magnitudes are smaller that a given threshold. This parameter defines the threshold for the right-hand side matrix (R) diagonal entries. If the parameter is set to STARNEIG_SCHUR_DEFAULT_THRESHOLD, then the implementation will determine a suitable threshold automatically. If the parameter is set to STARNEIG_SCHUR_NORM_STABLE_THRESHOLD, then the implementation will use the threshold $u|R|_F$, where u is the unit roundoff and $|R|_F$ is the Frobenius norm of the matrix R.

13.10.2.3 struct starneig_reorder_conf

Eigenvalue reordering configuration structure.

Data Fields

starneig_reorder_plan_t	plan	This parameter plan defines the used reordering plan. If the parameter is set to STARNEIG_REORDER_DEFAULT_PLAN, then the implementation will determine a suitable reordering plan automatically.
starneig_reorder_blueprint_t	blueprint	This parameter defines the used task insertion blueprint. If the parameter is set to STARNEIG_REORDER_DEFAULT_BLUEPRINT, then the implementation will determine a suitable task insertion blueprint automatically.
int	tile_size	The matrices are divided into square tiles. This parameter defines the used tile size. If the parameter is set to STARNEIG_REORDER_DEFAULT_TILE_SIZE, then the implementation will determine a suitable tile size automatically.
int	values_per_chain	The selected eigenvalues are processed in batches and each batch is assigned a window chain. This parameter defines the number of selected eigenvalues processed by each window chain. If the parameter is set to STARNEIG_REORDER_DEFAULT_VALUES_PER_CHAIN, then the implementation will determine a suitable value automatically.
int	window_size	The similarity similarity transformations are initially restricted to inside a small diagonal window and the accumulated transformation are applied only later as BLAS-3 updates. This parameter defines the size of the window. If the parameter is set to STARNEIG_REORDER_ROUNDED_WINDOW_SIZE, then • maximum window size is set to $2 *$ tile_size, • the windows are placed such that their upper left corners respect the boundaries of the underlying data tiles, and • the parameter values_per_chain is ignored. If the parameter is set to STARNEIG_REORDER_DEFAULT_WINDOW_SIZE, then the implementation will determine a suitable window size automatically.
int	small_window_size	Larger diagonal window are processed using even smaller diagonal windows in a recursive manner. This parameter defines the used small window size. If the parameter is set to STARNEIG_REORDER_DEFAULT_SMALL_WINDOW_SIZE, then the implementation will determine a suitable small window size automatically.

Data Fields

	int	small_window_threshold	Larger diagonal window are processed using even smaller diagonal windows in a recursive manner. This parameter defines the largest diagonal window that is pro- cessed in a scalar manner. If the parameter is set to STARNEIG_REORDER_DEFAULT_SMALL_WINDOW_THRESHOLD then the implementation will determine a suitable threshold automatically.
	int	update_width	The similarity similarity transformations are initially restricted to inside a small diagonal window and the accumulated transformation are applied only later as BLAS-3 updates. This parameter defines the width of each left-hand side update task. The value should be multiple of the tile size. If the parameter is set to STARNEIG_REORDER_DEFAULT_UPDATE_WIDTH, then the implementation will determine a suitable width automatically.
	int	update_height	The similarity similarity transformations are initially restricted to inside a small diagonal window and the accumulated transformation are applied only later as BLAS-3 updates. This parameter defines the height of each right-hand side update task. The value should be multiple of the tile size. If the parameter is set to STARNEIG_REORDER_DEFAULT_UPDATE_HEIGHT, then the implementation will determine a suitable height automatically.

13.10.2.4 struct starneig_eigenvectors_conf

Eigenvector computation configuration structure.

Data Fields

int	tile_size	The matrices are divided into tiles. This parameter defines the used tile size. If the parameter is set to STARNEIG_EIGENVECTORS_DEFAULT_TILE_SIZE, then the implementation will determine a suitable tile size automatically.

13.10.3 Enumeration Type Documentation

13.10.3.1 starneig_reorder_plan_t

enum starneig_reorder_plan_t

Reordering plan enumerator.

Eigenvalues that fall within a diagonal computation *window* are reordered such that all selected eigenvalues are moved to the upper left corner of the window. The corresponding orthogonal transformations are accumulated to separate accumulator matrix / matrices.

```
+--------+        +--------+
|¤ x x x x|       |$ x x x x|        $ selected eigenvalue
| $ x x x|        | $ x x x|        ¤ deselected eigenvalue
|   ¤ x x| ==> Q  |   ¤ x x| Q^T   x non-zero entry
|     ¤ x|        |     ¤ x|
|       $|        |       ¤|
+--------+        +--------+
```

A *window chain* comprises from multiple overlapping diagonal computation windows that are intended to be processed in a particular order. More precisely, the windows are placed such that the overlap between two windows is big enough to accommodate all selected eigenvalues that fall within the preceding windows. In this way, the windows can be processed in sequential order, starting from the bottom window, such that the reordering that takes place in one window always moves the preceding selected eigenvalues to the lower right corner of the next window. In the end, all selected that fall within the combined computation area of the chain are moved to the upper left corner of the topmost window.

An example showing how an eigenvalue can be moved six entries upwards by using three diagonal windows:

```
+-------+        +-------+        +-------+        +-------+
| x x x |  ===>  | x x x |  ===>  | x x x |  ===>  | #<--+ |
|   x-x-+---+     |   x-x-+---+     |   x-x-+---+     |   ¤-|-+---+
|   | x x x |     |   | x x x |     |   | #<--+ |     |   | ¤ ¤ ¤ |
+---+---x-x-+---+ +---+---x-x-+---+ +---+---¤-|-+---+ +---+---¤-¤-+---+
    |   | x x x |     |   | #<--+ |     |   | ¤ ¤ ¤ |     |   | ¤ ¤ ¤ |
    +---+---x x |     +---+---¤ | |     +---+---¤ ¤ |     +---+---¤ ¤ |
        |     # |         |     ¤ |         |     ¤ |         |     ¤ |
        +-------+         +-------+         +-------+         +-------+
```

The number of selected eigenvalues that can be moved by a single window chain is limited by the windows size. Thus, the whole reordering procedure usually involves multiple chains that must be be processed in a particular order. A *chain list* describes a list of chains that are intended to be processed together. Window chains that belong to different chain lists are processed separately.

A *plan* consists from one or more chain lists that are intended to be processed in a particular order.

STARNEIG_REORDER_ONE_PART_PLAN :

The first chain is placed in the upper left corner of the matrix and its size is chosen such that it contains a desired number of selected eigenvalues (starneig_reorder_conf::values_per_chain parameter). The next chain is places such that its upper left corner is located one entry after the location where the last selected eigenvalue, that falls within the first chain, would be after the reordering. The chain is sized such that the part of the chain, that does not intersect the first chain, contain the desired number of selected eigenvalues. This same procedure is repeated until all selected eigenvalues have been accounted for. All chains belong to the same chain lists and are intended to be processed sequentially.

An example showing the placement of the chains in a case where each chain wields two selected eigenvalues:

```
                                                        before:      after:
.-------+                                               .            0
| 0     | <-- chain 0                                   0 _          1_
|   .----------+                                         . .   ====> 2
| | . |        | <-- chain 1                             . . _       3_
+---|---1--------------+ . . . . . . . . . . . . 1_. .             4
    |   | .    |       | <-- chain 2                      . . _      5_
    |   | 2-----------------+                             2 . .      6
    |   | | . |       |     | <-- chain 3                 . . . _    7_
    +---|---|---3------------------------+ . . . . . . 3_. . .      8
        |   | | 4     |     |            | <-- chain 4    4 . .      9
        |   | | | .   |     |            |                . . .      .
        |   | | |  .  |     |            |                . . .      .
        +---|---|------5    |            |                . . 5_. .      .
            |   |       6   |            |                     6 .      .
            |   |        .  |            |                      . .      .
            |   |         . |            |                      . .      .
            +---|-------------7          |                      . 7_. .      .
                |           . |          |                        . .      .
                |            . |         |                         .      .
                |             8 |        |                          8      .
                +-----------------------9 . . . . . . . . . . 9      .
```

An example showing what happens when the first three chains are processed:

```
.-------+                           >0<------+
| 0      |                          |>1<      |
|   .-----------+                   |   #-----------+
|   | . |        |                  |   | # |        |
+---|---1--------------+            +---|---#--------------+
    |   | . |        |              |   | . |        |
    |   |   2   |      |    ===>     |   |   2   |      |    ===>
    |   |    . |      |              |   |    . |      |
    +---|-------3      |             +---|-------3      |
        |        4     |                 |        4     |
        |         .    |                 |         .    |
        |          . |                   |          . |
        +--------------5                 +--------------5
0-------+                           0-------+
| 1      |                          | 1      |
|   >2<----------+                  |   2-----------+
|   |>3<|        |                  |   | 3 |        |
+---|---#--------------+            +---|-->4<--------------+
    |   | # |        |              |   |>5<      |      |
    |   |   #   |      |    ===>     |   |   #   |      |
    |   |    # |      |              |   |    # |      |
    +---|-------#      |             +---|-------#      |
        |        4     |                 |        #     |
        |         .    |                 |         #    |
        |          . |                   |          # |
        +--------------5                 +--------------#
```

If necessary, each chain is re-sized to avoid splitting any 2×2 tiles.

Windows are placed such that the first window is located in the lower right corner of the computation area of the window chain. The last window is correspondingly placed in the upper left corner of the computation area.

If necessary, each window is re-sized to avoid splitting any 2×2 tiles.

STARNEIG_REORDER_MULTI_PART_PLAN :

A multi-part reordering plan is derived from an one-part reordering plan by splitting the chains into sub-chains as shown below:

```
Initial one-part plan:
  Chain 0: aaaaaa
  Chain 1:   bbbbbbbbbb              a,b,c,d,e diagonal computation window
  Chain 2:     cccccccccccccc
  Chain 3:       dddddddddddddddddddd
  Chain 4:         eeeeeeeeeeeeeeeeeeeee
Resulting multi-part plan:
  Chain 0: aaaaaa
  Chain 1:  ......bbbb
  Chain 2:  ..........cccc             chain list 0
  Chain 3:  ..............dddd
  Chain 4:  ..................eeee
  ---------------------------------------------------
  Chain 0:   bbbbbb....
  Chain 1:   ......cccc....            chain list 1
  Chain 2:   ..........dddd....
  Chain 3:   ..............eeee....
  ---------------------------------------------------
  Chain 0:    cccccc........
  Chain 1:    ......dddd........       chain list 2
  Chain 2:    ..........eeee........
  ---------------------------------------------------
  Chain 0:    dddddd............       chain list 3
  Chain 1:    ......eeee............
  ---------------------------------------------------
  Chain 0:    eeeeee...............    chain list 4
```

Note that the chains that belong to the same chain list are independent from each other and can therefore be processed in an arbitrary order.

Enumerator

STARNEIG_REORDER_DEFAULT_PLAN	Default plan.
STARNEIG_REORDER_ONE_PART_PLAN	One part plan.
STARNEIG_REORDER_MULTI_PART_PLAN	Multi part plan.

13.10.3.2 starneig_reorder_blueprint_t

enum starneig_reorder_blueprint_t

Task insertion blueprint.

A task insertion blueprint defines how a reordering plan is carried out.

Enumerator

STARNEIG_REORDER_DEFAULT_BLUEPRINT	Default blueprint.
STARNEIG_REORDER_DUMMY_INSERT_A	One-pass forward dummy blueprint. Processes the window chains in order starting from the topmost chain. All update tasks are inserted right after each window reordering task.
STARNEIG_REORDER_DUMMY_INSERT_B	Two-pass backward dummy blueprint. Processes the window chains in two phases starting from the bottommost chain. The window reordering tasks and the right-hand side update tasks are inserted during the first phase. Other update tasks are inserted during the second phase.
STARNEIG_REORDER_CHAIN_INSERT_A	One-pass forward chain blueprint. Processes the window chains in order starting from the topmost chain. The window reordering tasks and high priority right-hand side update tasks are inserted first. Other update tasks are inserted after them.
STARNEIG_REORDER_CHAIN_INSERT_B	Two-pass forward chain blueprint. Processes the window chains in two phases starting from the topmost chain. The window reordering tasks and high priority right-hand side update tasks are inserted first. The left- hand side updates are inserted after them. Other updates are inserted during the second phase.
STARNEIG_REORDER_CHAIN_INSERT_C	One-pass backward chain blueprint. Processes the window chains in order starting from the bottommost chain. The window reordering tasks and high priority right-hand side update tasks are inserted first. Other update tasks are inserted later.
STARNEIG_REORDER_CHAIN_INSERT_D	Two-pass backward chain blueprint. Processes the window chains in two phases starting from the bottommost chain. The window reordering tasks and high priority right-hand side update tasks are inserted first. Update tasks that are related to the Schur matrix are inserted later. Update tasks that are related to the orthogonal matrices are inserted during the second phase.
STARNEIG_REORDER_CHAIN_INSERT_E	Two-pass delayed backward chain blueprint. Processes the window chains in order starting from the bottommost chain. The window reordering tasks and high priority right-hand side update tasks are inserted first. Update tasks that are related to the Schur matrix are inserted later. Update tasks that are related to the orthogonal matrices are inserted only after all chain list have been processed.

Enumerator

STARNEIG_REORDER_CHAIN_INSERT_F	Three-pass delayed backward chain blueprint. Processes the window chains in two phases starting from the bottommost chain. The window reordering tasks and high priority right-hand side update tasks are inserted during the first phase. Update tasks that are related to the Schur matrix are inserted during the second phase. Update tasks that are related to the orthogonal matrices are inserted only after all chain list have been processed.

13.10.4 Function Documentation

13.10.4.1 starneig_hessenberg_init_conf()

```
void starneig_hessenberg_init_conf (
            struct starneig_hessenberg_conf * conf )
```

Initializes a Hessenberg reduction configuration structure with default parameters.

Parameters

out	*conf*	The Hessenberg reduction configuration structure.

13.10.4.2 starneig_schur_init_conf()

```
void starneig_schur_init_conf (
            struct starneig_schur_conf * conf )
```

Initializes a Schur reduction configuration structure with default parameters.

Parameters

out	*conf*	The Schur reduction configuration structure.

13.10.4.3 starneig_reorder_init_conf()

```
void starneig_reorder_init_conf (
            struct starneig_reorder_conf * conf )
```

Initializes an eigenvalue reordering configuration structure with default parameters.

Parameters

out	*conf*	The eigenvalue reordering configuration structure.

13.10.4.4   starneig_eigenvectors_init_conf()

```
void starneig_eigenvectors_init_conf (
            struct starneig_eigenvectors_conf * conf )
```

Initializes an eigenvectors configuration structure with default parameters.

Parameters

out	*conf*	The eigenvectors configuration structure.

13.11 ScaLAPACK compatibility / BLACS matrices

Data types and functions for BLACS formatted distributed matrices.

Data Structures

- struct starneig_blacs_descr

    BLACS descriptor. More...

BLACS contexts

- typedef int starneig_blacs_context_t

    BLACS context.
- starneig_blacs_context_t starneig_distr_to_blacs_context (starneig_distr_t distr)

    Convers a data distribution to a BLACS context.
- starneig_distr_t starneig_blacs_context_to_distr (starneig_blacs_context_t context)

    Convers a BLACS context to a data distribution.
- int starneig_distr_is_blacs_compatible (starneig_distr_t distr)

    Checks whether a data distribution is BLACS compatible.
- int starneig_distr_is_compatible_with (starneig_distr_t distr, starneig_blacs_context_t context)

    Checks whether a data distribution is compatible with a given BLACS context.

BLACS descriptors

- typedef struct starneig_blacs_descr starneig_blacs_descr_t

    BLACS descriptor.
- void starneig_blacs_create_matrix (int rows, int cols, int row_blksz, int col_blksz, starneig_datatype_t type, starneig_blacs_context_t context, starneig_blacs_descr_t ∗descr, void ∗∗local)

    Creates a BLACS matrix with uninitialized matrix elements.
- void starneig_create_blacs_matrix (int rows, int cols, int row_blksz, int col_blksz, starneig_datatype_t type, starneig_blacs_context_t context, starneig_blacs_descr_t ∗descr, void ∗∗local)

    Creates a BLACS matrix with uninitialized matrix elements. Deprecated.
- void starneig_blacs_destroy_matrix (starneig_blacs_descr_t ∗descr, void ∗∗local)

    Destroyes a BLACS matrix.
- void starneig_destroy_blacs_matrix (starneig_blacs_descr_t ∗descr, void ∗∗local)

    Destroyes a BLACS matrix. Deprecated.
- void starneig_distr_matrix_to_blacs_descr (starneig_distr_matrix_t matrix, starneig_blacs_context_t context, starneig_blacs_descr_t ∗descr, void ∗∗local)

    Convers a distributed matrix to a BLACS descriptor and a matching local array.
- starneig_distr_matrix_t starneig_blacs_descr_to_distr_matrix (starneig_datatype_t type, starneig_distr_t distr, starneig_blacs_descr_t ∗descr, void ∗local)

    Convers a BLACS descriptor and a matching local array to a distributed matrix.
- int starneig_distr_matrix_is_blacs_compatible (starneig_distr_matrix_t matrix)

    Checks whether a distributed matrix is BLACS compatible.
- int starneig_distr_matrix_is_compatible_with (starneig_distr_matrix_t matrix, starneig_blacs_context_t context)

    Checks whether a distributed matrix is compatible with a given BLACS context.

13.11.1 Detailed Description

Data types and functions for BLACS formatted distributed matrices.

13.11.2 Data Structure Documentation

13.11.2.1 struct starneig_blacs_descr

BLACS descriptor.

Data Fields

int	type	The descriptor type.
starneig_blacs_context_t	context	The related BLACS context.
int	m	The number of (global) rows in the matrix.
int	n	The number of (global) columns in the matrix.
int	sm	The number of rows in a distribution block.
int	sn	The number of columns in a distribution block.
int	rsrc	The process grid row over which the first row is distributed.
int	csrc	The process grid column over which the first column is distributed.
int	lld	The leading dimension of the local array.

13.11.3 Function Documentation

13.11.3.1 starneig_distr_to_blacs_context()

starneig_blacs_context_t starneig_distr_to_blacs_context (
          starneig_distr_t *distr*)

Convers a data distribution to a BLACS context.

Attention

> The data distribution must describe a two-dimensional block cyclic distribution.

Parameters

in	*distr*	The data distribution.

Returns

> The BLACS context.

13.11.3.2 starneig_blacs_context_to_distr()

starneig_distr_t starneig_blacs_context_to_distr (
            starneig_blacs_context_t *context*)

Convers a BLACS context to a data distribution.

Parameters

in	*context*	The BLACS context.

Returns

    The data distribution.

13.11.3.3 starneig_distr_is_blacs_compatible()

int starneig_distr_is_blacs_compatible (
            starneig_distr_t *distr*)

Checks whether a data distribution is BLACS compatible.

Parameters

in	*distr*	The data distribution.

Returns

    Non-zero if the data distribution matrix is BLACS compatible.

13.11.3.4 starneig_distr_is_compatible_with()

int starneig_distr_is_compatible_with (
            starneig_distr_t *distr,*
            starneig_blacs_context_t *context*)

Checks whether a data distribution is compatible with a given BLACS context.

Parameters

in	*distr*	The data distribution.
in	*context*	The BLACS context.

Returns

Non-zero if the data distribution compatible with the BLACS context.

13.11.3.5 starneig_blacs_create_matrix()

```
void starneig_blacs_create_matrix (
            int rows,
            int cols,
            int row_blksz,
            int col_blksz,
            starneig_datatype_t type,
            starneig_blacs_context_t context,
            starneig_blacs_descr_t * descr,
            void ** local )
```

Creates a BLACS matrix with uninitialized matrix elements.

Parameters

in	*rows*	The number of (global) rows in the matrix.
in	*cols*	The number of (global) columns in the matrix.
in	*row_blksz*	The number of rows in a distribution block. Can be set to -1 in which case the library decides the value.
in	*col_blksz*	The number of columns in a distribution block. Can be set to -1 in which case the library decides the value.
in	*type*	The matrix element data type.
in	*context*	The BLACS context.
out	*descr*	The BLACS descriptor.
out	*local*	A pointer to the local array.

13.11.3.6 starneig_create_blacs_matrix()

```
void starneig_create_blacs_matrix (
            int rows,
            int cols,
            int row_blksz,
            int col_blksz,
            starneig_datatype_t type,
            starneig_blacs_context_t context,
            starneig_blacs_descr_t * descr,
            void ** local )
```

Creates a BLACS matrix with uninitialized matrix elements. Deprecated.

Deprecated The starneig_create_blacs_matrix() function has been replaced with the starneig_blacs_create_matrix() function. This function will be removed in a future release of the library.

13.11.3.7   starneig_blacs_destroy_matrix()

```
void starneig_blacs_destroy_matrix (
            starneig_blacs_descr_t * descr,
            void ** local )
```

Destroyes a BLACS matrix.

Parameters

in,out	*descr*	The BLACS descriptor.
in,out	*local*	A pointer to the local array.

13.11.3.8   starneig_destroy_blacs_matrix()

```
void starneig_destroy_blacs_matrix (
            starneig_blacs_descr_t * descr,
            void ** local )
```

Destroyes a BLACS matrix. Deprecated.

Deprecated The starneig_destroy_blacs_matrix() function has been replaced with the starneig_blacs_destroy_matrix() function. This function will be removed in a future release of the library.

13.11.3.9   starneig_distr_matrix_to_blacs_descr()

```
void starneig_distr_matrix_to_blacs_descr (
            starneig_distr_matrix_t matrix,
            starneig_blacs_context_t context,
            starneig_blacs_descr_t * descr,
            void ** local )
```

Convers a distributed matrix to a BLACS descriptor and a matching local array.

This function creates a wrapper object. The contents of the distributed matrix may be modified by the functions that use the wrapper object.
```
starneig_distr_matrix_t dA = starneig_distr_matrix_create(...);
...
starneig_distr_t distr = starneig_distr_matrix_get_distr(A);
starneig_blacs_context_t context = starneig_distr_to_blacs_context(distr);
starneig_blacs_descr_t descr_a;
double *local_a;
starneig_distr_matrix_to_blacs_descr(
    dA, context, &descr_a, (void **)&local_a);
```

Parameters

in	*matrix*	The distributed matrix.
in	*context*	The BLACS context. The context must have been converted from the same data distribution the distributed matrix is using or vice versa.
out	*descr*	The BLACS descriptor.
out	*local*	A pointer to the local array.

13.11.3.10   starneig_blacs_descr_to_distr_matrix()

starneig_distr_matrix_t starneig_blacs_descr_to_distr_matrix (
        starneig_datatype_t *type,*
        starneig_distr_t *distr,*
        starneig_blacs_descr_t * *descr,*
        void * *local*)

Convers a BLACS descriptor and a matching local array to a distributed matrix.

This function creates a wrapper object. The contents of the local array may be modified by the functions that use the wrapper object. The starneig_distr_matrix_destroy() function does not de-initilize the BLACS descriptor nor free the local array.

```
starneig_blacs_context_t context;
starneig_blacs_descr_t descr_a;
double *local_a;
...
starneig_distr_t distr = starneig_blacs_context_to_distr(context);
starneig_distr_matrix_t dA = starneig_blacs_descr_to_distr_matrix(
    STARNEIG_REAL_DOUBLE, distr, descr_a, (void *)local_a);
```

Parameters

in	*type*	The matrix element data type.
in	*distr*	The data distribution. The data distribution must have been converted from the same BLACS context the BLACS descriptor is using or vice versa.
in	*descr*	The BLACS descriptor.
in	*local*	A pointer to the local array.

Returns

    The distributed matrix.

13.11.3.11   starneig_distr_matrix_is_blacs_compatible()

int starneig_distr_matrix_is_blacs_compatible (
        starneig_distr_matrix_t *matrix*)

Checks whether a distributed matrix is BLACS compatible.

Parameters

in	*matrix*	The distributed matrix.

Returns

    Non-zero if the distributed matrix is BLACS compatible.

13.11.3.12 starneig_distr_matrix_is_compatible_with()

```
int starneig_distr_matrix_is_compatible_with (
            starneig_distr_matrix_t matrix,
            starneig_blacs_context_t context )
```

Checks whether a distributed matrix is compatible with a given BLACS context.

Parameters

in	*matrix*	The distributed matrix.
in	*context*	The BLACS context.

Returns

Non-zero if the distributed matrix compatible with the BLACS context.

13.12   ScaLAPACK compatibility / BLACS helpers

Data types and helper functions for BLACS.

Functions

- void starneig_blacs_pinfo (int ∗my_rank, int ∗rank_count)

    Queries process rank information.
- int starneig_blacs_get (starneig_blacs_context_t context, starneig_blacs_query_id_t query)

    Returns BLACS context's internal defaults.
- starneig_blacs_context_t starneig_blacs_gridinit (starneig_blacs_context_t system_context, char ∗order, int rows, int cols)

    Initializes a BLACS process grid.
- void starneig_blacs_gridinfo (starneig_blacs_context_t context, int ∗rows, int ∗cols, int ∗row, int ∗col)

    Queries BLACS process grid information.
- void starneig_blacs_pcoord (starneig_blacs_context_t context, int process, int ∗row, int ∗col)

    Queries BLACS process grid coordinates.
- void starneig_blacs_gridexit (starneig_blacs_context_t context)

    Releases process grid specific resources.
- void starneig_blacs_exit (int cont)

    Releases all contexts and related resources.
- int starneig_blacs_numroc (int n, int nb, int iproc, int isrcproc, int nprocs)

    Computes the number of matrix rows/columns owned by a given process.
- int starneig_numroc (int n, int nb, int iproc, int isrcproc, int nprocs)

    Computes the number of matrix rows/columns owned by a given process. Deprecated.
- int starneig_blacs_descinit (struct starneig_blacs_descr ∗descr, int m, int n, int sm, int sn, int irsrc, int icsrc, starneig_blacs_context_t context, int ld)

    Initializes a BLACS descriptor.
- int starneig_descinit (struct starneig_blacs_descr ∗descr, int m, int n, int sm, int sn, int irsrc, int icsrc, starneig_blacs_context_t context, int ld)

    Initializes a BLACS descriptor. Deprecated.

Query indeces

- typedef int starneig_blacs_query_id_t

    Data type for blacs_get() function query id.
- #define STARNEIG_BLACS_GET_DEFAULT_CONTEXT 0

    Query id for getting the default system context.

13.12.1   Detailed Description

Data types and helper functions for BLACS.

13.12.2   Function Documentation

13.12.2.1   starneig_blacs_pinfo()

```
void starneig_blacs_pinfo (
            int * my_rank,
            int * rank_count )
```

Queries process rank information.

Parameters

out	*my_rank*	An unique process id (rank).
out	*rank_count*	The total number of processes (ranks) available.

13.12.2.2 starneig_blacs_get()

```
int starneig_blacs_get (
            starneig_blacs_context_t context,
            starneig_blacs_query_id_t query )
```

Returns BLACS context's internal defaults.

Parameters

in	*context*	The BLACS context.
in	*query*	The query id.

Returns

The internal default value that matches the given query id.

13.12.2.3 starneig_blacs_gridinit()

```
starneig_blacs_context_t starneig_blacs_gridinit (
            starneig_blacs_context_t system_context,
            char * order,
            int rows,
            int cols )
```

Initializes a BLACS process grid.

Parameters

in	*system_context*	The system BLACS context to be used in creating the process grid.
in	*order*	The process mapping order. "R" : Use row-major natural ordering. "C" : Use column-major natural ordering. ELSE: Use row-major natural ordering.
in	*rows*	The number of rows in the process grid.
in	*cols*	The number of columns in the process grid.

Returns

A handle to the created BLACS context.

13.12.2.4 starneig_blacs_gridinfo()

```
void starneig_blacs_gridinfo (
            starneig_blacs_context_t context,
            int * rows,
            int * cols,
            int * row,
            int * col )
```

Queries BLACS process grid information.

Parameters

in	*context*	The BLACS context.
out	*rows*	The number of rows in the process grid.
out	*cols*	The number of columns in the process grid.
out	*row*	The row coordinate of the calling process.
out	*col*	The column coordinate of the calling process.

13.12.2.5 starneig_blacs_pcoord()

```
void starneig_blacs_pcoord (
            starneig_blacs_context_t context,
            int process,
            int * row,
            int * col )
```

Queries BLACS process grid coordinates.

Parameters

in	*context*	The BLACS context.
in	*process*	The process id (rank).
out	*row*	The row coordinate of the process.
out	*col*	The column coordinate of the process.

13.12.2.6 starneig_blacs_gridexit()

```
void starneig_blacs_gridexit (
            starneig_blacs_context_t context )
```

Releases process grid specific resources.

Parameters

in	*context*	The BLACS context.

13.12.2.7 starneig_blacs_exit()

```
void starneig_blacs_exit (
            int cont )
```

Releases all contexts and related resources.

Parameters

in	*cont*	The continue flag.

13.12.2.8 starneig_blacs_numroc()

```
int starneig_blacs_numroc (
            int n,
            int nb,
            int iproc,
            int isrcproc,
            int nprocs )
```

Computes the number of matrix rows/columns owned by a given process.

Parameters

in	*n*	The number of rows/columns in the distributed matrix.
in	*nb*	The block size.
in	*iproc*	The coordinate of the process whose local array row or column is to be determined.
in	*isrcproc*	The coordinate of the process that possesses the first row or column of the distributed matrix.
in	*nprocs*	The total number processes over which the matrix is distributed.

Returns

The number of rows/columns owned by the process.

13.12.2.9 starneig_numroc()

```
int starneig_numroc (
            int n,
            int nb,
            int iproc,
            int isrcproc,
            int nprocs )
```

Computes the number of matrix rows/columns owned by a given process. Deprecated.

Deprecated The starneig_numroc() function has been replaced with the starneig_blacs_numroc() function. This function will be removed in a future release of the library.

13.12.2.10 starneig_blacs_descinit()

```
int starneig_blacs_descinit (
            struct starneig_blacs_descr * descr,
            int m,
            int n,
            int sm,
            int sn,
            int irsrc,
            int icsrc,
            starneig_blacs_context_t context,
            int ld )
```

Initializes a BLACS descriptor.

Parameters

out	*descr*	The matrix descriptor.
in	*m*	The number of rows in the matrix.
in	*n*	The number of columns in the matrix.
in	*sm*	The number of rows in a distributed block.
in	*sn*	The number of columns in a distributed block.
in	*irsrc*	The process grid row over which the first row is distributed.
in	*icsrc*	The process grid column over which the first column is distributed.
in	*context*	The BLACS context.
in	*ld*	The local array leading dimension.

Returns

Zero if the initialization was successful, non-zero otherwise.

13.12.2.11 starneig_descinit()

```
int starneig_descinit (
            struct starneig_blacs_descr * descr,
            int m,
            int n,
            int sm,
            int sn,
            int irsrc,
            int icsrc,
            starneig_blacs_context_t context,
            int ld )
```

Initializes a BLACS descriptor. Deprecated.

Deprecated The starneig_descinit() function has been replaced with the starneig_blacs_descinit() function. This function will be removed in a future release of the library.

Chapter 14

File Documentation

14.1 blacs_helpers.h File Reference

This file contains various BLACS helper functions.

```
#include <starneig/configuration.h>
#include <starneig/blacs_matrix.h>
```

Functions

- void starneig_blacs_pinfo (int ∗my_rank, int ∗rank_count)

  Queries process rank information.
- int starneig_blacs_get (starneig_blacs_context_t context, starneig_blacs_query_id_t query)

  Returns BLACS context's internal defaults.
- starneig_blacs_context_t starneig_blacs_gridinit (starneig_blacs_context_t system_context, char ∗order, int rows, int cols)

  Initializes a BLACS process grid.
- void starneig_blacs_gridinfo (starneig_blacs_context_t context, int ∗rows, int ∗cols, int ∗row, int ∗col)

  Queries BLACS process grid information.
- void starneig_blacs_pcoord (starneig_blacs_context_t context, int process, int ∗row, int ∗col)

  Queries BLACS process grid coordinates.
- void starneig_blacs_gridexit (starneig_blacs_context_t context)

  Releases process grid specific resources.
- void starneig_blacs_exit (int cont)

  Releases all contexts and related resources.
- int starneig_blacs_numroc (int n, int nb, int iproc, int isrcproc, int nprocs)

  Computes the number of matrix rows/columns owned by a given process.
- int starneig_numroc (int n, int nb, int iproc, int isrcproc, int nprocs)

  Computes the number of matrix rows/columns owned by a given process. Deprecated.
- int starneig_blacs_descinit (struct starneig_blacs_descr ∗descr, int m, int n, int sm, int sn, int irsrc, int icsrc, starneig_blacs_context_t context, int ld)

  Initializes a BLACS descriptor.
- int starneig_descinit (struct starneig_blacs_descr ∗descr, int m, int n, int sm, int sn, int irsrc, int icsrc, starneig_blacs_context_t context, int ld)

  Initializes a BLACS descriptor. Deprecated.

Query indeces

- #define STARNEIG_BLACS_GET_DEFAULT_CONTEXT 0

    Query id for getting the default system context.

- typedef int starneig_blacs_query_id_t

    Data type for blacs_get() function query id.

14.1.1   Detailed Description

This file contains various BLACS helper functions.

Author

   Mirko Myllykoski (`mirkom@cs.umu.se`), Umeå University

14.2   blacs_matrix.h File Reference

This file contains data types and functions for BLACS formatted distributed matrices.

```
#include <starneig/configuration.h>
#include <starneig/distr_matrix.h>
```

Data Structures

- struct starneig_blacs_descr

    BLACS descriptor. More...

BLACS contexts

- typedef int starneig_blacs_context_t

    BLACS context.

- starneig_blacs_context_t starneig_distr_to_blacs_context (starneig_distr_t distr)

    Convers a data distribution to a BLACS context.

- starneig_distr_t starneig_blacs_context_to_distr (starneig_blacs_context_t context)

    Convers a BLACS context to a data distribution.

- int starneig_distr_is_blacs_compatible (starneig_distr_t distr)

    Checks whether a data distribution is BLACS compatible.

- int starneig_distr_is_compatible_with (starneig_distr_t distr, starneig_blacs_context_t context)

    Checks whether a data distribution is compatible with a given BLACS context.

BLACS descriptors

- typedef struct starneig_blacs_descr starneig_blacs_descr_t

    BLACS descriptor.
- void starneig_blacs_create_matrix (int rows, int cols, int row_blksz, int col_blksz, starneig_datatype_t type, starneig_blacs_context_t context, starneig_blacs_descr_t ∗descr, void ∗∗local)

    Creates a BLACS matrix with uninitialized matrix elements.
- void starneig_create_blacs_matrix (int rows, int cols, int row_blksz, int col_blksz, starneig_datatype_t type, starneig_blacs_context_t context, starneig_blacs_descr_t ∗descr, void ∗∗local)

    Creates a BLACS matrix with uninitialized matrix elements. Deprecated.
- void starneig_blacs_destroy_matrix (starneig_blacs_descr_t ∗descr, void ∗∗local)

    Destroyes a BLACS matrix.
- void starneig_destroy_blacs_matrix (starneig_blacs_descr_t ∗descr, void ∗∗local)

    Destroyes a BLACS matrix. Deprecated.
- void starneig_distr_matrix_to_blacs_descr (starneig_distr_matrix_t matrix, starneig_blacs_context_t context, starneig_blacs_descr_t ∗descr, void ∗∗local)

    Convers a distributed matrix to a BLACS descriptor and a matching local array.
- starneig_distr_matrix_t starneig_blacs_descr_to_distr_matrix (starneig_datatype_t type, starneig_distr_t distr, starneig_blacs_descr_t ∗descr, void ∗local)

    Convers a BLACS descriptor and a matching local array to a distributed matrix.
- int starneig_distr_matrix_is_blacs_compatible (starneig_distr_matrix_t matrix)

    Checks whether a distributed matrix is BLACS compatible.
- int starneig_distr_matrix_is_compatible_with (starneig_distr_matrix_t matrix, starneig_blacs_context_t context)

    Checks whether a distributed matrix is compatible with a given BLACS context.

14.2.1 Detailed Description

This file contains data types and functions for BLACS formatted distributed matrices.

Author

Mirko Myllykoski (mirkom@cs.umu.se), Umeå University

14.3 configuration.h File Reference

This file contains StarNEig library configuration.

Macros

- #define STARNEIG_VERSION_MAJOR

    StarNEig major version.
- #define STARNEIG_VERSION_MINOR

    StarNEig minor version.
- #define STARNEIG_VERSION_PATCH

    StarNEig patch level.
- #define STARNEIG_ENABLE_MPI

    MPI support enabled.

- #define STARNEIG_ENABLE_CUDA

    CUDA support enabled.
- #define STARNEIG_ENABLE_BLACS

    BLACS support enabled.
- #define STARNEIG_SEP_DM_HESSENBERG

    Distributed memory Hessenberg reduction enabled.
- #define STARNEIG_GEP_DM_HESSENBERGTRIANGULAR

    Distributed memory Hessenberg-triangular reduction enabled.
- #define STARNEIG_SEP_DM_REDUCE

    Full distributed memory reduction enabled.
- #define STARNEIG_GEP_DM_REDUCE

    Full distributed memory generalized reduction enabled.

14.3.1 Detailed Description

This file contains StarNEig library configuration.

Author

        Mirko Myllykoski (mirkom@cs.umu.se), Umeå University

14.4 distr_helpers.h File Reference

This file contains generic distributed memory interface functions.

```
#include <starneig/configuration.h>
#include <stddef.h>
#include <mpi.h>
```

Functions

MPI communicator

- void starneig_mpi_set_comm (MPI_Comm comm)
    Sets a MPI communicator for the library.
- MPI_Comm starneig_mpi_get_comm ()
    Returns the library MPI communicator.

Broadcast

- void starneig_mpi_broadcast (int root, size_t size, void ∗buffer)
    Broadcast a buffer.
- void starneig_broadcast (int root, size_t size, void ∗buffer)
    Broadcast a buffer. Deprecated.

14.4.1 Detailed Description

This file contains generic distributed memory interface functions.

Author

Mirko Myllykoski (mirkom@cs.umu.se), Umeå University

14.5 distr_matrix.h File Reference

This file contains data types and functions for distributed matrices.

```
#include <starneig/configuration.h>
#include <stddef.h>
```

Data Structures

- struct starneig_distr_block

    Distributed block. More...

Functions

- void **starneig_broadcast** (int root, size_t size, void ∗buffer)

Query functions

- void   starneig_distr_matrix_get_blocks   (starneig_distr_matrix_t   matrix,   struct   starneig_distr_block ∗∗blocks, int ∗num_blocks)

    Returns the locally owned distributed blocks.
- starneig_distr_t starneig_distr_matrix_get_distr (starneig_distr_matrix_t matrix)

    Returns the distribution that is associated with a distributed matrix.
- starneig_datatype_t starneig_distr_matrix_get_datatype (starneig_distr_matrix_t matrix)

    Returns the matrix element data type.
- size_t starneig_distr_matrix_get_elemsize (starneig_distr_matrix_t matrix)

    Returns the matrix element size.
- int starneig_distr_matrix_get_rows (starneig_distr_matrix_t matrix)

    Returns the number of (global) rows.
- int starneig_distr_matrix_get_cols (starneig_distr_matrix_t matrix)

    Returns the number of (global) columns.
- int starneig_distr_matrix_get_row_blksz (starneig_distr_matrix_t matrix)

    Returns the number of rows in a distribution block.
- int starneig_distr_matrix_get_col_blksz (starneig_distr_matrix_t matrix)

    Returns the number of columns in a distribution block.

Data distributions

- enum   starneig_distr_order_t   {   STARNEIG_ORDER_DEFAULT,   STARNEIG_ORDER_ROW_MAJOR, STARNEIG_ORDER_COL_MAJOR }

  Process mapping order.
- typedef struct starneig_distr ∗ starneig_distr_t

  Data distribution.
- starneig_distr_t starneig_distr_init ()

  Creates a default data distribution.
- starneig_distr_t starneig_distr_init_mesh (int rows, int cols, starneig_distr_order_t order)

  Creates a two-dimensional block cyclic data distribution.
- starneig_distr_t starneig_distr_init_func (int(∗func)(int row, int col, void ∗arg), void ∗arg, size_t arg_size)

  Creates a distribution using a data distribution function.
- starneig_distr_t starneig_distr_duplicate (starneig_distr_t distr)

  Duplicates a data distribution.
- void starneig_distr_destroy (starneig_distr_t distr)

  Destroys a data distribution.

Distributed matrices

- enum starneig_datatype_t { STARNEIG_REAL_DOUBLE }

  Distributed matrix element data type.
- typedef struct starneig_distr_matrix ∗ starneig_distr_matrix_t

  Distributed matrix.
- starneig_distr_matrix_t starneig_distr_matrix_create (int rows, int cols, int row_blksz, int col_blksz, starneig_datatype_t type, starneig_distr_t distr)

  Creates a distributed matrix with uninitialized matrix elements.
- starneig_distr_matrix_t starneig_distr_matrix_create_local (int rows, int cols, starneig_datatype_t type, int owner, double ∗A, int ldA)

  Creates a single-owner distributed matrix from a local matrix.
- void starneig_distr_matrix_destroy (starneig_distr_matrix_t matrix)

  Destroys a distributed matrix.
- void starneig_distr_matrix_copy (starneig_distr_matrix_t source, starneig_distr_matrix_t dest)

  Copies the contents of a distributed matrix to a second distributed matrix.
- void starneig_distr_matrix_copy_region (int sr, int sc, int dr, int dc, int rows, int cols, starneig_distr_matrix_t source, starneig_distr_matrix_t dest)

  Copies region of a distributed matrix to a second distributed matrix.

14.5.1   Detailed Description

This file contains data types and functions for distributed matrices.

Author

    Mirko Myllykoski (`mirkom@cs.umu.se`), Umeå University

    Lars Karlsson (`larsk@cs.umu.se`), Umeå University

14.6   error.h File Reference

This file contains the library error codes.

```
#include <starneig/configuration.h>
```

Macros

- #define STARNEIG_SUCCESS 0

    Success.
- #define STARNEIG_GENERIC_ERROR 1

    Reneric error.
- #define STARNEIG_NOT_INITIALIZED 2

    Not initialized.
- #define STARNEIG_INVALID_CONFIGURATION 3

    Invalid configuration.
- #define STARNEIG_INVALID_ARGUMENTS 4

    Invalid argument.
- #define STARNEIG_INVALID_DISTR_MATRIX 5

    Invalid distributed matrix.
- #define STARNEIG_DID_NOT_CONVERGE 6

    Did not converge.
- #define STARNEIG_PARTIAL_REORDERING 7

    Partial reordering.
- #define STARNEIG_CLOSE_EIGENVALUES 8

    Close eigenvalues.

Typedefs

- typedef int starneig_error_t

    Interface function return value data type.

14.6.1   Detailed Description

This file contains the library error codes.

Author

   Mirko Myllykoski (mirkom@cs.umu.se), Umeå University

14.7   expert.h File Reference

This file contains configuration structures and functions for the expert interface functions.

```
#include <starneig/configuration.h>
```

Data Structures

- struct starneig_hessenberg_conf

  Hessenberg reduction configuration structure. More...
- struct starneig_schur_conf

  Schur reduction configuration structure. More...
- struct starneig_reorder_conf

  Eigenvalue reordering configuration structure. More...
- struct starneig_eigenvectors_conf

  Eigenvector computation configuration structure. More...

Hessenberg reduction

- #define STARNEIG_HESSENBERG_DEFAULT_TILE_SIZE -1

  Default tile size.
- #define STARNEIG_HESSENBERG_DEFAULT_PANEL_WIDTH -1

  Default panel width.
- #define STARNEIG_HESSENBERG_DEFAULT_PARALLEL_WORKER_SIZE -1

  Default parallel worker size.
- void starneig_hessenberg_init_conf (struct starneig_hessenberg_conf ∗conf)

  Initializes a Hessenberg reduction configuration structure with default parameters.

Schur reduction

- #define STARNEIG_SCHUR_DEFAULT_INTERATION_LIMIT -1

  Default iteration limit.
- #define STARNEIG_SCHUR_DEFAULT_TILE_SIZE -1

  Default tile size.
- #define STARNEIG_SCHUR_DEFAULT_SMALL_LIMIT -1

  Default sequential QR limit.
- #define STARNEIG_SCHUR_DEFAULT_AED_WINDOW_SIZE -1

  Default AED window size.
- #define STARNEIG_SCHUR_DEFAULT_AED_NIBBLE -1

  Default nibble value.
- #define STARNEIG_SCHUR_DEFAULT_AED_PARALLEL_SOFT_LIMIT -1

  Default soft sequential AED limit.
- #define STARNEIG_SCHUR_DEFAULT_AED_PARALLEL_HARD_LIMIT -1

  Default hard sequential AED limit.
- #define STARNEIG_SCHUR_DEFAULT_SHIFT_COUNT -1

  Default shift count.
- #define STARNEIG_SCHUR_DEFAULT_WINDOW_SIZE -1

  Default bulge chasing window size.
- #define STARNEIG_SCHUR_ROUNDED_WINDOW_SIZE -2

  Rounded bulge chasing window.
- #define STARNEIG_SCHUR_DEFAULT_SHIFTS_PER_WINDOW -1

  Default number of shifts per bulge chasing window.
- #define STARNEIG_SCHUR_DEFAULT_UPDATE_WIDTH -1

  Default left-hand side update width.
- #define STARNEIG_SCHUR_DEFAULT_UPDATE_HEIGHT -1

       Default right-hand side update height.

- #define STARNEIG_SCHUR_DEFAULT_THRESHOLD -1

       Default deflation threshold.

- #define STARNEIG_SCHUR_NORM_STABLE_THRESHOLD -2

       Norm stable deflation threshold.

- #define STARNEIG_SCHUR_LAPACK_THRESHOLD -3

       LAPACK-style deflation threshold.

- void starneig_schur_init_conf (struct starneig_schur_conf ∗conf)

       Initializes a Schur reduction configuration structure with default parameters.

Eigenvalue reordering

- #define STARNEIG_REORDER_DEFAULT_UPDATE_WIDTH -1

       Default left-hand side update task width.

- #define STARNEIG_REORDER_DEFAULT_UPDATE_HEIGHT -1

       Default right-hand side update task height.

- #define STARNEIG_REORDER_DEFAULT_TILE_SIZE -1

       Default tile size.

- #define STARNEIG_REORDER_DEFAULT_VALUES_PER_CHAIN -1

       Default number of selected eigenvalues per window.

- #define STARNEIG_REORDER_DEFAULT_WINDOW_SIZE -1

       Default default window size.

- #define STARNEIG_REORDER_ROUNDED_WINDOW_SIZE -2

       Default rounded window size.

- #define STARNEIG_REORDER_DEFAULT_SMALL_WINDOW_SIZE -1

       Default small window size.

- #define STARNEIG_REORDER_DEFAULT_SMALL_WINDOW_THRESHOLD -1

       Default small window threshold.

- enum starneig_reorder_plan_t { STARNEIG_REORDER_DEFAULT_PLAN = 1, STARNEIG_REORDER_ONE_PART_PLAN = 2, STARNEIG_REORDER_MULTI_PART_PLAN = 3 }

       Reordering plan enumerator.

- enum starneig_reorder_blueprint_t {
STARNEIG_REORDER_DEFAULT_BLUEPRINT = 1, STARNEIG_REORDER_DUMMY_INSERT_A = 2,
STARNEIG_REORDER_DUMMY_INSERT_B = 3, STARNEIG_REORDER_CHAIN_INSERT_A = 4,
STARNEIG_REORDER_CHAIN_INSERT_B = 5, STARNEIG_REORDER_CHAIN_INSERT_C = 6,
STARNEIG_REORDER_CHAIN_INSERT_D = 7, STARNEIG_REORDER_CHAIN_INSERT_E = 8,
STARNEIG_REORDER_CHAIN_INSERT_F = 9 }

       Task insertion blueprint.

- void starneig_reorder_init_conf (struct starneig_reorder_conf ∗conf)

       Initializes an eigenvalue reordering configuration structure with default parameters.

Eigenvectors

- #define STARNEIG_EIGENVECTORS_DEFAULT_TILE_SIZE -1

       Default tile size.

- void starneig_eigenvectors_init_conf (struct starneig_eigenvectors_conf ∗conf)

       Initializes an eigenvectors configuration structure with default parameters.

14.7.1 Detailed Description

This file contains configuration structures and functions for the expert interface functions.

Author

Mirko Myllykoski (mirkom@cs.umu.se), Umeå University

Angelika Schwarz (angies@cs.umu.se), Umeå University

14.8 gep_dm.h File Reference

This file contains distributed memory interface functions for generalized eigenvalue problems.

```
#include <starneig/configuration.h>
#include <starneig/error.h>
#include <starneig/expert.h>
#include <starneig/distr_matrix.h>
```

Functions

Computational functions

- starneig_error_t starneig_GEP_DM_HessenbergTriangular (starneig_distr_matrix_t A, starneig_distr_matrix_t B, starneig_distr_matrix_t Q, starneig_distr_matrix_t Z)

  Computes a Hessenberg-triangular decomposition of a general matrix pencil.
- starneig_error_t starneig_GEP_DM_Schur (starneig_distr_matrix_t H, starneig_distr_matrix_t R, starneig_distr_matrix_t Q, starneig_distr_matrix_t Z, double real[ ], double imag[ ], double beta[ ])

  Computes a generalized Schur decomposition given a Hessenberg-triangular decomposition.
- starneig_error_t starneig_GEP_DM_ReorderSchur (int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t T, starneig_distr_matrix_t Q, starneig_distr_matrix_t Z, double real[ ], double imag[ ], double beta[ ])

  Reorders selected generalized eigenvalues to the top left corner of a generalized Schur decomposition.
- starneig_error_t starneig_GEP_DM_Reduce (starneig_distr_matrix_t A, starneig_distr_matrix_t B, starneig_distr_matrix_t Q, starneig_distr_matrix_t Z, double real[ ], double imag[ ], double beta[ ], int(∗predicate)(double real, double imag, double beta, void ∗arg), void ∗arg, int selected[ ], int ∗num_↩ selected)

  Computes a (reordered) generalized Schur decomposition given a general matrix pencil.
- starneig_error_t starneig_GEP_DM_Eigenvectors (int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t T, starneig_distr_matrix_t Z, starneig_distr_matrix_t X)

  Computes a generalized eigenvector for each selected generalized eigenvalue.

Helper functions

- starneig_error_t starneig_GEP_DM_Select (starneig_distr_matrix_t S, starneig_distr_matrix_t T, int(∗predicate)(double real, double imag, double beta, void ∗arg), void ∗arg, int selected[ ], int ∗num_↩ selected)

  Generates a selection array for a Schur-triangular matrix pencil using a user-supplied predicate function.

Expert computational functions

- starneig_error_t starneig_GEP_DM_Schur_expert (struct starneig_schur_conf ∗conf, starneig_distr_matrix_t H, starneig_distr_matrix_t R, starneig_distr_matrix_t Q, starneig_distr_matrix_t Z, double real[ ], double imag[ ])

Computes a generalized Schur decomposition given a Hessenberg-triangular decomposition.
- starneig_error_t starneig_GEP_DM_ReorderSchur_expert (struct starneig_reorder_conf ∗conf, int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t T, starneig_distr_matrix_t Q, starneig_distr_matrix_t Z, double real[ ], double imag[ ], double beta[ ])

  Reorders selected generalized eigenvalues to the top left corner of a generalized Schur decomposition.
- starneig_error_t starneig_GEP_DM_Eigenvectors_expert (struct starneig_eigenvectors_conf ∗conf, int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t T, starneig_distr_matrix_t Z, starneig_distr_matrix_t X)

  Computes a generalized eigenvector for each selected generalized eigenvalue.

14.8.1 Detailed Description

This file contains distributed memory interface functions for generalized eigenvalue problems.

Author

Mirko Myllykoski (mirkom@cs.umu.se), Umeå University

Lars Karlsson (larsk@cs.umu.se), Umeå University

14.9 gep_sm.h File Reference

This file contains shared memory interface functions for generalized eigenvalue problems.

```
#include <starneig/configuration.h>
#include <starneig/error.h>
#include <starneig/expert.h>
```

Functions

Computational functions

- starneig_error_t starneig_GEP_SM_HessenbergTriangular (int n, double A[ ], int ldA, double B[ ], int ldB, double Q[ ], int ldQ, double Z[ ], int ldZ)

  Computes a Hessenberg-triangular decomposition of a general matrix pencil.
- starneig_error_t starneig_GEP_SM_Schur (int n, double H[ ], int ldH, double R[ ], int ldR, double Q[ ], int ldQ, double Z[ ], int ldZ, double real[ ], double imag[ ], double beta[ ])

  Computes a generalized Schur decomposition given a Hessenberg-triangular decomposition.
- starneig_error_t starneig_GEP_SM_ReorderSchur (int n, int selected[ ], double S[ ], int ldS, double T[ ], int ldT, double Q[ ], int ldQ, double Z[ ], int ldZ, double real[ ], double imag[ ], double beta[ ])

  Reorders selected generalized eigenvalues to the top left corner of a generalized Schur decomposition.
- starneig_error_t starneig_GEP_SM_Reduce (int n, double A[ ], int ldA, double B[ ], int ldB, double Q[ ], int ldQ, double Z[ ], int ldZ, double real[ ], double imag[ ], double beta[ ], int(∗predicate)(double real, double imag, double beta, void ∗arg), void ∗arg, int selected[ ], int ∗num_selected)

  Computes a (reordered) generalized Schur decomposition given a general matrix pencil.
- starneig_error_t starneig_GEP_SM_Eigenvectors (int n, int selected[ ], double S[ ], int ldS, double T[ ], int ldT, double Z[ ], int ldZ, double X[ ], int ldX)

  Computes a generalized eigenvector for each selected generalized eigenvalue.

Helper functions

- starneig_error_t starneig_GEP_SM_Select (int n, double S[ ], int ldS, double T[ ], int ldT, int(∗predicate)(double real, double imag, double beta, void ∗arg), void ∗arg, int selected[ ], int ∗num_selected)

    Generates a selection array for a Schur-triangular matrix pencil using a user-supplied predicate function.

Expert computational functions

- starneig_error_t starneig_GEP_SM_Schur_expert (struct starneig_schur_conf ∗conf, int n, double H[ ], int ldH, double R[ ], int ldR, double Q[ ], int ldQ, double Z[ ], int ldZ, double real[ ], double imag[ ], double beta[ ])

    Computes a generalized Schur decomposition given a Hessenberg-triangular decomposition.

- starneig_error_t starneig_GEP_SM_ReorderSchur_expert (struct starneig_reorder_conf ∗conf, int n, int selected[ ], double S[ ], int ldS, double T[ ], int ldT, double Q[ ], int ldQ, double Z[ ], int ldZ, double real[ ], double imag[ ], double beta[ ])

    Reorders selected eigenvalues to the top left corner of a generalized Schur decomposition.

- starneig_error_t starneig_GEP_SM_Eigenvectors_expert (struct starneig_eigenvectors_conf ∗conf, int n, int selected[ ], double S[ ], int ldS, double T[ ], int ldT, double Z[ ], int ldZ, double X[ ], int ldX)

    Computes a generalized eigenvector for each selected generalized eigenvalue.

14.9.1 Detailed Description

This file contains shared memory interface functions for generalized eigenvalue problems.

Author

Mirko Myllykoski (mirkom@cs.umu.se), Umeå University

Lars Karlsson (larsk@cs.umu.se), Umeå University

14.10 node.h File Reference

This file contains interface to configure the intra-node execution environment.

```
#include <starneig/configuration.h>
#include <stddef.h>
#include <mpi.h>
```

Functions

- void starneig_node_init (int cores, int gpus, starneig_flag_t flags)

    Initializes the intra-node execution environment.

- int starneig_node_initialized ()

    Checks whether the intra-node execution environment is initialized.

- int starneig_node_get_cores ()

    Returns the number of cores (threads) per MPI rank.

- void starneig_node_set_cores (int cores)

    Changes the number of CPUs cores (threads) to use per MPI rank.

- int starneig_node_get_gpus ()

    Returns the number of GPUs per MPI rank.

- void starneig_node_set_gpus (int gpus)

    Changes the number of GPUs to use per MPI rank.

- void starneig_node_finalize ()

Deallocates resources associated with the intra-node configuration.
- void **starneig_mpi_set_comm** (MPI_Comm comm)
- MPI_Comm **starneig_mpi_get_comm** ()

Pinned host memory

- void starneig_node_enable_pinning ()
    Enable CUDA host memory pinning.
- void starneig_node_disable_pinning ()
    Disables CUDA host memory pinning.

Library initialization flags

- #define STARNEIG_USE_ALL -1

    Use all resources.
- #define STARNEIG_DEFAULT 0x0

    Default mode.
- #define STARNEIG_HINT_SM 0x0

    Shared memory mode.
- #define STARNEIG_HINT_DM 0x1

    Distributed memory mode.
- #define STARNEIG_FXT_DISABLE 0x2

    No FxT traces mode.
- #define STARNEIG_AWAKE_WORKERS 0x4

    Awake worker mode.
- #define STARNEIG_AWAKE_MPI_WORKER 0x8

    Awake MPI worker mode.
- #define STARNEIG_FAST_DM (STARNEIG_HINT_DM | STARNEIG_AWAKE_WORKERS | STARNEIG_AWAKE_MPI_WORK

    Fast distributed memory mode.
- #define STARNEIG_NO_VERBOSE 0x10

    No verbose mode.
- #define STARNEIG_NO_MESSAGES (STARNEIG_NO_VERBOSE | 0x20)

    No messages mode.
- typedef unsigned starneig_flag_t

    Library initialization flag data type.

14.10.1 Detailed Description

This file contains interface to configure the intra-node execution environment.

Author

Mirko Myllykoski (mirkom@cs.umu.se), Umeå University

Lars Karlsson (larsk@cs.umu.se), Umeå University

14.11 sep_dm.h File Reference

This file contains distributed memory interface functions for standard eigenvalue problems.

```
#include <starneig/configuration.h>
#include <starneig/error.h>
#include <starneig/expert.h>
#include <starneig/distr_matrix.h>
```

Functions

Computational functions

- starneig_error_t starneig_SEP_DM_Hessenberg (starneig_distr_matrix_t A, starneig_distr_matrix_t Q)

  Computes a Hessenberg decomposition of a general matrix.
- starneig_error_t starneig_SEP_DM_Schur (starneig_distr_matrix_t H, starneig_distr_matrix_t Q, double real[ ], double imag[ ])

  Computes a Schur decomposition given a Hessenberg decomposition.
- starneig_error_t starneig_SEP_DM_ReorderSchur (int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t Q, double real[ ], double imag[ ])

  Reorders selected eigenvalues to the top left corner of a Schur decomposition.
- starneig_error_t starneig_SEP_DM_Reduce (starneig_distr_matrix_t A, starneig_distr_matrix_t Q, double real[ ], double imag[ ], int(∗predicate)(double real, double imag, void ∗arg), void ∗arg, int selected[ ], int ∗num_selected)

  Computes a (reordered) Schur decomposition of a general matrix.
- starneig_error_t starneig_SEP_DM_Eigenvectors (int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t Q, starneig_distr_matrix_t X)

  Computes an eigenvector for each selected eigenvalue.

Helper functions

- starneig_error_t starneig_SEP_DM_Select (starneig_distr_matrix_t S, int(∗predicate)(double real, double imag, void ∗arg), void ∗arg, int selected[ ], int ∗num_selected)

  Generates a selection array for a Schur matrix using a user-supplied predicate function.

Expert computational functions

- starneig_error_t starneig_SEP_DM_Schur_expert (struct starneig_schur_conf ∗conf, starneig_distr_matrix_t H, starneig_distr_matrix_t Q, double real[ ], double imag[ ])

  Computes a Schur decomposition given a Hessenberg decomposition.
- starneig_error_t starneig_SEP_DM_ReorderSchur_expert (struct starneig_reorder_conf ∗conf, int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t Q, double real[ ], double imag[ ])

  Reorders selected eigenvalues to the top left corner of a Schur decomposition.
- starneig_error_t starneig_SEP_DM_Eigenvectors_expert (struct starneig_eigenvectors_conf ∗conf, int selected[ ], starneig_distr_matrix_t S, starneig_distr_matrix_t Q, starneig_distr_matrix_t X)

  Computes an eigenvector for each selected eigenvalue.

14.11.1 Detailed Description

This file contains distributed memory interface functions for standard eigenvalue problems.

Author

Mirko Myllykoski (mirkom@cs.umu.se), Umeå University

Lars Karlsson (larsk@cs.umu.se), Umeå University

14.12 sep_sm.h File Reference

This file contains shared memory interface functions for standard eigenvalue problems.

```
#include <starneig/configuration.h>
#include <starneig/error.h>
#include <starneig/expert.h>
```

Functions

Computational functions

- starneig_error_t starneig_SEP_SM_Hessenberg (int n, double A[ ], int ldA, double Q[ ], int ldQ)

  Computes a Hessenberg decomposition of a general matrix.
- starneig_error_t starneig_SEP_SM_Schur (int n, double H[ ], int ldH, double Q[ ], int ldQ, double real[ ], double imag[ ])

  Computes a Schur decomposition given a Hessenberg decomposition.
- starneig_error_t starneig_SEP_SM_ReorderSchur (int n, int selected[ ], double S[ ], int ldS, double Q[ ], int ldQ, double real[ ], double imag[ ])

  Reorders selected eigenvalues to the top left corner of a Schur decomposition.
- starneig_error_t starneig_SEP_SM_Reduce (int n, double A[ ], int ldA, double Q[ ], int ldQ, double real[ ], double imag[ ], int(∗predicate)(double real, double imag, void ∗arg), void ∗arg, int selected[ ], int ∗num_↩ selected)

  Computes a (reordered) Schur decomposition of a general matrix.
- starneig_error_t starneig_SEP_SM_Eigenvectors (int n, int selected[ ], double S[ ], int ldS, double Q[ ], int ldQ, double X[ ], int ldX)

  Computes an eigenvector for each selected eigenvalue.

Helper functions

- starneig_error_t starneig_SEP_SM_Select (int n, double S[ ], int ldS, int(∗predicate)(double real, double imag, void ∗arg), void ∗arg, int selected[ ], int ∗num_selected)

  Generates a selection array for a Schur matrix using a user-supplied predicate function.

Expert computational functions

- starneig_error_t starneig_SEP_SM_Hessenberg_expert (struct starneig_hessenberg_conf ∗conf, int n, int begin, int end, double A[ ], int ldA, double Q[ ], int ldQ)

  Computes a Hessenberg decomposition of a general matrix.
- starneig_error_t starneig_SEP_SM_Schur_expert (struct starneig_schur_conf ∗conf, int n, double H[ ], int ldH, double Q[ ], int ldQ, double real[ ], double imag[ ])

  Computes a Schur decomposition given a Hessenberg decomposition.
- starneig_error_t starneig_SEP_SM_ReorderSchur_expert (struct starneig_reorder_conf ∗conf, int n, int selected[ ], double S[ ], int ldS, double Q[ ], int ldQ, double real[ ], double imag[ ])

  Reorders selected eigenvalues to the top left corner of a Schur decomposition.
- starneig_error_t starneig_SEP_SM_Eigenvectors_expert (struct starneig_eigenvectors_conf ∗conf, int n, int selected[ ], double S[ ], int ldS, double Q[ ], int ldQ, double X[ ], int ldX)

  Computes an eigenvector for each selected eigenvalue.

14.12.1 Detailed Description

This file contains shared memory interface functions for standard eigenvalue problems.

Author

Mirko Myllykoski (mirkom@cs.umu.se), Umeå University

Lars Karlsson (larsk@cs.umu.se), Umeå University

14.13 starneig.h File Reference

This file includes most StarNEig header files.

```
#include <starneig/configuration.h>
#include <starneig/node.h>
#include <starneig/gep_sm.h>
#include <starneig/sep_sm.h>
#include <starneig/distr_helpers.h>
#include <starneig/gep_dm.h>
#include <starneig/sep_dm.h>
#include <starneig/blacs_helpers.h>
#include <starneig/blacs_matrix.h>
```

14.13.1 Detailed Description

This file includes most StarNEig header files.

Author

> Mirko Myllykoski (mirkom@cs.umu.se), Umeå University

Chapter 15

Example Documentation

15.1 gep_dm_full_chain.c

```c
#include "validate.h"
#include <stdlib.h>
#include <stdio.h>
#include <time.h>
#include <mpi.h>
#include <starneig/starneig.h>
// a predicate function that selects all finate eigenvalues that have positive
// a real part
static int predicate(double real, double imag, double beta, void *arg)
{
    if (0.0 < real && beta != 0.0)
        return 1;
    return 0;
}
int main(int argc, char **argv)
{
    const int n = 3000; // matrix dimension
    const int root = 0; // root rank
    // initialize MPI
    int thread_support;
    MPI_Init_thread(
        &argc, (char ***)&argv, MPI_THREAD_MULTIPLE, &thread_support);
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    // the root node initializes the matrices locally
    int ldA = 0, ldB = 0, ldQ = 0, ldZ = 0, ldC = 0, ldD = 0;
    double *A = NULL, *B = NULL, *Q = NULL, *Z = NULL, *C = NULL, *D = NULL;
    if (world_rank == root) {
        srand((unsigned) time(NULL));
        // generate a full random matrix A and a copy C
        ldA = ((n/8)+1)*8, ldC = ((n/8)+1)*8;
        A = malloc(n*ldA*sizeof(double));
        C = malloc(n*ldC*sizeof(double));
        for (int j = 0; j < n; j++)
            for (int i = 0; i < n; i++)
                A[j*ldA+i] = C[j*ldC+i] = 2.0*rand()/RAND_MAX - 1.0;
        // generate a full random matrix B and a copy D
        ldB = ((n/8)+1)*8, ldD = ((n/8)+1)*8;
        B = malloc(n*ldB*sizeof(double));
        D = malloc(n*ldD*sizeof(double));
        for (int j = 0; j < n; j++)
            for (int i = 0; i < n; i++)
                B[j*ldB+i] = D[j*ldD+i] = 2.0*rand()/RAND_MAX - 1.0;
        // generate an identity matrix Q
        ldQ = ((n/8)+1)*8;
        Q = malloc(n*ldA*sizeof(double));
        for (int j = 0; j < n; j++)
            for (int i = 0; i < n; i++)
                Q[j*ldQ+i] = i == j ? 1.0 : 0.0;
        // generate an identity matrix Z
        ldZ = ((n/8)+1)*8;
        Z = malloc(n*ldZ*sizeof(double));
        for (int j = 0; j < n; j++)
            for (int i = 0; i < n; i++)
                Z[j*ldZ+i] = i == j ? 1.0 : 0.0;
    }
    // allocate space for the eigenvalues and the eigenvalue selection vector
    double *real = malloc(n*sizeof(double));
```

```c
double *imag = malloc(n*sizeof(double));
double *beta = malloc(n*sizeof(double));
int *select = malloc(n*sizeof(int));
// Initialize the StarNEig library using all available CPU cores and
// GPUs. The STARNEIG_FAST_DM flag indicates that the library should
// initialize itself for distributed memory computations and keep StarPU
// worker threads and StarPU-MPI communication thread awake between
// interface function calls.
starneig_node_init(STARNEIG_USE_ALL, STARNEIG_USE_ALL, STARNEIG_FAST_DM);
// create a two-dimensional block cyclic distribution with column-major
// ordering
starneig_distr_t distr = starneig_distr_init_mesh(
    -1, -1, STARNEIG_ORDER_COL_MAJOR);
// Convert the local matrix A to a distributed matrix lA that is owned by
// the root node. This is done in-place, i.e., the matrices A and lA point
// to the same data.
starneig_distr_matrix_t lA = starneig_distr_matrix_create_local(
    n, n, STARNEIG_REAL_DOUBLE, root, A, ldA);
// create a distributed matrix dA using the earlier created data
// distribution and default distributed block size
starneig_distr_matrix_t dA =
    starneig_distr_matrix_create(n, n, -1, -1, STARNEIG_REAL_DOUBLE, distr);
// copy the local matrix lA to the distributed matrix dA (scatter)
starneig_distr_matrix_copy(lA, dA);
// scatter the matrix B
starneig_distr_matrix_t lB = starneig_distr_matrix_create_local(
    n, n, STARNEIG_REAL_DOUBLE, root, B, ldB);
starneig_distr_matrix_t dB =
    starneig_distr_matrix_create(n, n, -1, -1, STARNEIG_REAL_DOUBLE, distr);
starneig_distr_matrix_copy(lB, dB);
// scatter the matrix Q
starneig_distr_matrix_t lQ = starneig_distr_matrix_create_local(
    n, n, STARNEIG_REAL_DOUBLE, root, Q, ldQ);
starneig_distr_matrix_t dQ =
    starneig_distr_matrix_create(n, n, -1, -1, STARNEIG_REAL_DOUBLE, distr);
starneig_distr_matrix_copy(lQ, dQ);
// scatter the matrix Z
starneig_distr_matrix_t lZ = starneig_distr_matrix_create_local(
    n, n, STARNEIG_REAL_DOUBLE, root, Z, ldZ);
starneig_distr_matrix_t dZ =
    starneig_distr_matrix_create(n, n, -1, -1, STARNEIG_REAL_DOUBLE, distr);
starneig_distr_matrix_copy(lZ, dZ);
// reduce the dense-dense matrix pair (A,B) to Hessenberg-triangular form
printf("Hessenberg-triangular reduction...\n");
starneig_GEP_DM_HessenbergTriangular(dA, dB, dQ, dZ);
// reduce the Hessenberg-triangular matrix pair (A,B) to generalized Schur
// form
printf("Schur reduction...\n");
starneig_GEP_DM_Schur(dA, dB, dQ, dZ, real, imag, beta);
// select eigenvalues that have positive a real part
int num_selected;
starneig_GEP_DM_Select(dA, dB, &predicate, NULL, select, &num_selected);
printf("Selected %d eigenvalues out of %d.\n", num_selected, n);
// reorder selected eigenvalues to the upper left corner of the generalized
// Schur form (A,B)
printf("Reordering...\n");
starneig_GEP_DM_ReorderSchur(select, dA, dB, dQ, dZ, real, imag, beta);
// copy the distributed matrix dA back to the local matrix lA (gather)
starneig_distr_matrix_copy(dA, lA);
// free the distributed matrix lA (matrix A is not freed)
starneig_distr_matrix_destroy(lA);
// free the distributed matrix dA (all local resources are freed)
starneig_distr_matrix_destroy(dA);
// gather the matrix B
starneig_distr_matrix_copy(dB, lB);
starneig_distr_matrix_destroy(lB);
starneig_distr_matrix_destroy(dB);
// gather the matrix Q
starneig_distr_matrix_copy(dQ, lQ);
starneig_distr_matrix_destroy(lQ);
starneig_distr_matrix_destroy(dQ);
// gather the matrix Z
starneig_distr_matrix_copy(dZ, lZ);
starneig_distr_matrix_destroy(lZ);
starneig_distr_matrix_destroy(dZ);
// free the data distribution
starneig_distr_destroy(distr);
// de-initialize the StarNEig library
starneig_node_finalize();
// de-initialize MPI
MPI_Finalize();
if (world_rank == root) {
    // check residual || Q A Z^T - C ||_F / || C ||_F
    check_residual(n, ldQ, ldA, ldZ, ldC, Q, A, Z, C);
    // check residual || Q B Z^T - D ||_F / || D ||_F
    check_residual(n, ldQ, ldB, ldZ, ldD, Q, B, Z, D);
    // check residual || Q Q^T - I ||_F / || I ||_F
```

```c
        check_orthogonality(n, ldQ, Q);
        // check residual || Z Z^T - I ||_F / || I ||_F
        check_orthogonality(n, ldZ, Z);
    }
    // cleanup
    free(A);
    free(C);
    free(B);
    free(D);
    free(Q);
    free(Z);
    free(real);
    free(imag);
    free(beta);
    free(select);
    return 0;
}
```

15.2 gep_sm_eigenvectors.c

```c
#include "validate.h"
#include <stdlib.h>
#include <stdio.h>
#include <time.h>
#include <starneig/starneig.h>
// a predicate function that selects all finite eigenvalues that have positive
// a real part
static int predicate(double real, double imag, double beta, void *arg)
{
    if (0.0 < real && beta != 0.0)
        return 1;
    return 0;
}
int main()
{
    const int n = 3000; // matrix dimension
    srand((unsigned) time(NULL));
    // generate a full random matrix A and a copy C
    int ldA = ((n/8)+1)*8, ldC = ((n/8)+1)*8;
    double *A = malloc(n*ldA*sizeof(double));
    double *C = malloc(n*ldC*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            A[j*ldA+i] = C[j*ldC+i] = 2.0*rand()/RAND_MAX - 1.0;
    // generate a full random matrix B and a copy D
    int ldB = ((n/8)+1)*8, ldD = ((n/8)+1)*8;
    double *B = malloc(n*ldB*sizeof(double));
    double *D = malloc(n*ldD*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            B[j*ldB+i] = D[j*ldD+i] = 2.0*rand()/RAND_MAX - 1.0;
    // generate an identity matrix Q
    int ldQ = ((n/8)+1)*8;
    double *Q = malloc(n*ldA*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            Q[j*ldQ+i] = i == j ? 1.0 : 0.0;
    // generate an identity matrix Z
    int ldZ = ((n/8)+1)*8;
    double *Z = malloc(n*ldZ*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            Z[j*ldZ+i] = i == j ? 1.0 : 0.0;
    double *X = NULL; int ldX = 0;
    // allocate space for the eigenvalues and the eigenvalue selection vector
    double *real = malloc(n*sizeof(double));
    double *imag = malloc(n*sizeof(double));
    double *beta = malloc(n*sizeof(double));
    int *select = malloc(n*sizeof(int));
    // Initialize the StarNEig library using all available CPU cores and
    // GPUs. The STARNEIG_HINT_SM flag indicates that the library should
    // initialize itself for shared memory computations.
    starneig_node_init(STARNEIG_USE_ALL, STARNEIG_USE_ALL, STARNEIG_HINT_SM);
    // reduce the dense-dense matrix pair (A,B) to generalized Schur form
    // (skip reordering)
    printf("Reduce...\n");
    starneig_GEP_SM_Reduce(
        n, A, ldA, B, ldB, Q, ldQ, Z, ldZ, real, imag, beta,
        NULL, NULL, NULL, NULL);
    // select eigenvalues that have positive a real part and allocate space for
    // the eigenvectors
    int num_selected;
    starneig_GEP_SM_Select(
```

```
        n, A, ldA, B, ldB, &predicate, NULL, select, &num_selected);
    printf("Selected %d eigenvalues out of %d.\n", num_selected, n);
    ldX = ((n/8)+1)*8;
    X = malloc(num_selected*ldX*sizeof(double));
    // compute a selected set of eigenvectors
    printf("Eigenvectors...\n");
    starneig_GEP_SM_Eigenvectors(n, select, A, ldA, B, ldB, Q, ldQ, X, ldX);
    // de-initialize the StarNEig library
    starneig_node_finalize();
    // check residual || Q A Z^T - C ||_F / || C ||_F
    check_residual(n, ldQ, ldA, ldZ, ldC, Q, A, Z, C);
    // check residual || Q B Z^T - D ||_F / || D ||_F
    check_residual(n, ldQ, ldB, ldZ, ldD, Q, B, Z, D);
    // check residual || Q Q^T - I ||_F / || I ||_F
    check_orthogonality(n, ldQ, Q);
    // check residual || Z Z^T - I ||_F / || I ||_F
    check_orthogonality(n, ldZ, Z);
    // cleanup
    free(A);
    free(C);
    free(B);
    free(D);
    free(Q);
    free(Z);
    free(X);
    free(real);
    free(imag);
    free(beta);
    free(select);
    return 0;
}
```

15.3   gep_sm_full_chain.c

```
#include "validate.h"
#include <stdlib.h>
#include <stdio.h>
#include <time.h>
#include <starneig/starneig.h>
// a predicate function that selects all finate eigenvalues that have positive
// a real part
static int predicate(double real, double imag, double beta, void *arg)
{
    if (0.0 < real && beta != 0.0)
        return 1;
    return 0;
}
int main()
{
    const int n = 3000; // matrix dimension
    srand((unsigned) time(NULL));
    // generate a full random matrix A and a copy C
    int ldA = ((n/8)+1)*8, ldC = ((n/8)+1)*8;
    double *A = malloc(n*ldA*sizeof(double));
    double *C = malloc(n*ldC*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            A[j*ldA+i] = C[j*ldC+i] = 2.0*rand()/RAND_MAX - 1.0;
    // generate a full random matrix B and a copy D
    int ldB = ((n/8)+1)*8, ldD = ((n/8)+1)*8;
    double *B = malloc(n*ldB*sizeof(double));
    double *D = malloc(n*ldD*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            B[j*ldB+i] = D[j*ldD+i] = 2.0*rand()/RAND_MAX - 1.0;
    // generate an identity matrix Q
    int ldQ = ((n/8)+1)*8;
    double *Q = malloc(n*ldA*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            Q[j*ldQ+i] = i == j ? 1.0 : 0.0;
    // generate an identity matrix Z
    int ldZ = ((n/8)+1)*8;
    double *Z = malloc(n*ldZ*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            Z[j*ldZ+i] = i == j ? 1.0 : 0.0;
    // allocate space for the eigenvalues and the eigenvalue selection vector
    double *real = malloc(n*sizeof(double));
    double *imag = malloc(n*sizeof(double));
    double *beta = malloc(n*sizeof(double));
    int *select = malloc(n*sizeof(int));
    // Initialize the StarNEig library using all available CPU cores and
```

```c
    // GPUs. The STARNEIG_HINT_SM flag indicates that the library should
    // initialize itself for shared memory computations and the
    // STARNEIG_AWAKE_WORKERS indicates that the library should should keep
    // StarPU worker threads awake between interface function calls.
    starneig_node_init(STARNEIG_USE_ALL, STARNEIG_USE_ALL,
        STARNEIG_HINT_SM | STARNEIG_AWAKE_WORKERS);
    // reduce the dense-dense matrix pair (A,B) to Hessenberg-triangular form
    printf("Hessenberg-triangular reduction...\n");
    starneig_GEP_SM_HessenbergTriangular(n, A, ldA, B, ldB, Q, ldQ, Z, ldZ);
    // reduce the Hessenberg-triangular matrix pair (A,B) to generalized Schur
    // form
    printf("Schur reduction...\n");
    starneig_GEP_SM_Schur(n, A, ldA, B, ldB, Q, ldQ, Z, ldZ, real, imag, beta);
    // select eigenvalues that have positive a real part
    int num_selected;
    starneig_GEP_SM_Select(
        n, A, ldA, B, ldB, &predicate, NULL, select, &num_selected);
    printf("Selected %d eigenvalues out of %d.\n", num_selected, n);
    // reorder selected eigenvalues to the upper left corner of the generalized
    // Schur form (A,B)
    printf("Reordering...\n");
    starneig_GEP_SM_ReorderSchur(
        n, select, A, ldA, B, ldB, Q, ldQ, Z, ldZ, real, imag, beta);
    // de-initialize the StarNEig library
    starneig_node_finalize();
    // check residual || Q A Z^T - C ||_F / || C ||_F
    check_residual(n, ldQ, ldA, ldZ, ldC, Q, A, Z, C);
    // check residual || Q B Z^T - D ||_F / || D ||_F
    check_residual(n, ldQ, ldB, ldZ, ldD, Q, B, Z, D);
    // check residual || Q Q^T - I ||_F / || I ||_F
    check_orthogonality(n, ldQ, Q);
    // check residual || Z Z^T - I ||_F / || I ||_F
    check_orthogonality(n, ldZ, Z);
    // cleanup
    free(A);
    free(C);
    free(B);
    free(D);
    free(Q);
    free(Z);
    free(real);
    free(imag);
    free(beta);
    free(select);
    return 0;
}
```

15.4 sep_dm_full_chain.c

```c
#include "validate.h"
#include <stdlib.h>
#include <stdio.h>
#include <time.h>
#include <mpi.h>
#include <starneig/starneig.h>
// a predicate function that selects all eigenvalues that have positive a real
// part
static int predicate(double real, double imag, void *arg)
{
    if (0.0 < real)
        return 1;
    return 0;
}
int main(int argc, char **argv)
{
    const int n = 3000; // matrix dimension
    const int root = 0; // root rank
    // initialize MPI
    int thread_support;
    MPI_Init_thread(
        &argc, (char ***)&argv, MPI_THREAD_MULTIPLE, &thread_support);
    int world_rank;
    MPI_Comm_rank(MPI_COMM_WORLD, &world_rank);
    // the root node initializes the matrices locally
    int ldA = 0, ldQ = 0, ldC = 0;
    double *A = NULL, *Q = NULL, *C = NULL;
    if (world_rank == root) {
        srand((unsigned) time(NULL));
        // generate a full random matrix A and a copy C
        ldA = ((n/8)+1)*8; ldC = ((n/8)+1)*8;
        A = malloc(n*ldA*sizeof(double));
        C = malloc(n*ldC*sizeof(double));
        for (int j = 0; j < n; j++)
```

```c
            for (int i = 0; i < n; i++)
                A[j*ldA+i] = C[j*ldC+i] = 2.0*rand()/RAND_MAX - 1.0;
        // generate an identity matrix Q
        ldQ = ((n/8)+1)*8;
        Q = malloc(n*ldA*sizeof(double));
        for (int j = 0; j < n; j++)
            for (int i = 0; i < n; i++)
                Q[j*ldQ+i] = i == j ? 1.0 : 0.0;
    }
    // allocate space for the eigenvalues and the eigenvalue selection vector
    double *real = malloc(n*sizeof(double));
    double *imag = malloc(n*sizeof(double));
    int *select = malloc(n*sizeof(int));
    // Initialize the StarNEig library using all available CPU cores and
    // GPUs. The STARNEIG_HINT_DM flag indicates that the library should
    // initialize itself for distributed memory computations.
    starneig_node_init(STARNEIG_USE_ALL, STARNEIG_USE_ALL, STARNEIG_HINT_DM);
    // create a two-dimensional block cyclic distribution with row-major
    // ordering
    starneig_distr_t distr = starneig_distr_init_mesh(
        -1, -1, STARNEIG_ORDER_ROW_MAJOR);
    // Convert the local matrix A to a distributed matrix lA that is owned by
    // the root node. This is done in-place, i.e., the matrices A and lA point
    // to the same data.
    starneig_distr_matrix_t lA = starneig_distr_matrix_create_local(
        n, n, STARNEIG_REAL_DOUBLE, root, A, ldA);
    // create a distributed matrix dA using default data distribution and
    // distributed block size
    starneig_distr_matrix_t dA =
        starneig_distr_matrix_create(n, n, -1, -1, STARNEIG_REAL_DOUBLE, distr);
    // copy the local matrix lA to the distributed matrix dA (scatter)
    starneig_distr_matrix_copy(lA, dA);
    // scatter the matrix Q
    starneig_distr_matrix_t lQ = starneig_distr_matrix_create_local(
        n, n, STARNEIG_REAL_DOUBLE, root, Q, ldQ);
    starneig_distr_matrix_t dQ =
        starneig_distr_matrix_create(n, n, -1, -1, STARNEIG_REAL_DOUBLE, distr);
    starneig_distr_matrix_copy(lQ, dQ);
    // reduce the full matrix dA to upper Hessenberg form
    printf("Hessenberg reduction...\n");
    starneig_SEP_DM_Hessenberg(dA, dQ);
    // reduce the upper Hessenberg matrix dA to Schur form
    printf("Schur reduction...\n");
    starneig_SEP_DM_Schur(dA, dQ, real, imag);
    // select eigenvalues that have positive a real part
    int num_selected;
    starneig_SEP_DM_Select(dA, &predicate, NULL, select, &num_selected);
    printf("Selected %d eigenvalues out of %d.\n", num_selected, n);
    // reorder the selected eigenvalues to the upper left corner of the matrix
    // dA
    printf("Reordering...\n");
    starneig_SEP_DM_ReorderSchur(select, dA, dQ, real, imag);
    // copy the distributed matrix dA back to the local matrix lA (gather)
    starneig_distr_matrix_copy(dA, lA);
    // free the distributed matrix lA (matrix A is not freed)
    starneig_distr_matrix_destroy(lA);
    // free the distributed matrix dA (all local resources are freed)
    starneig_distr_matrix_destroy(dA);
    // gather the matrix Q
    starneig_distr_matrix_copy(dQ, lQ);
    starneig_distr_matrix_destroy(lQ);
    starneig_distr_matrix_destroy(dQ);
    // free the data distribution
    starneig_distr_destroy(distr);
    // de-initialize the StarNEig library
    starneig_node_finalize();
    // de-initialize MPI
    MPI_Finalize();
    if (world_rank == root) {
        // check residual || Q A Q^T - C ||_F / || C ||_F
        check_residual(n, ldQ, ldA, ldQ, ldC, Q, A, Q, C);
        // check residual || Q Q^T - I ||_F / || I ||_F
        check_orthogonality(n, ldQ, Q);
    }
    // cleanup
    free(A);
    free(C);
    free(Q);
    free(real);
    free(imag);
    free(select);
    return 0;
}
```

15.5 sep_sm_eigenvectors.c

```c
#include "validate.h"
#include <stdlib.h>
#include <stdio.h>
#include <time.h>
#include <starneig/starneig.h>
// a predicate function that selects all eigenvalues that have positive a real
// part
static int predicate(double real, double imag, void *arg)
{
    if (0.0 < real)
        return 1;
    return 0;
}
int main()
{
    const int n = 3000; // matrix dimension
    srand((unsigned) time(NULL));
    // generate a full random matrix A and a copy C
    int ldA = ((n/8)+1)*8, ldC = ((n/8)+1)*8;
    double *A = malloc(n*ldA*sizeof(double));
    double *C = malloc(n*ldC*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            A[j*ldA+i] = C[j*ldC+i] = 2.0*rand()/RAND_MAX - 1.0;
    // generate an identity matrix Q
    int ldQ = ((n/8)+1)*8;
    double *Q = malloc(n*ldA*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            Q[j*ldQ+i] = i == j ? 1.0 : 0.0;
    double *X = NULL; int ldX = 0;
    // allocate space for the eigenvalues and the eigenvector selection vector
    double *real = malloc(n*sizeof(double));
    double *imag = malloc(n*sizeof(double));
    int *select = malloc(n*sizeof(int));
    // Initialize the StarNEig library using all available CPU cores and
    // GPUs. The STARNEIG_HINT_SM flag indicates that the library should
    // initialize itself for shared memory computations.
    starneig_node_init(STARNEIG_USE_ALL, STARNEIG_USE_ALL, STARNEIG_HINT_SM);
    // reduce the full matrix matrix A to Schur form (skip reordering)
    printf("Reduce...\n");
    starneig_SEP_SM_Reduce(
        n, A, ldA, Q, ldQ, real, imag, NULL, NULL, NULL, NULL);
    // select eigenvalues that have positive a real part and allocate space for
    // the eigenvectors
    int num_selected;
    starneig_SEP_SM_Select(n, A, ldA, &predicate, NULL, select, &num_selected);
    printf("Selected %d eigenvalues out of %d.\n", num_selected, n);
    ldX = ((n/8)+1)*8;
    X = malloc(num_selected*ldX*sizeof(double));
    // compute a selected set of eigenvectors
    printf("Eigenvectors...\n");
    starneig_SEP_SM_Eigenvectors(n, select, A, ldA, Q, ldQ, X, ldX);
    // de-initialize the StarNEig library
    starneig_node_finalize();
    // check residual || Q A Q^T - C ||_F / || C ||_F
    check_residual(n, ldQ, ldA, ldQ, ldC, Q, A, Q, C);
    // check residual || Q Q^T - I ||_F / || I ||_F
    check_orthogonality(n, ldQ, Q);
    // cleanup
    free(A);
    free(C);
    free(Q);
    free(X);
    free(real);
    free(imag);
    free(select);
    return 0;
}
```

15.6 sep_sm_full_chain.c

```c
#include "validate.h"
#include <stdlib.h>
#include <stdio.h>
#include <time.h>
#include <starneig/starneig.h>
// a predicate function that selects all eigenvalues that have positive a real
// part
static int predicate(double real, double imag, void *arg)
{
```

```c
        if (0.0 < real)
            return 1;
        return 0;
}
int main()
{
    const int n = 3000; // matrix dimension
    srand((unsigned) time(NULL));
    // generate a full random matrix A and a copy C
    int ldA = ((n/8)+1)*8, ldC = ((n/8)+1)*8;
    double *A = malloc(n*ldA*sizeof(double));
    double *C = malloc(n*ldC*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            A[j*ldA+i] = C[j*ldC+i] = 2.0*rand()/RAND_MAX - 1.0;
    // generate an identity matrix Q
    int ldQ = ((n/8)+1)*8;
    double *Q = malloc(n*ldA*sizeof(double));
    for (int j = 0; j < n; j++)
        for (int i = 0; i < n; i++)
            Q[j*ldQ+i] = i == j ? 1.0 : 0.0;
    // allocate space for the eigenvalues and the eigenvalue selection vector
    double *real = malloc(n*sizeof(double));
    double *imag = malloc(n*sizeof(double));
    int *select = malloc(n*sizeof(int));
    // Initialize the StarNEig library using all available CPU cores and
    // GPUs. The STARNEIG_HINT_SM flag indicates that the library should
    // initialize itself for shared memory computations.
    starneig_node_init(STARNEIG_USE_ALL, STARNEIG_USE_ALL, STARNEIG_HINT_SM);
    // reduce the full matrix matrix A to upper Hessenberg form
    printf("Hessenberg reduction...\n");
    starneig_SEP_SM_Hessenberg(n, A, ldA, Q, ldQ);
    // reduce the upper Hessenberg matrix A to Schur form
    printf("Schur reduction...\n");
    starneig_SEP_SM_Schur(n, A, ldA, Q, ldQ, real, imag);
    // select eigenvalues that have positive a real part
    int num_selected;
    starneig_SEP_SM_Select(n, A, ldA, &predicate, NULL, select, &num_selected);
    printf("Selected %d eigenvalues out of %d.\n", num_selected, n);
    // reorder the selected eigenvalues to the upper left corner of the matrix A
    printf("Reordering...\n");
    starneig_SEP_SM_ReorderSchur(n, select, A, ldA, Q, ldQ, real, imag);
    // de-initialize the StarNEig library
    starneig_node_finalize();
    // check residual || Q A Q^T - C ||_F / || C ||_F
    check_residual(n, ldQ, ldA, ldQ, ldC, Q, A, Q, C);
    // check residual || Q Q^T - I ||_F / || I ||_F
    check_orthogonality(n, ldQ, Q);
    // cleanup
    free(A);
    free(C);
    free(Q);
    free(real);
    free(imag);
    free(select);
    return 0;
}
double *A = malloc(n*ldA*sizeof(double));
```

Index